

# Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

14-513

18-613

# Virtual Memory: Systems

Assignment Project Exam Help

15-213/18-213/14-513/15-513/18-613:

Introduction to Com

18<sup>th</sup> Lecture, Octobe

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Announcements

## ■ Lab 5 (malloclab)

- Checkpoint due Thu, Oct. 29, 11:59pm ET

## ■ Written Assignment 7 peer grading

- Due Wed, Nov. 4, 11:59pm ET

## ■ Written Assign

- Due Wed, Nov.

## ■ Recitation on Malloclab (part II)

- Mon, Nov. 2. Slides are already posted

## ■ U.S. Election Day is Tues, Nov.3

- If eligible, go VOTE!
- Skip class if need be (NO QUIZ on TUES!)

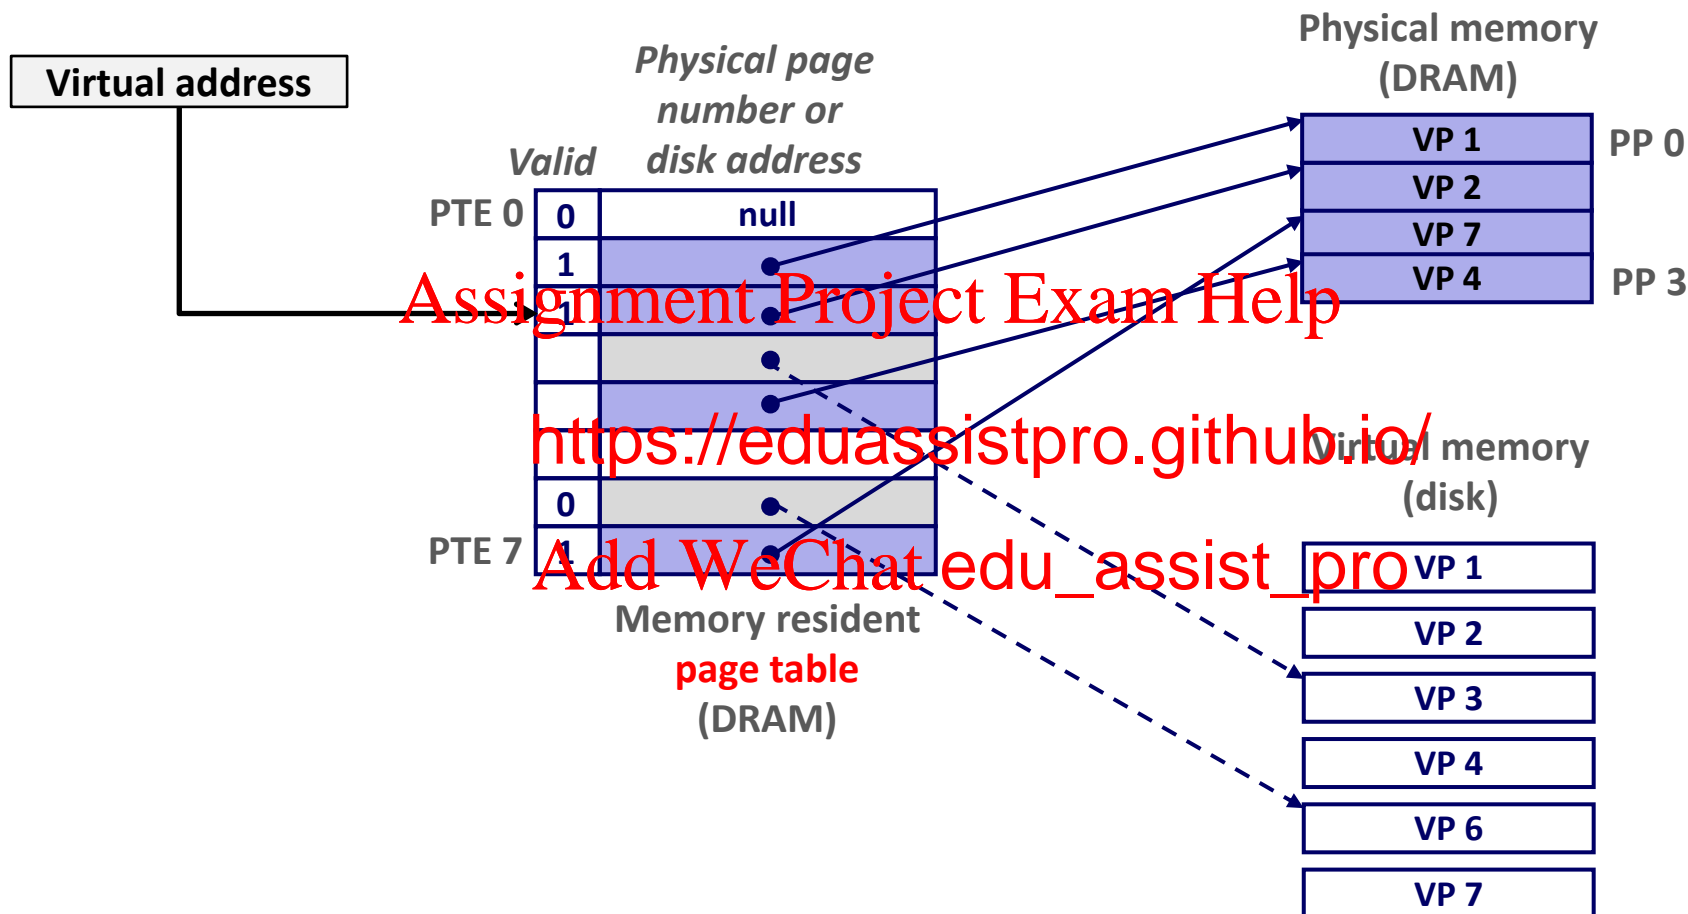


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu\_assist\_pro

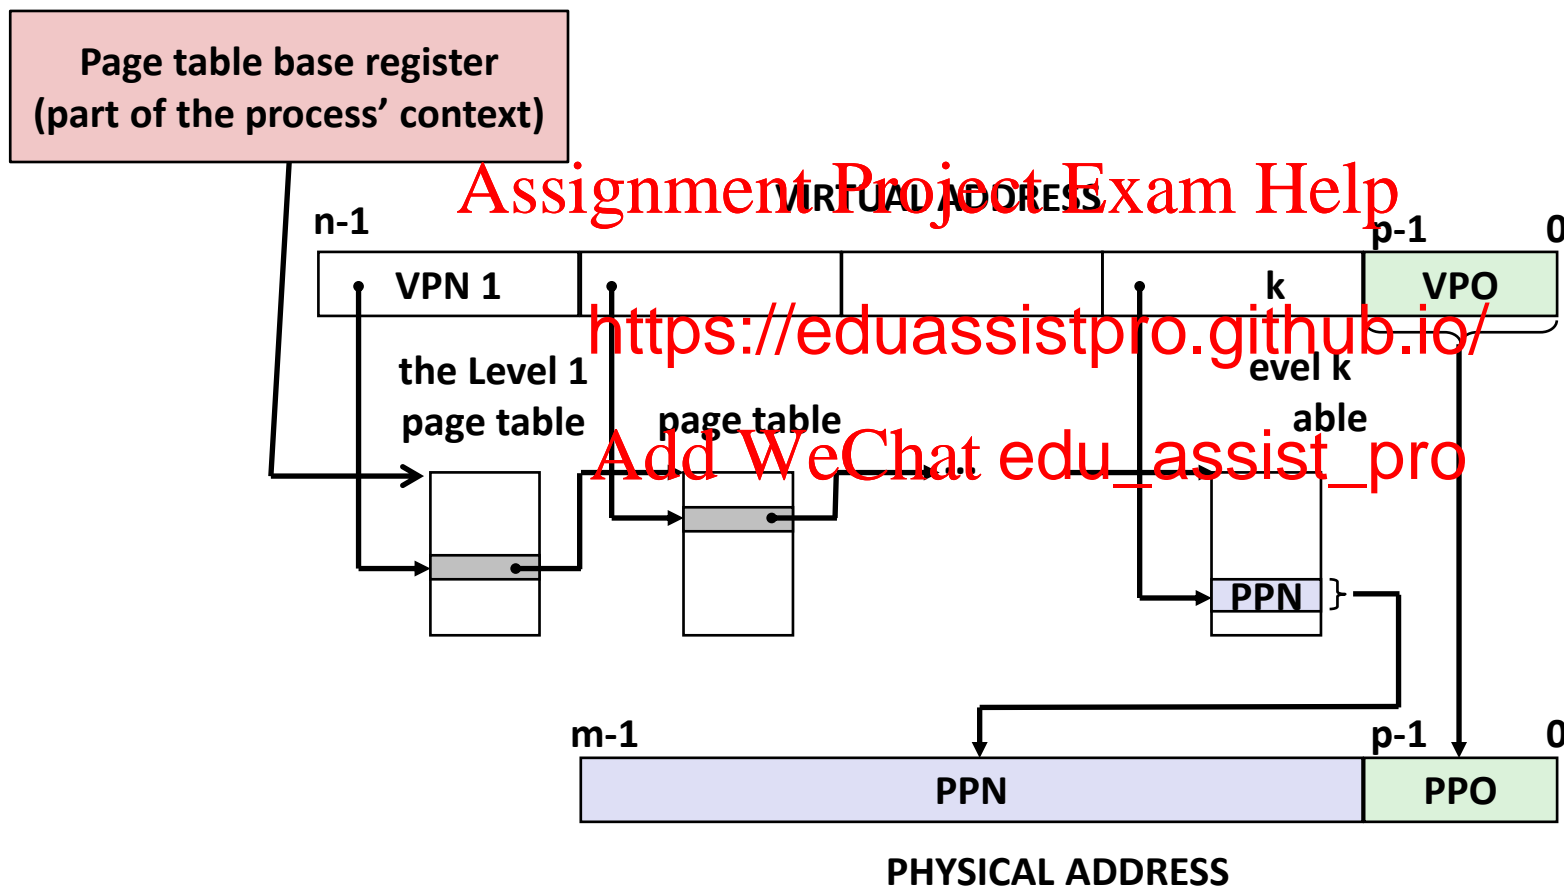
# Review: Virtual Memory & Physical Memory



- A **page table** contains page table entries (PTEs) that map virtual pages to physical pages.

# Translating with a k-level Page Table

- Having multiple levels greatly reduces page table size



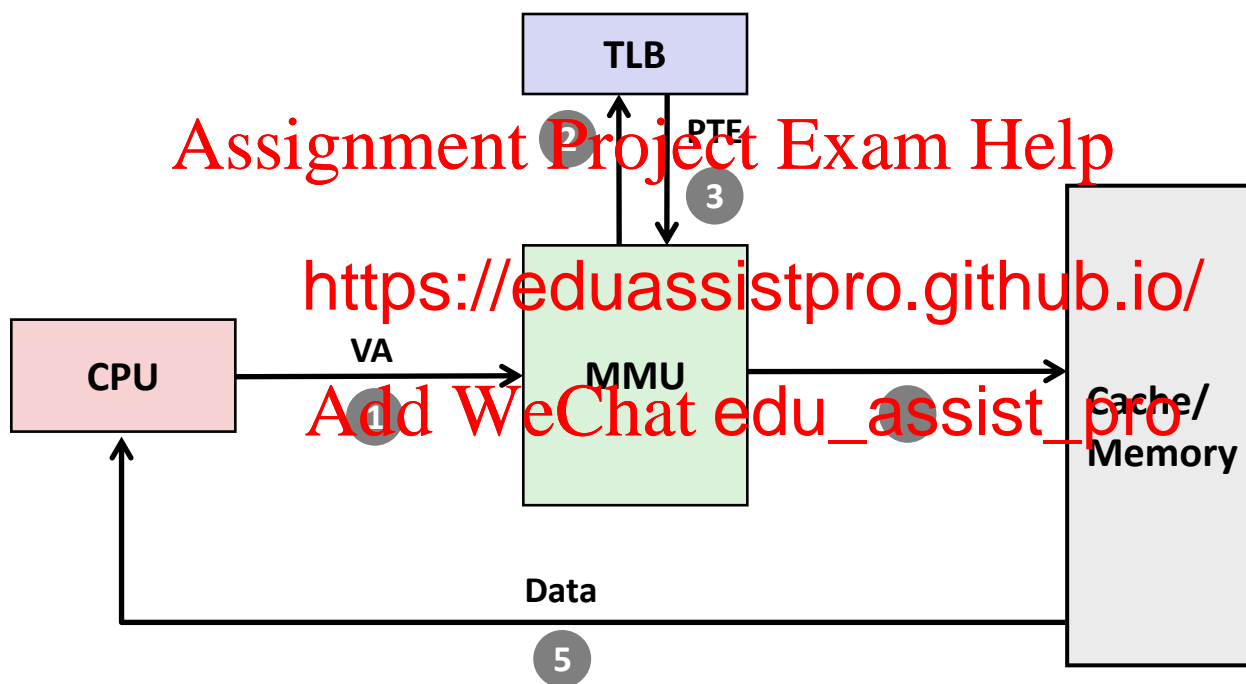
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Translation Lookaside Buffer (TLB)

- A small cache of page table entries with fast access by MMU

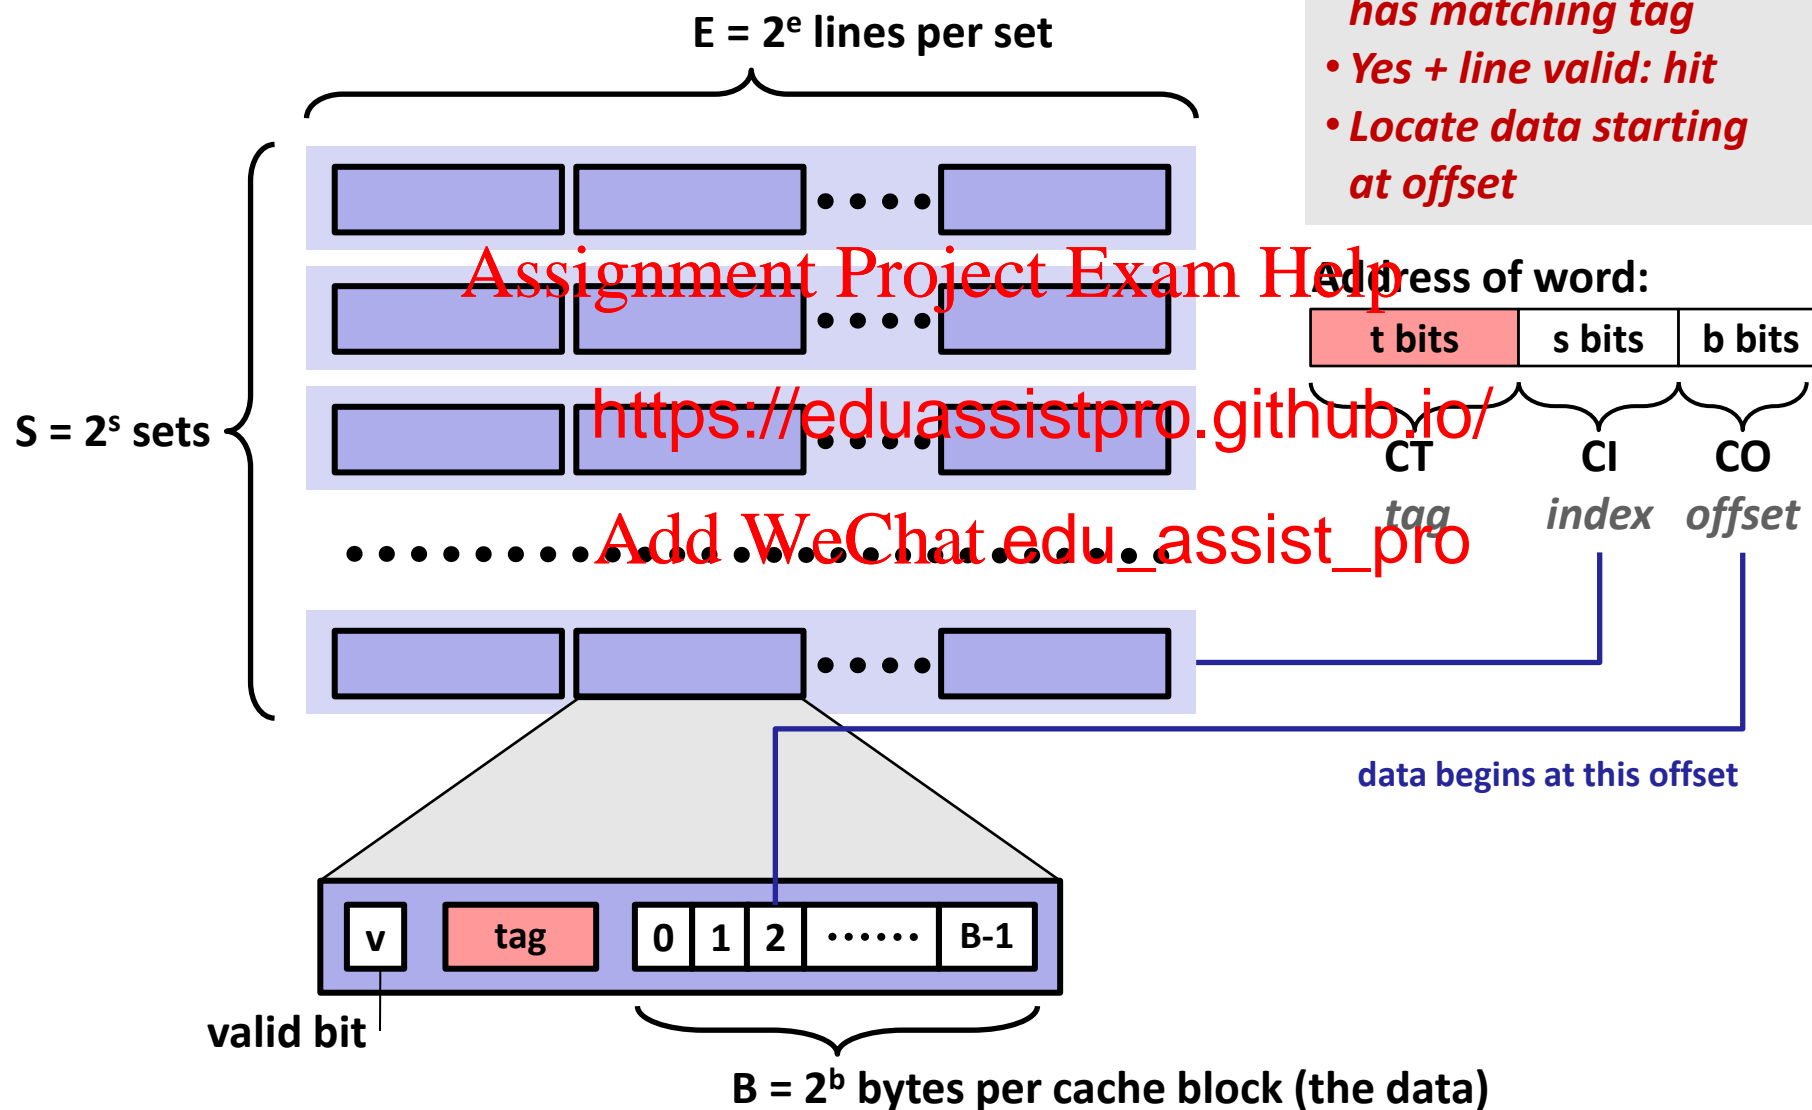


Typically, a **TLB hit** eliminates the  $k$  memory accesses required to do a page table lookup.

# Recall: Set Associative Cache

## Steps for a READ:

- *Locate set*
- *Check if any line in set has matching tag*
- *Yes + line valid: hit*
- *Locate data starting at offset*



# Review of Symbols

## ■ Basic Parameters

- $N = 2^n$ : Number of addresses in virtual address space
- $M = 2^m$ : Number of addresses in physical address space
- $P = 2^p$ : Page size (bytes)

## ■ Components of the **virtual address**

- TLBI: TLB index
- TLBT: TLB tag
- VPO: Virtual page offset
- VPN: Virtual page number

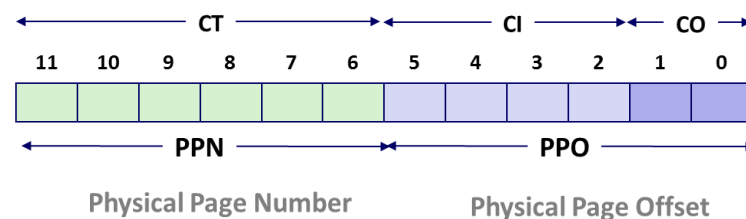
## ■ Components of the **physical address (PA)**

- PPO: Physical page offset (same as VPO)
- PPN: Physical page number
- CO: Byte offset within cache line
- CI: Cache index
- CT: Cache tag

Virtual Page Number

Virtual Page Offset

(bits per field for our simple example)



Physical Page Number

Physical Page Offset

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Today

- Simple memory system example CSAPP 9.6.4
- Case study: Core i7/Linux memory system CSAPP 9.7
- Memory mapping CSAPP 9.8

Assignment Project Exam Help

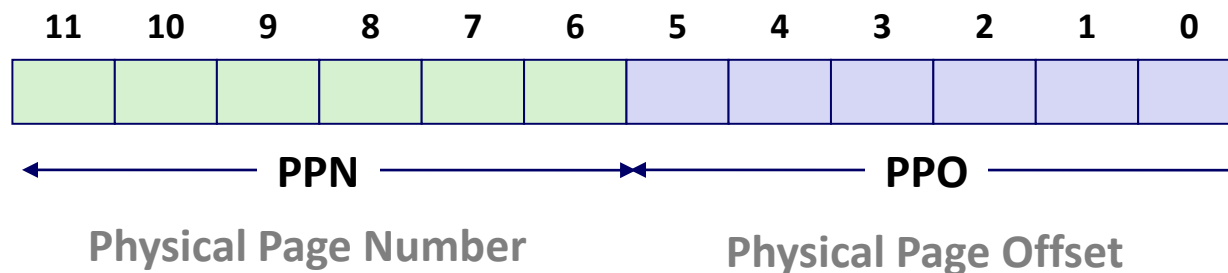
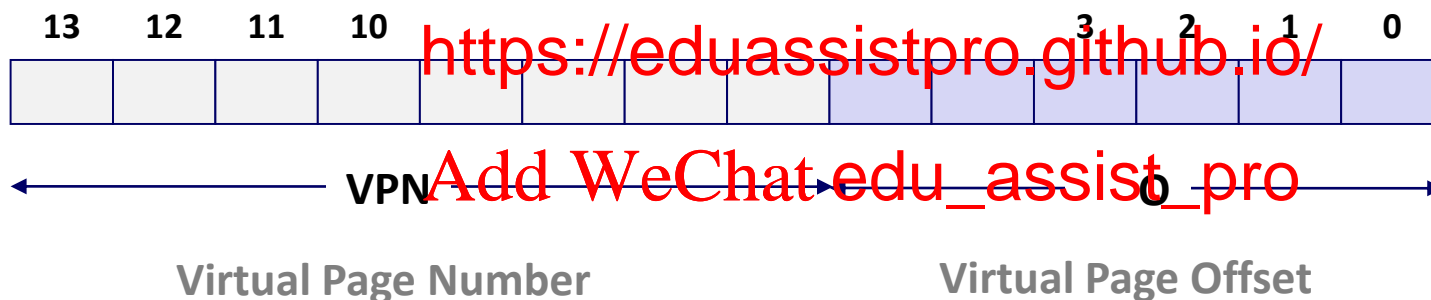
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Simple Memory System Example

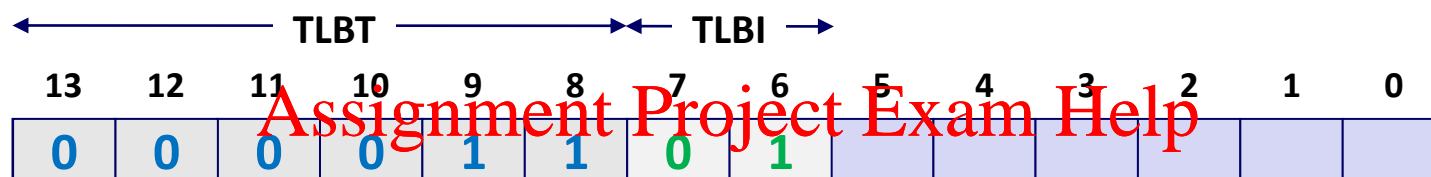
## ■ Addressing

- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



# Simple Memory System TLB

- 16 entries
- 4-way associative



VPN = 0b1101 = 0xD

## Translation Lookaside Buffer (TLB)

| Set | Tag | PPN | Valid | Tag | PPN | Valid | Tag | PPN | Valid | Tag | PPN | Valid |
|-----|-----|-----|-------|-----|-----|-------|-----|-----|-------|-----|-----|-------|
| 0   | 03  | –   | 0     | 09  | 0D  | 1     | 00  | –   | 0     | 07  | 02  | 1     |
| 1   | 03  | 2D  | 1     | 02  | –   | 0     | 04  | –   | 0     | 0A  | –   | 0     |
| 2   | 02  | –   | 0     | 08  | –   | 0     | 06  | –   | 0     | 03  | –   | 0     |
| 3   | 07  | –   | 0     | 03  | 0D  | 1     | 0A  | 34  | 1     | 02  | –   | 0     |

# Simple Memory System Page Table

Only showing the first 16 entries (out of 256)

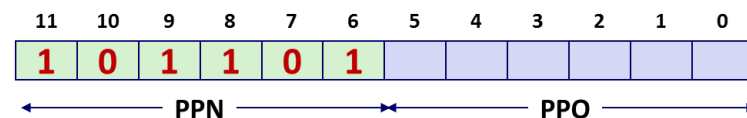
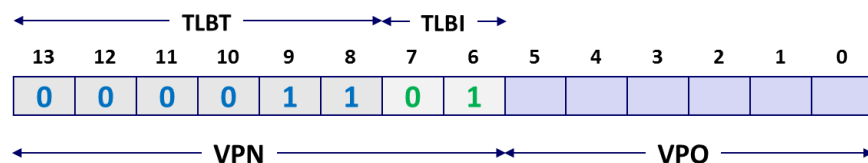
| VPN | PPN | Valid |
|-----|-----|-------|
| 00  | 28  | 1     |
| 01  | —   | 0     |
| 02  |     |       |
| 03  |     |       |
| 04  | —   | 0     |
| 05  | 16  | 1     |
| 06  | —   | 0     |
| 07  | —   | 0     |

| VPN | PPN | Valid |
|-----|-----|-------|
| 08  | 13  | 1     |
| 09  | 17  | 1     |
|     |     | 1     |
|     |     | 0     |
|     |     | 0     |
|     |     | 1     |
| 0E  | 11  | 1     |
| 0F  | 0D  | 1     |

<https://eduassistpro.github.io/>

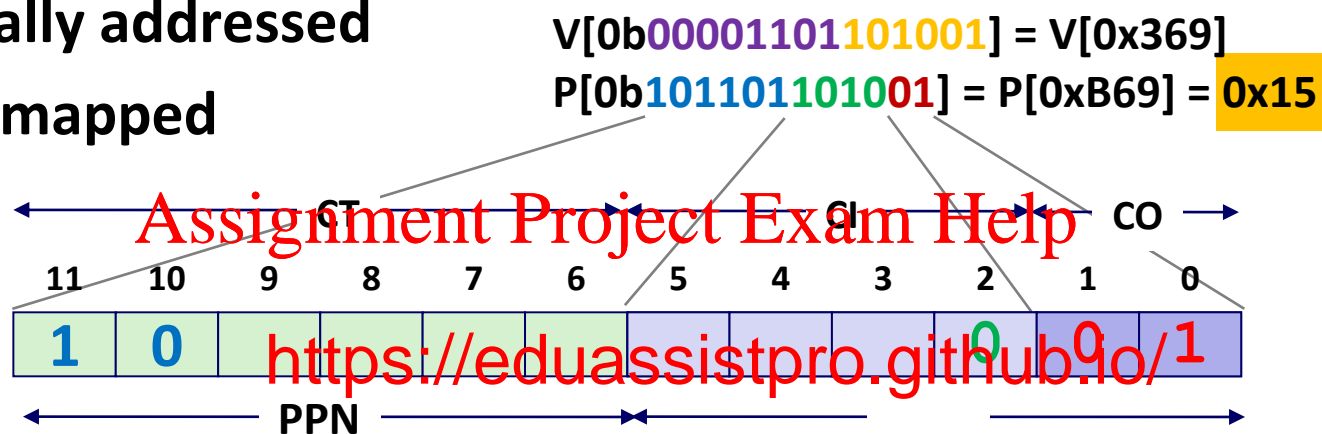
Add WeChat edu\_assist\_pro

0x0D → 0x2D



# Simple Memory System Cache

- 16 lines, 4-byte cache line size
- Physically addressed
- Direct mapped



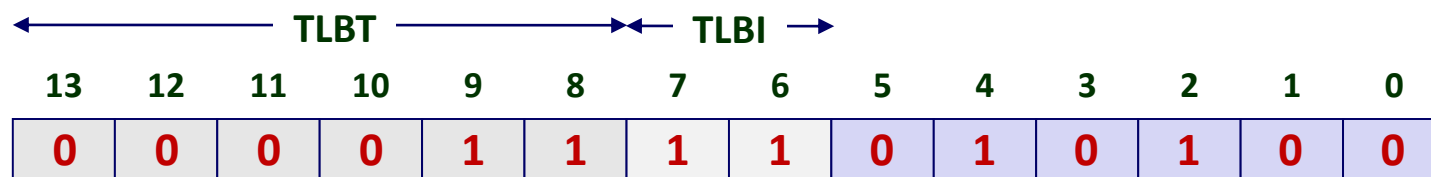
Add WeChat edu\_assist\_pro

| Idx | Tag | Valid | B0 | B1 | B2 | B3 |
|-----|-----|-------|----|----|----|----|
| 0   | 19  | 1     | 99 | 11 | 23 | 11 |
| 1   | 15  | 0     | —  | —  | —  | —  |
| 2   | 1B  | 1     | 00 | 02 | 04 | 08 |
| 3   | 36  | 0     | —  | —  | —  | —  |
| 4   | 32  | 1     | 43 | 6D | 8F | 09 |
| 5   | 0D  | 1     | 36 | 72 | F0 | 1D |
| 6   | 31  | 0     | —  | —  | —  | —  |
| 7   | 16  | 1     | 11 | C2 | DF | 03 |

|   |    | id | B0 | B1 | B2 | B3 |
|---|----|----|----|----|----|----|
| 8 | 24 | 1  | 3A | 00 | 51 | 89 |
| 9 | 2D | 0  | —  | —  | —  | —  |
| A | 2D | 1  | 93 | 15 | DA | 3B |
| B | 0B | 0  | —  | —  | —  | —  |
| C | 12 | 0  | —  | —  | —  | —  |
| D | 16 | 1  | 04 | 96 | 34 | 15 |
| E | 13 | 1  | 83 | 77 | 1B | D3 |
| F | 14 | 0  | —  | —  | —  | —  |

# Address Translation Example

Virtual Address: 0x03D4



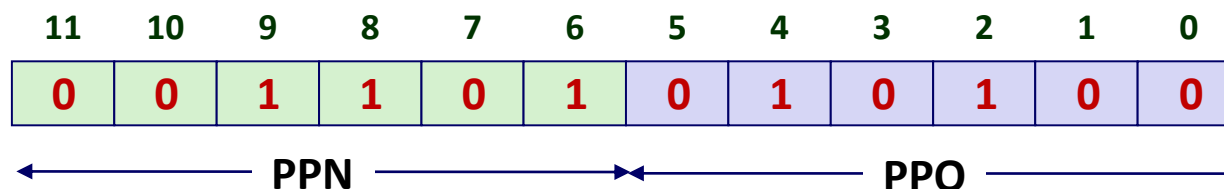
Assignment Project Exam Help

VPN: 0x0F TLBI: 0x Page Fault? N PPN: 0x0D

TLB

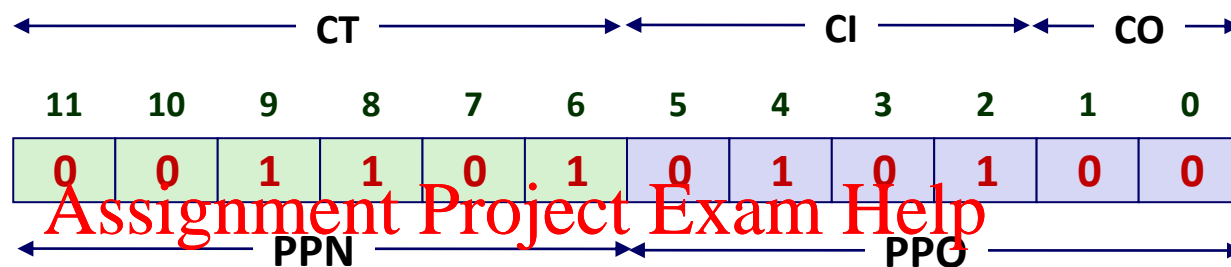
| Set | Tag | PPN | Valid | Tag | PPN | Valid |    |    | Valid | Tag | PPN | Valid |
|-----|-----|-----|-------|-----|-----|-------|----|----|-------|-----|-----|-------|
| 0   | 03  | –   | 0     | 09  | 0D  | 1     |    |    |       | 07  | 02  | 1     |
| 1   | 03  | 2D  | 1     | 02  | –   | 0     | 04 | –  | 0     | 0A  | –   | 0     |
| 2   | 02  | –   | 0     | 08  | –   | 0     | 06 | –  | 0     | 03  | –   | 0     |
| 3   | 07  | –   | 0     | 03  | 0D  | 1     | 0A | 34 | 1     | 02  | –   | 0     |

Physical Address



# Address Translation Example

## Physical Address

CO 0CI 0x5
<https://eduassistpro.github.io/>

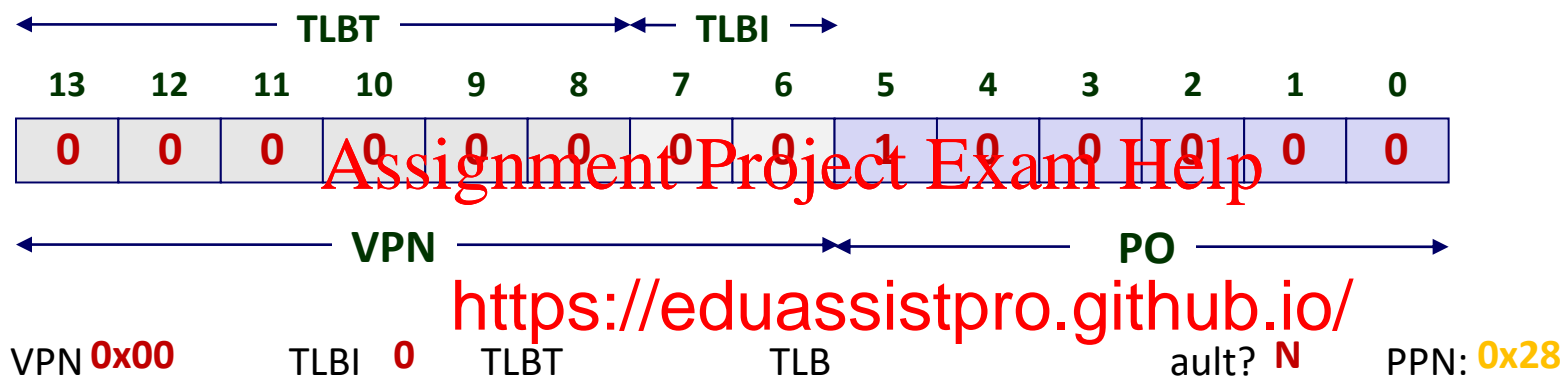
## Cache

| Idx | Tag | Valid | B0 | B1 | B2 | B3 |
|-----|-----|-------|----|----|----|----|
| 0   | 19  | 1     | 99 | 11 | 23 | 11 |
| 1   | 15  | 0     | —  | —  | —  | —  |
| 2   | 1B  | 1     | 00 | 02 | 04 | 08 |
| 3   | 36  | 0     | —  | —  | —  | —  |
| 4   | 32  | 1     | 43 | 6D | 8F | 09 |
| 5   | 0D  | 1     | 36 | 72 | F0 | 1D |
| 6   | 31  | 0     | —  | —  | —  | —  |
| 7   | 16  | 1     | 11 | C2 | DF | 03 |

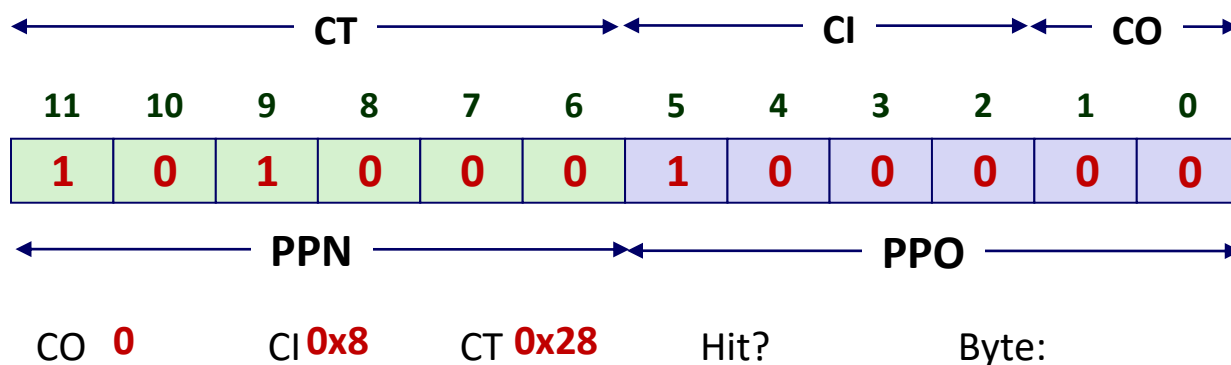
| Idx | Tag | Valid | B0 | B1 | B2 | B3 |
|-----|-----|-------|----|----|----|----|
| 8   | 24  | 1     | 3A | 00 | 51 | 89 |
| 9   | 2D  | 0     | —  | —  | —  | —  |
| A   | 2D  | 1     | 93 | 15 | DA | 3B |
| B   | 0B  | 0     | —  | —  | —  | —  |
| C   | 12  | 0     | —  | —  | —  | —  |
| D   | 16  | 1     | 04 | 96 | 34 | 15 |
| E   | 13  | 1     | 83 | 77 | 1B | D3 |
| F   | 14  | 0     | —  | —  | —  | —  |

# Address Translation Example: TLB/Cache Miss

Virtual Address: 0x0020



Physical Address



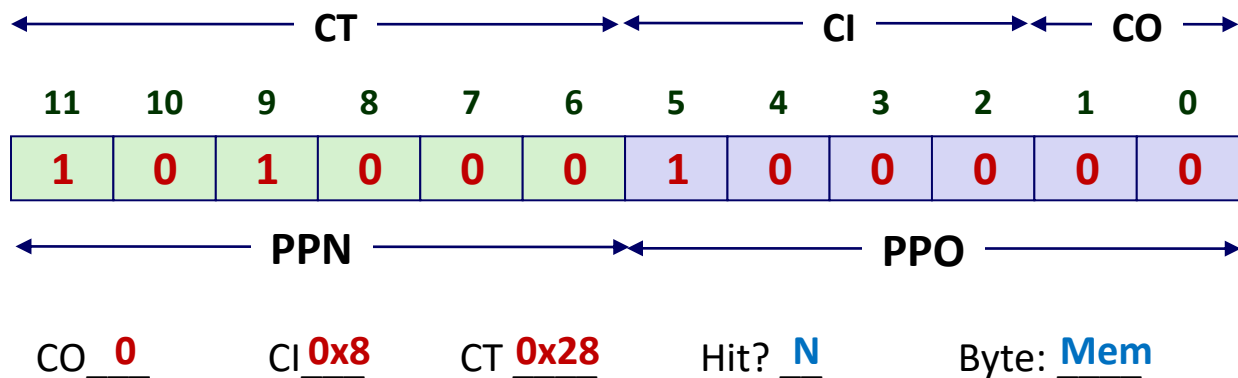
Page table

| VPN | PPN | Valid |
|-----|-----|-------|
| 00  | 28  | 1     |
| 01  | –   | 0     |
| 02  | 33  | 1     |
| 03  | 02  | 1     |
| 04  | –   | 0     |
| 05  | 16  | 1     |
| 06  | –   | 0     |
| 07  | –   | 0     |



| Idx | Tag | Valid | B0 | B1 | B2 | B3 |
|-----|-----|-------|----|----|----|----|
| 0   | 19  | 1     | 99 | 11 | 23 | 11 |
| 1   | 15  | 0     | –  | –  | –  | –  |
| 2   | 1B  | 1     | 00 | 02 | 04 | 08 |
| 3   | 36  | 0     | –  | –  | –  | –  |
| 4   | 32  | 1     | 43 |    |    |    |
| 5   | 0D  | 1     | 36 |    |    |    |
| 6   | 31  | 0     | –  | –  | –  | –  |
| 7   | 16  | 1     | 11 | C2 | 0F | 03 |

## Physical Address



# Quiz Time!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Check out:

Add WeChat edu\_assist\_pro

<https://canvas.cmu.edu/courses/17808>

# Today

- Simple memory system example
- **Case study: Core i7/Linux memory system**
- Memory mapping

Assignment Project Exam Help

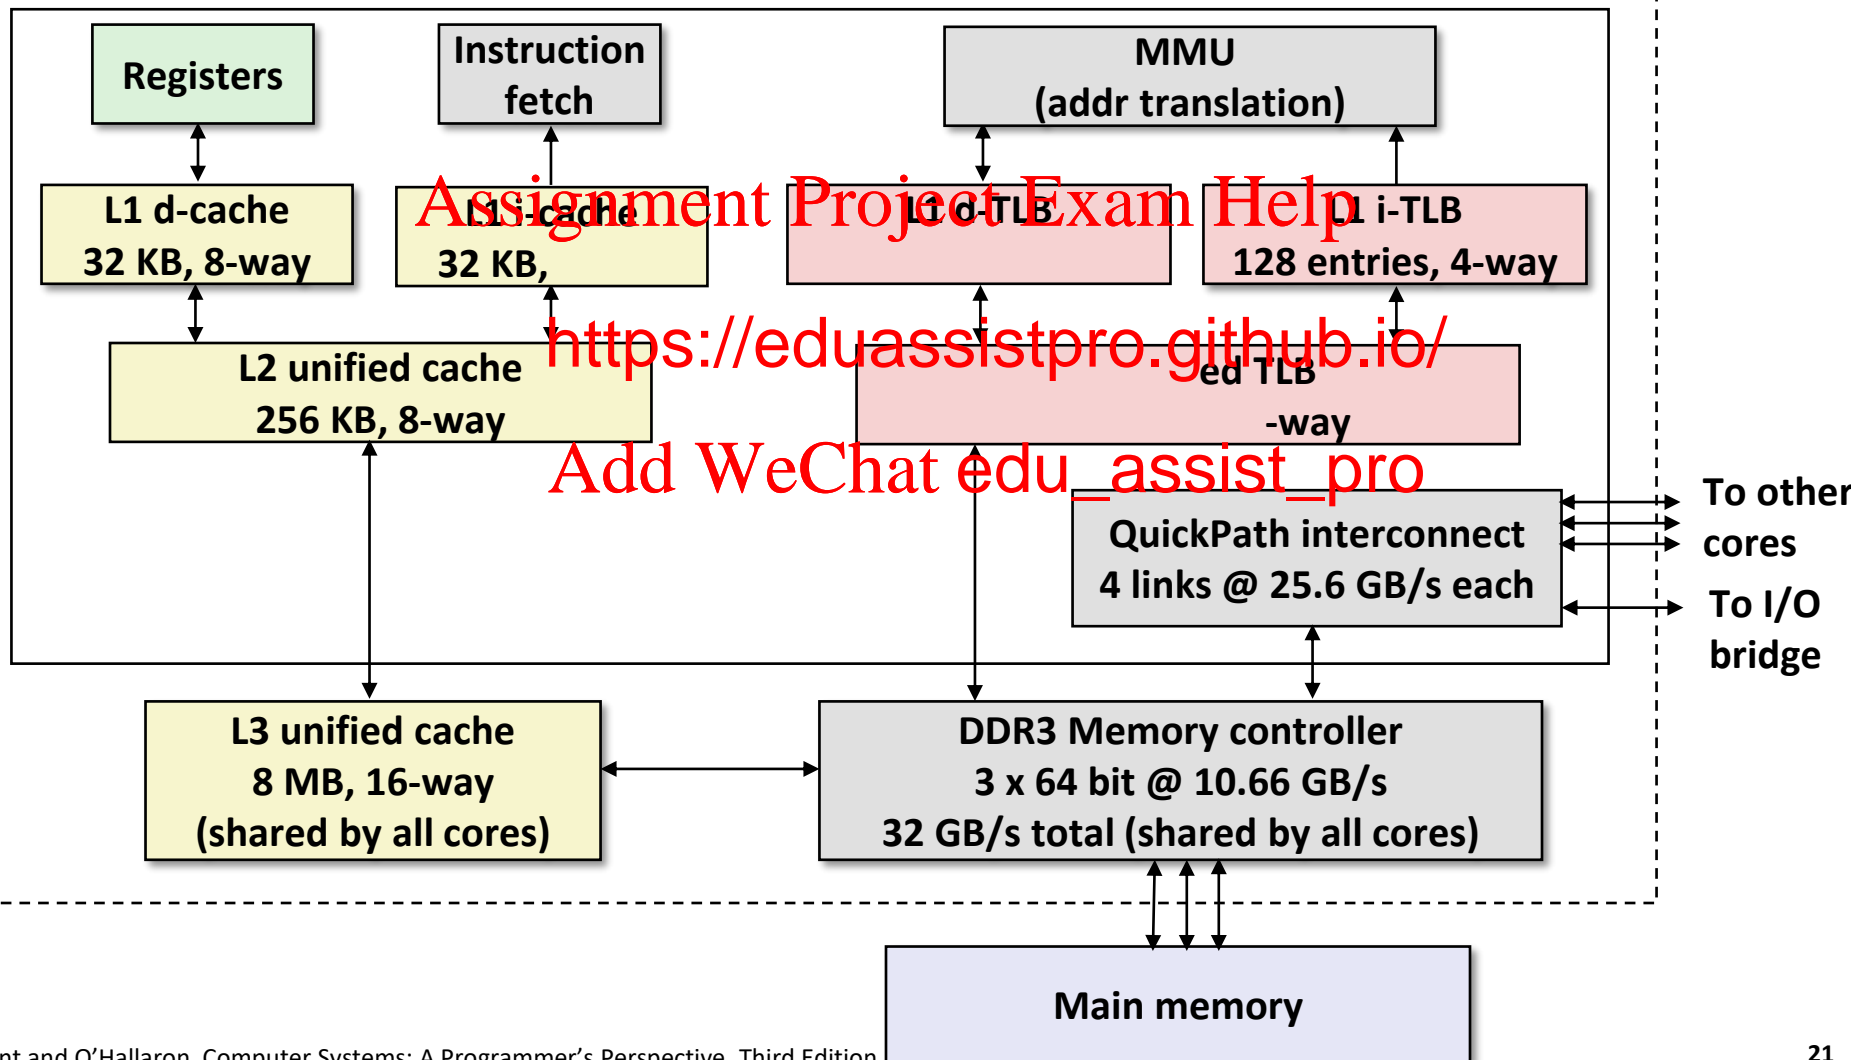
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

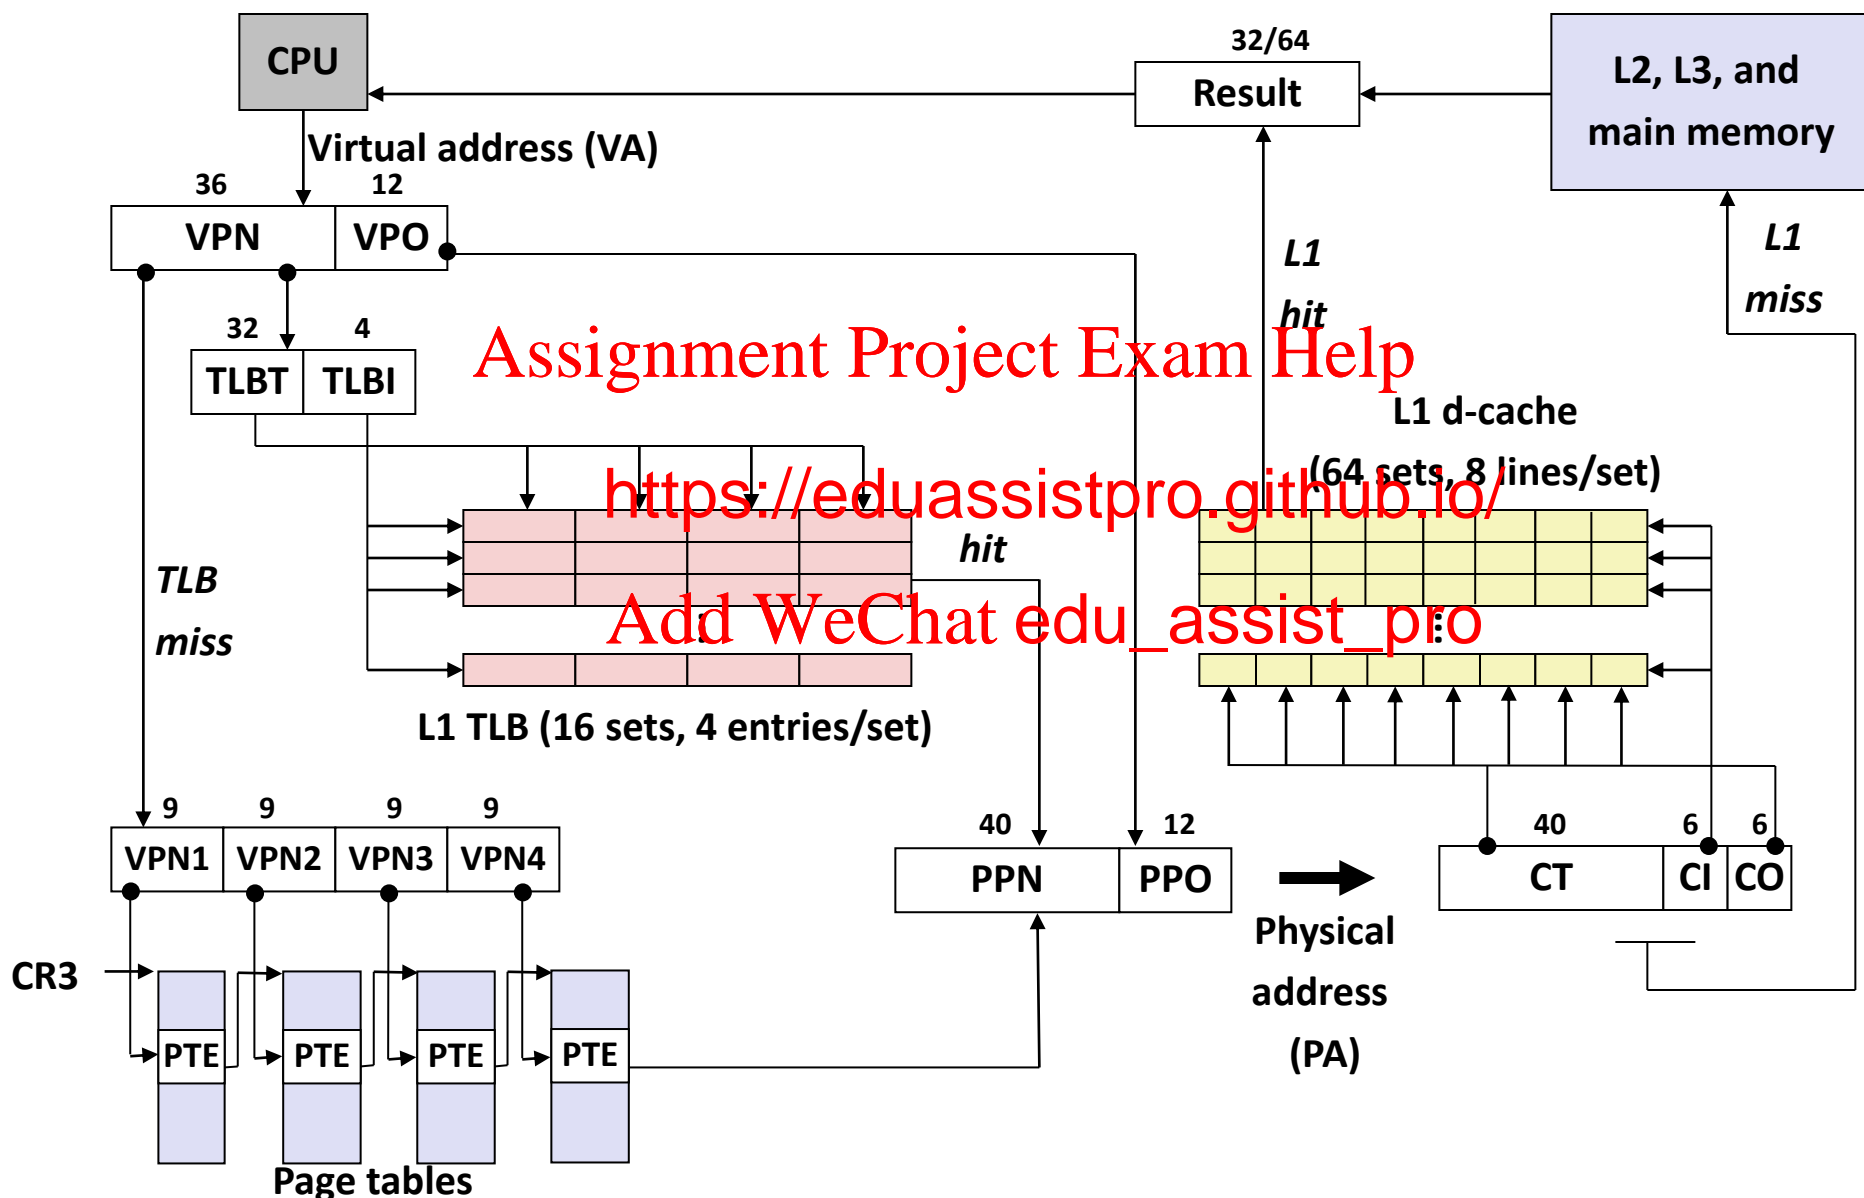
# Intel Core i7 Memory System

Processor package

Core x4



# End-to-end Core i7 Address Translation



# Core i7 Level 1-3 Page Table Entries

| 63                                             | 62     | 52                               | 51 | 12 | 11 | 9      | 8 | 7  | 6 | 5 | 4  | 3  | 2   | 1   | 0   |
|------------------------------------------------|--------|----------------------------------|----|----|----|--------|---|----|---|---|----|----|-----|-----|-----|
| XD                                             | Unused | Page table physical base address |    |    |    | Unused | G | PS |   | A | CD | WT | U/S | R/W | P=1 |
| Available for OS (page table location on disk) |        |                                  |    |    |    |        |   |    |   |   |    |    |     |     | P=0 |

## Assignment Project Exam Help

Each entry references a 4K child page table. Significant fields:

**P:** Child page table present in <https://eduassistpro.github.io/>

**R/W:** Read-only or read-write access access permissions for the pages.

**U/S:** user or supervisor (kernel) mode access permissions for the pages. [Add WeChat edu\\_assist\\_pro](#)

**WT:** Write-through or write-back cache policy for the child page table.

**A:** Reference bit (set by MMU on reads and writes, cleared by software).

**PS:** Page size either 4 KB or 4 MB (defined for Level 1 PTEs only).

**Page table physical base address:** 40 most significant bits of physical page table address (forces page tables to be 4KB aligned)

**XD:** Disable or enable instruction fetches from all pages reachable from this PTE.

# Core i7 Level 4 Page Table Entries

| 63                                       | 62     | 52                         | 51 | 12 | 11 | 9      | 8 | 7 | 6 | 5 | 4  | 3  | 2   | 1   | 0   |
|------------------------------------------|--------|----------------------------|----|----|----|--------|---|---|---|---|----|----|-----|-----|-----|
| XD                                       | Unused | Page physical base address |    |    |    | Unused | G |   | D | A | CD | WT | U/S | R/W | P=1 |
| Available for OS (page location on disk) |        |                            |    |    |    |        |   |   |   |   |    |    |     |     | P=0 |

Assignment Project Exam Help

Each entry references a 4K child page. Significant fields:

**P:** Child page is present in memory <https://eduassistpro.github.io/>

**R/W:** Read-only or read-write access permission for

**U/S:** User or supervisor mode access [Add WeChat edu\\_assist\\_pro](#)

**WT:** Write-through or write-back cache policy for this page

**A:** Reference bit (set by MMU on reads and writes, cleared by software)

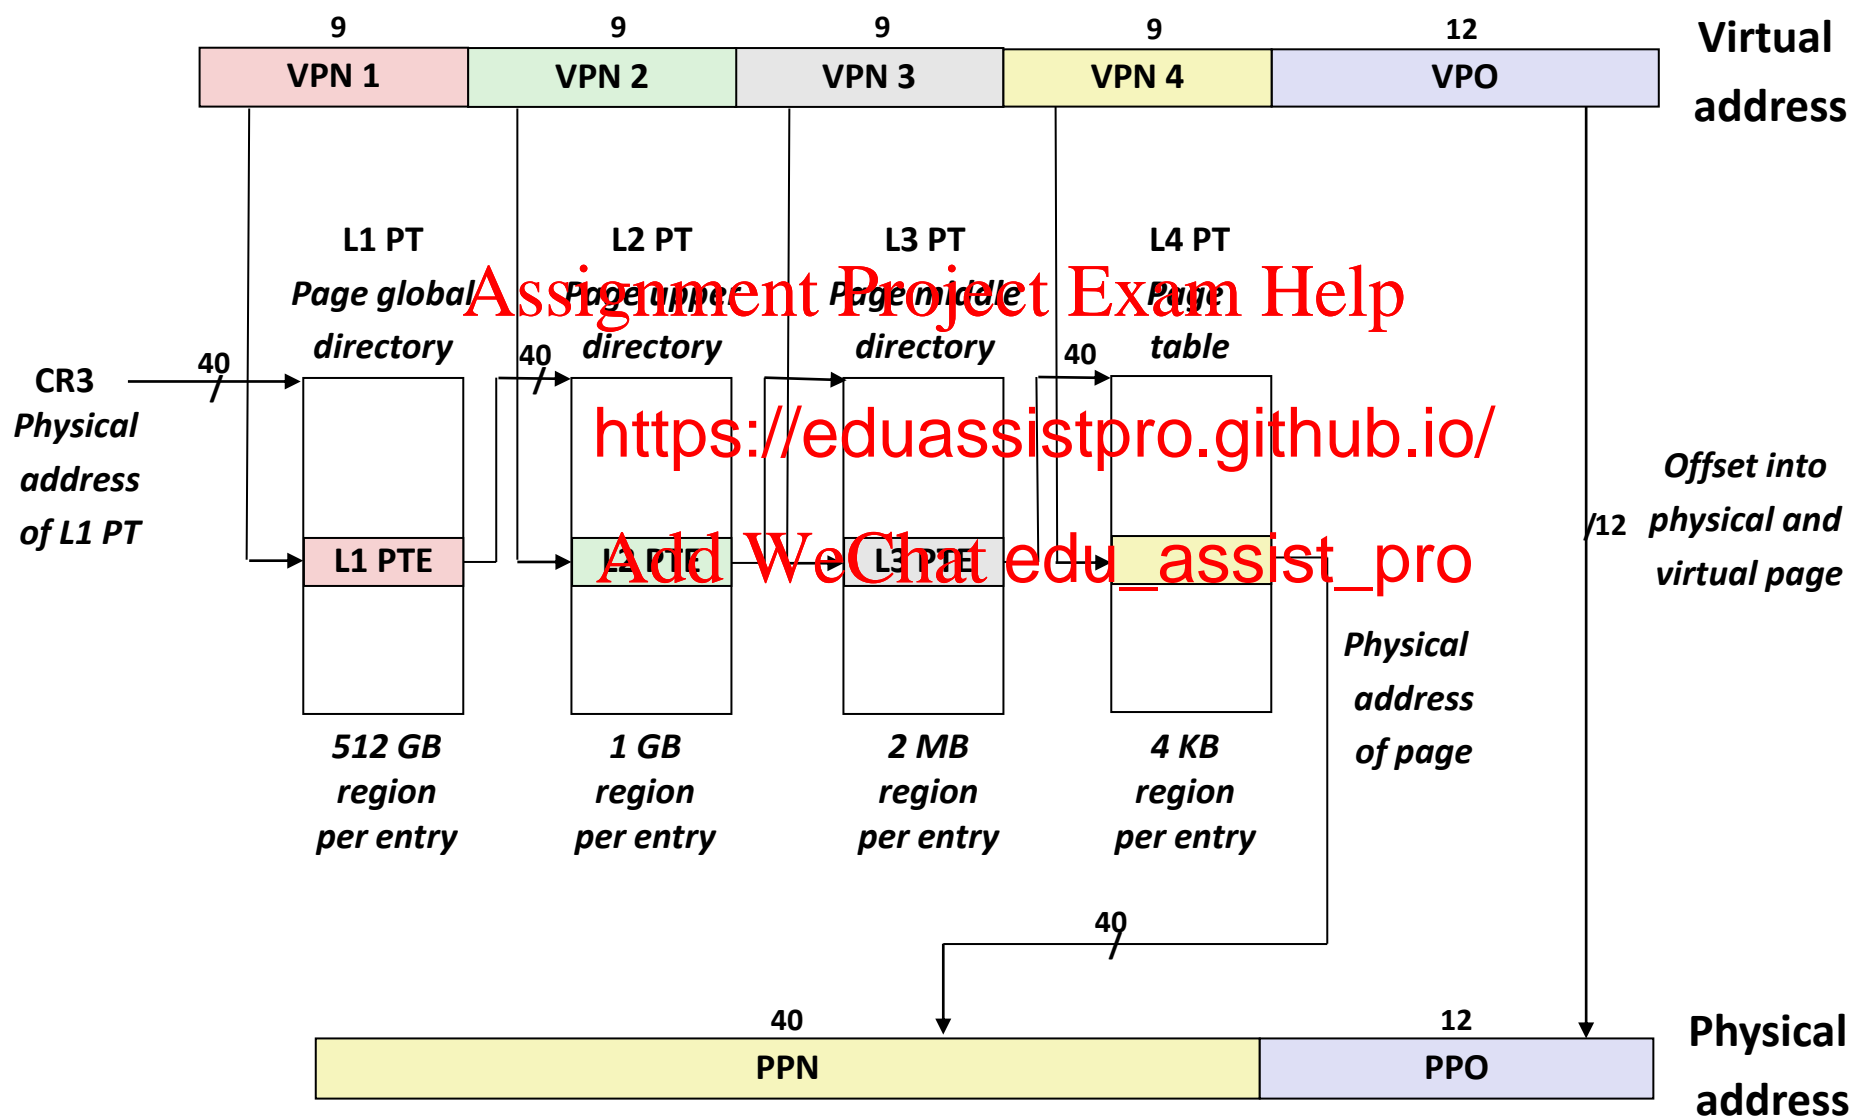
**D:** Dirty bit (set by MMU on writes, cleared by software)

**G:** Global page (don't evict from TLB on task switch)

**Page physical base address:** 40 most significant bits of physical page address  
(forces pages to be 4KB aligned)

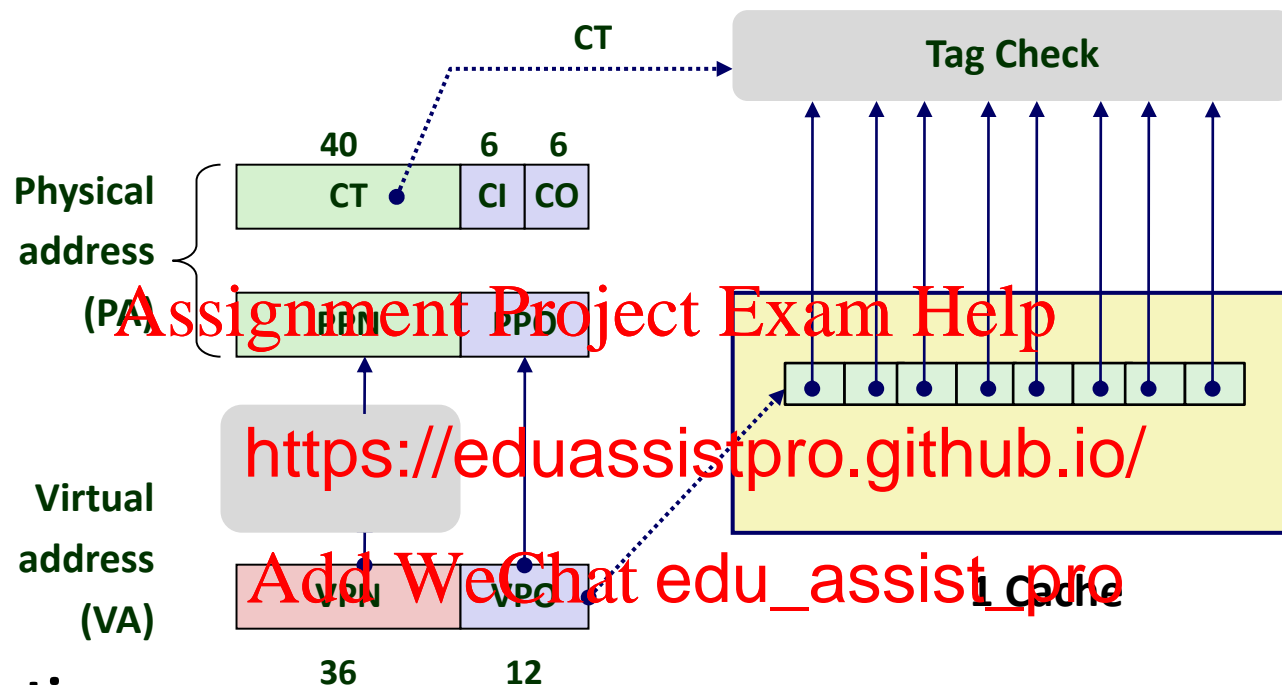
**XD:** Disable or enable instruction fetches from this page.

# Core i7 Page Table Translation





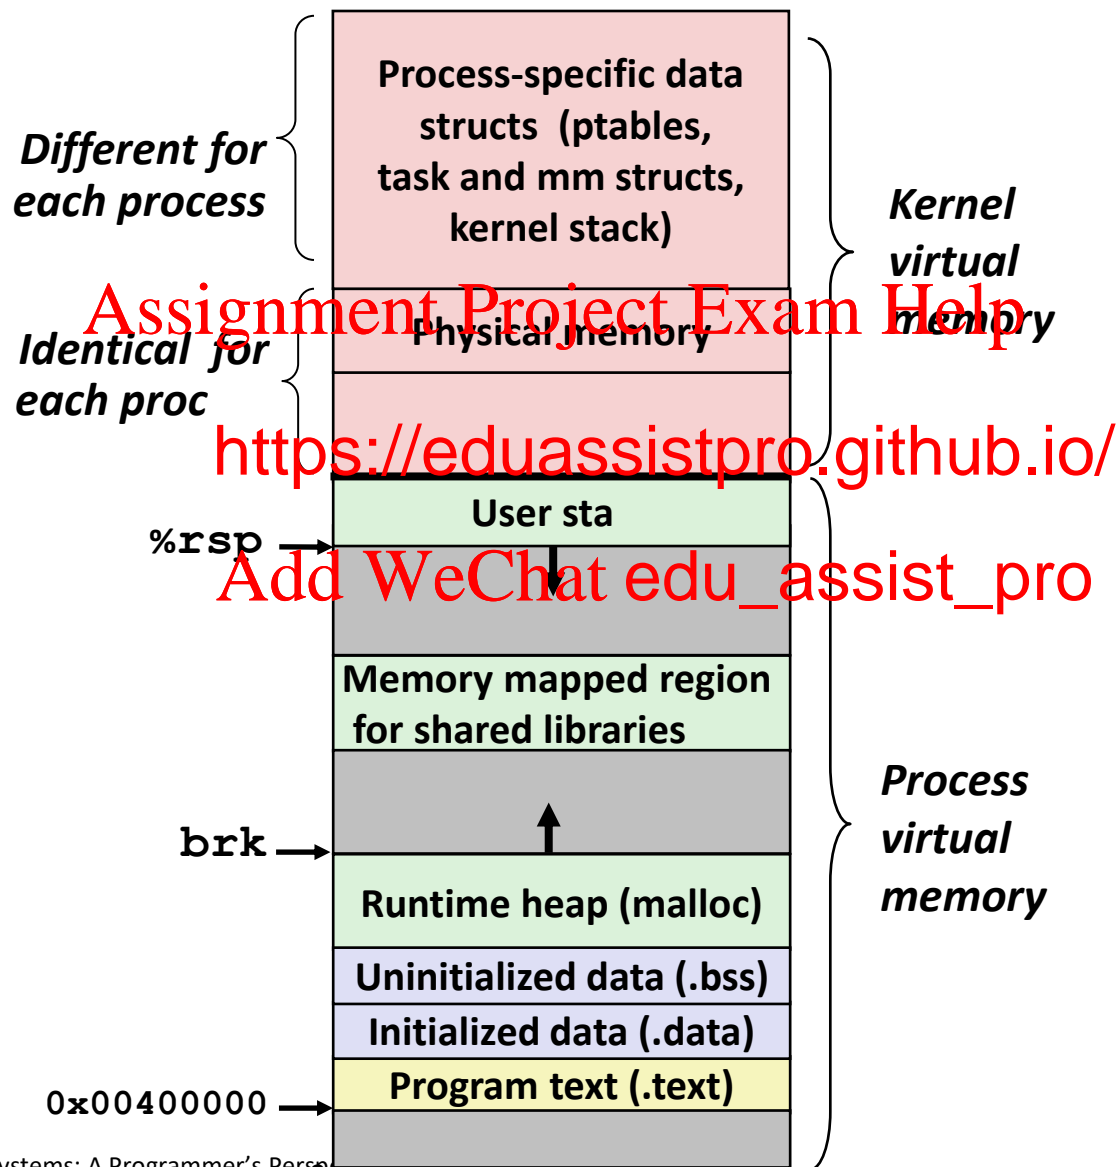
# Cute Trick for Speeding Up L1 Access



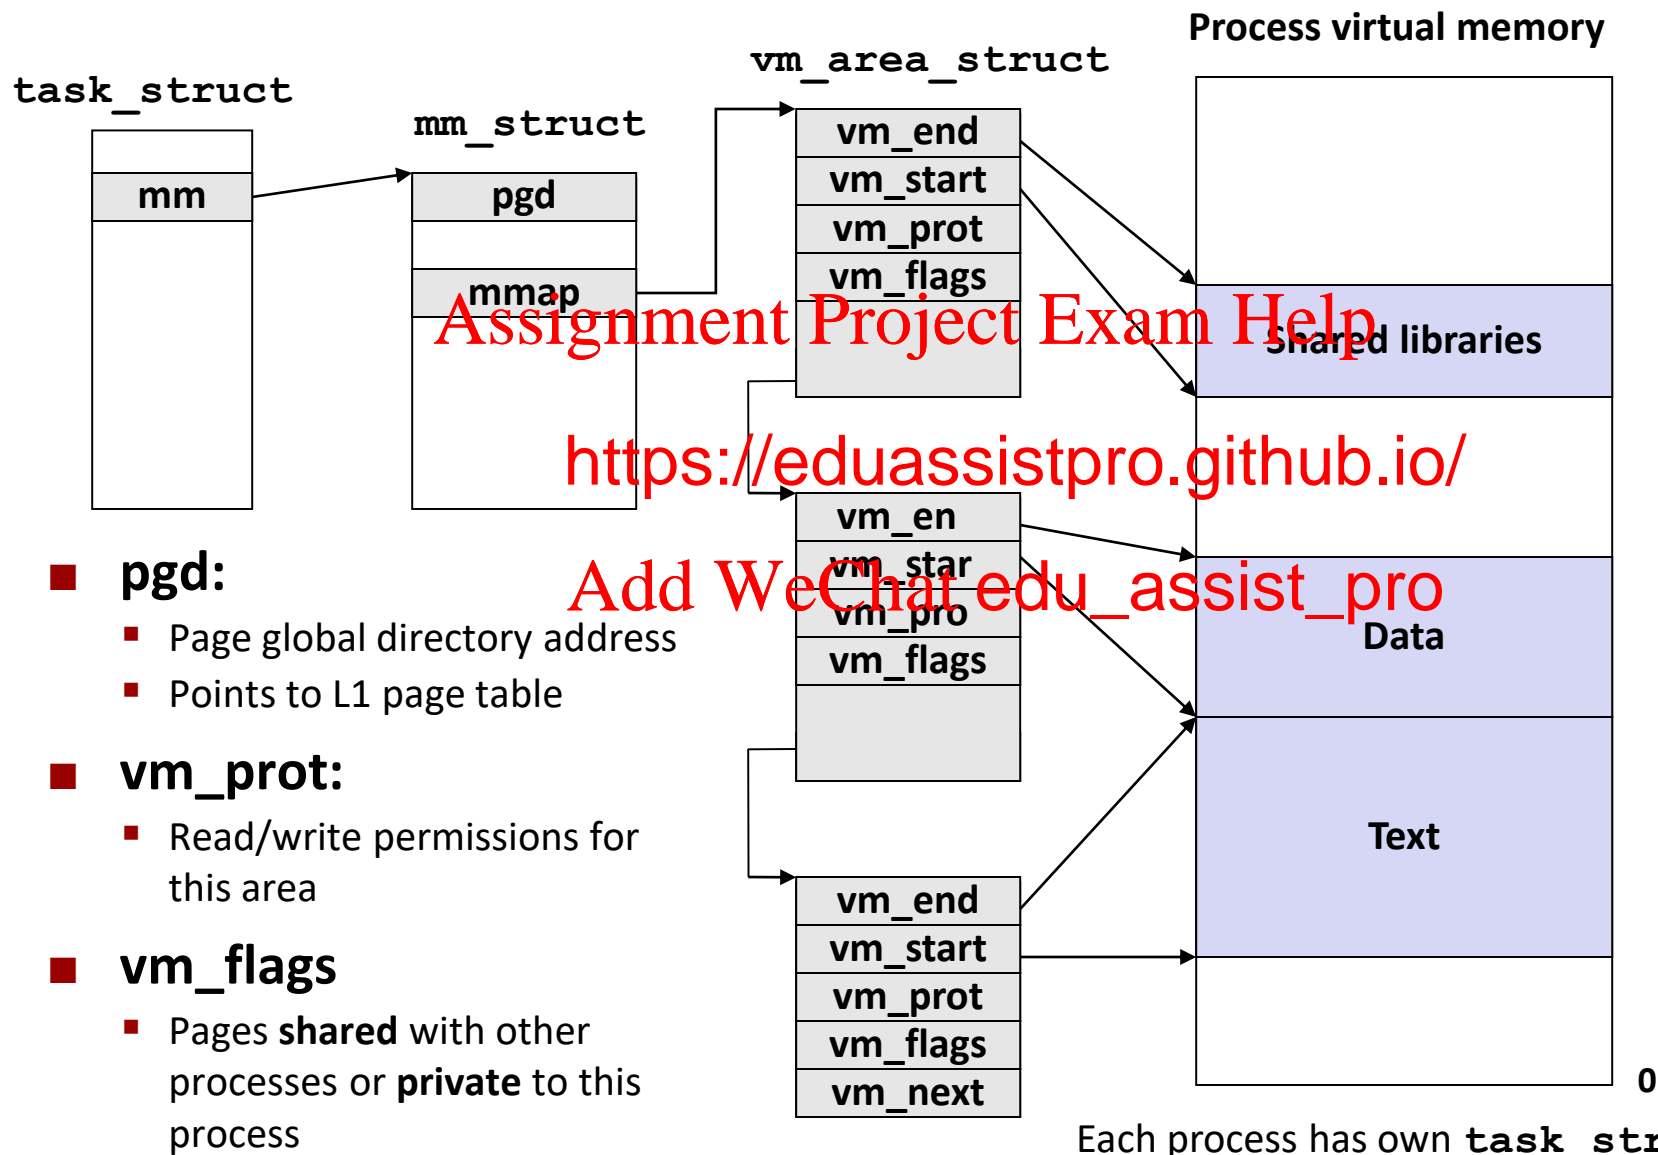
## ■ Observation

- Bits that determine CI identical in virtual and physical address
- Can index into cache while address translation taking place
- Generally we hit in TLB, so PPN bits (CT bits) available quickly
- ***“Virtually indexed, physically tagged”***
- Cache carefully sized to make this possible

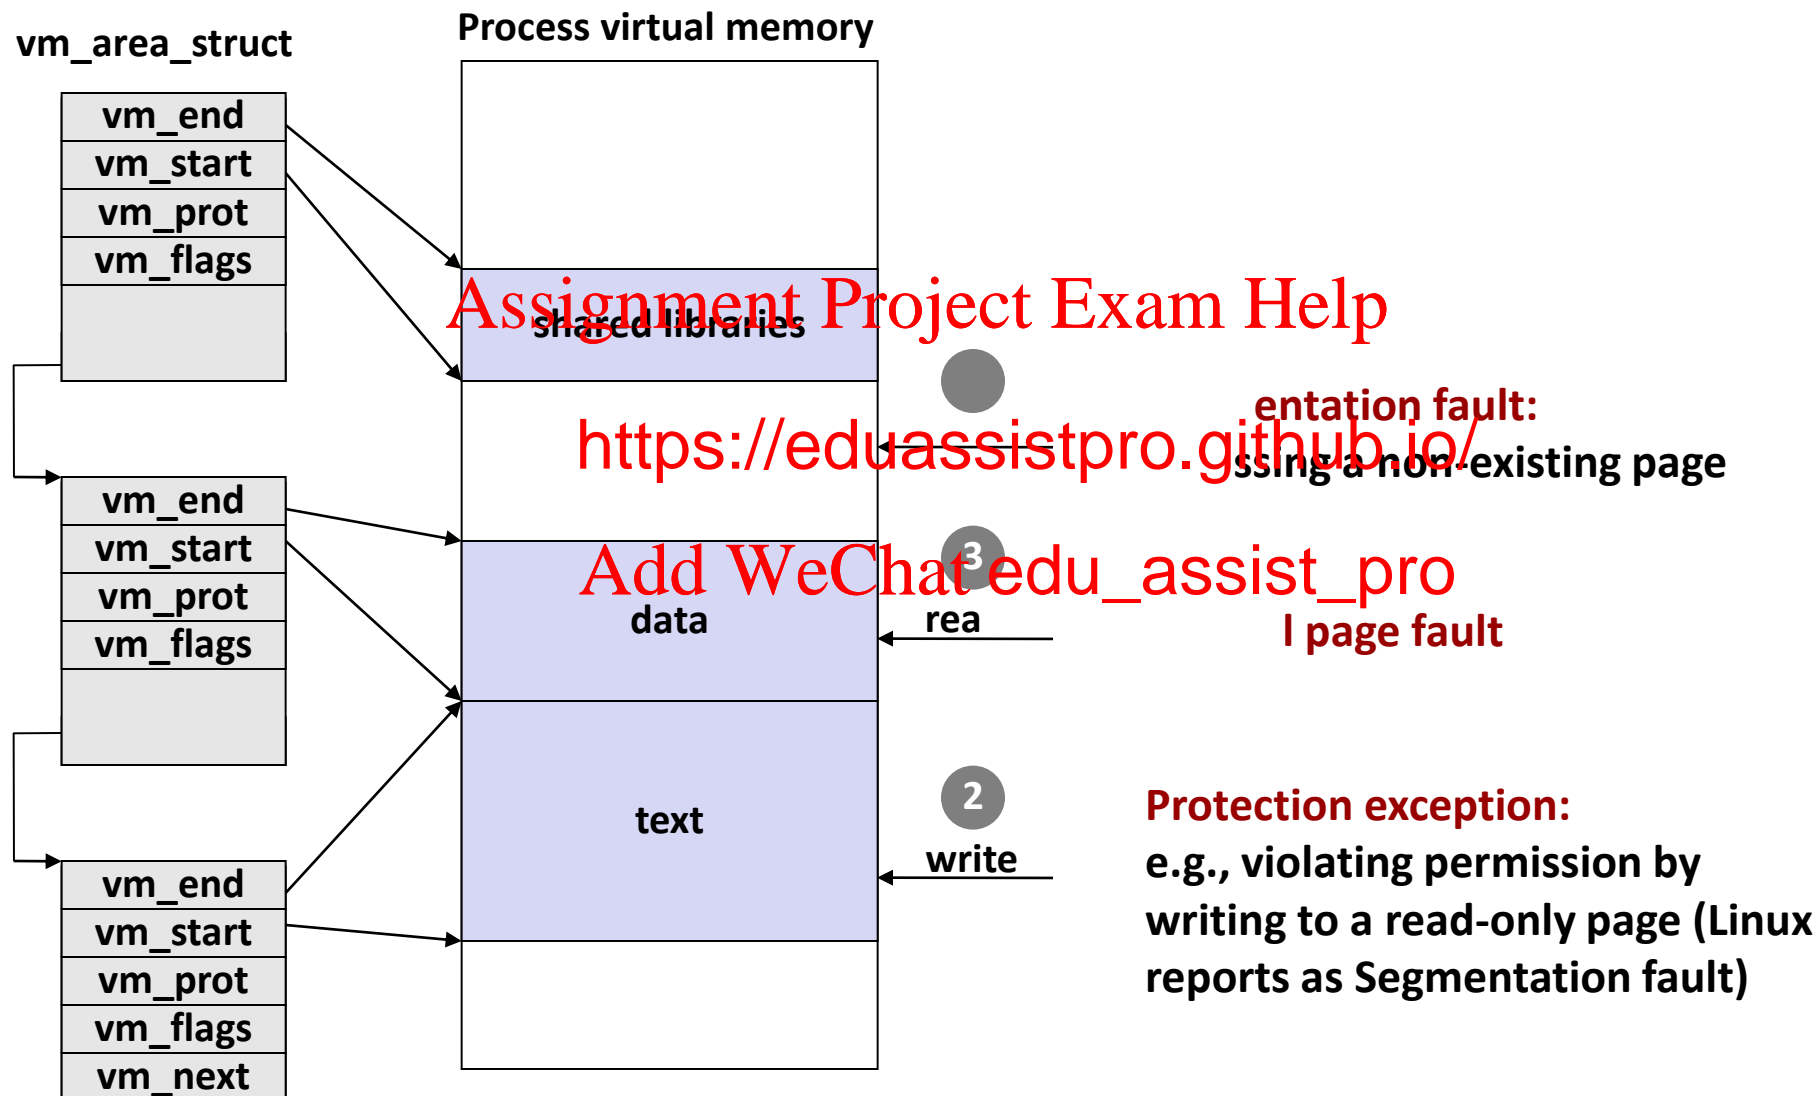
# Virtual Address Space of a Linux Process



# Linux Organizes VM as Collection of “Areas”



# Linux Page Fault Handling



# Today

- Simple memory system example
- Case study: Core i7/Linux memory system
- **Memory mapping**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

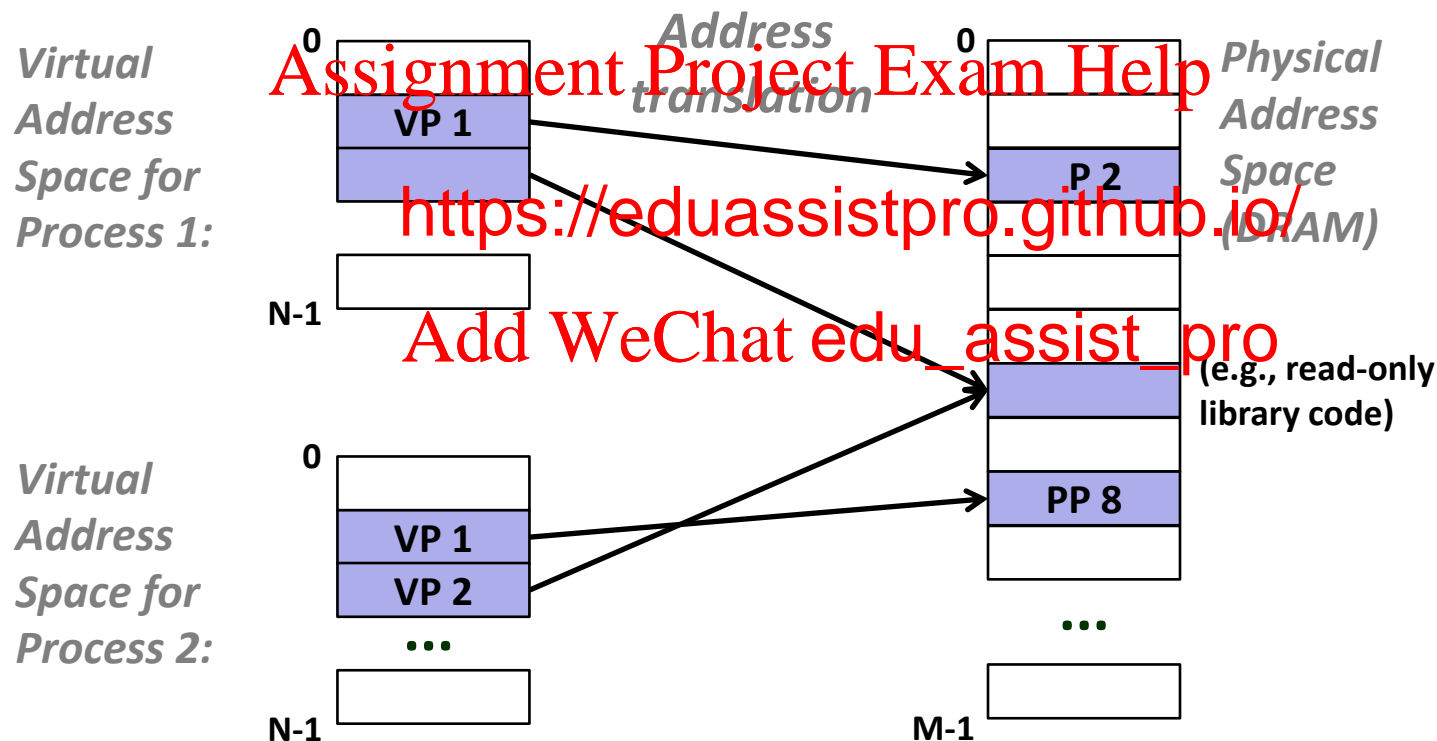
Add WeChat edu\_assist\_pro

# Memory Mapping

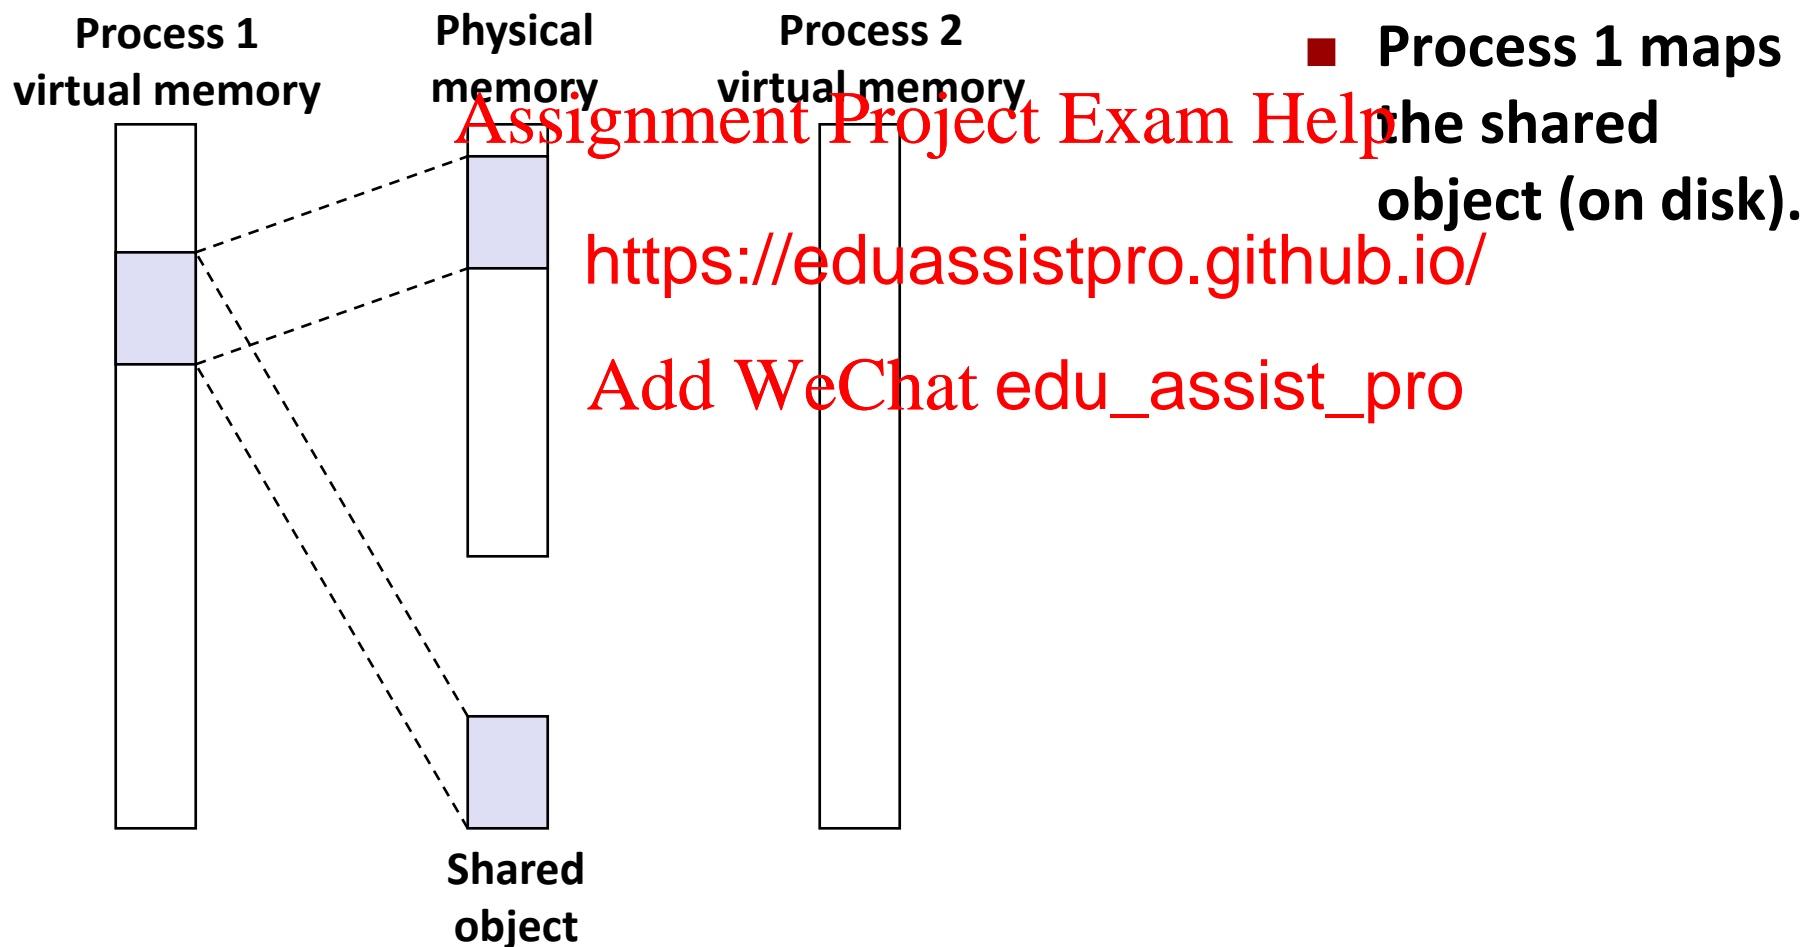
- VM areas initialized by associating them with disk objects.
  - Called *memory mapping*
- Area can be **backed by file** (i.e., get its initial values from) :
  - *Regular file* on **file**
    - Initial page **file**
  - *Anonymous file* (e.g., nothing)
    - First fault will allocate a physic's (*demand-zero page*)
    - Once the page is written to (*dirtied*), it is like any other page
- Dirty pages are copied back and forth between memory and a special *swap file*.

# Review: Memory Management & Protection

- Code and data can be isolated or shared among processes

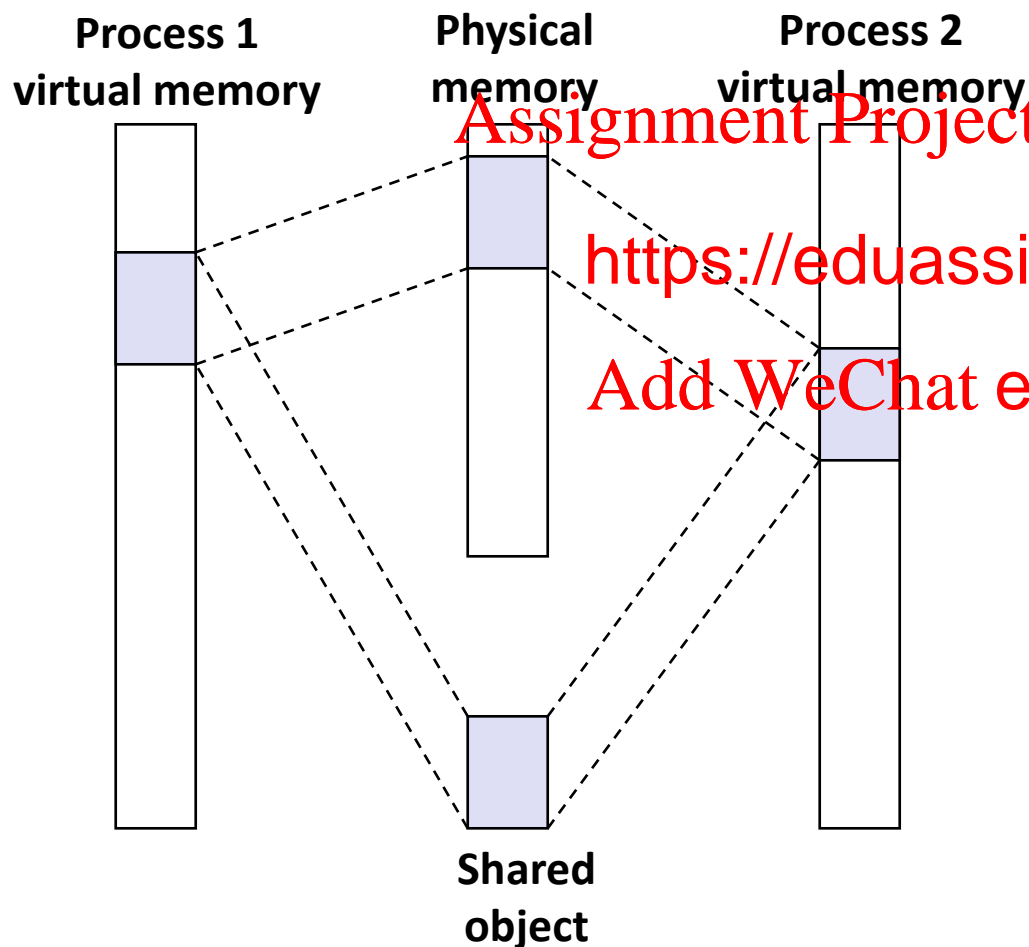


# Sharing Revisited: Shared Objects





# Sharing Revisited: Shared Objects

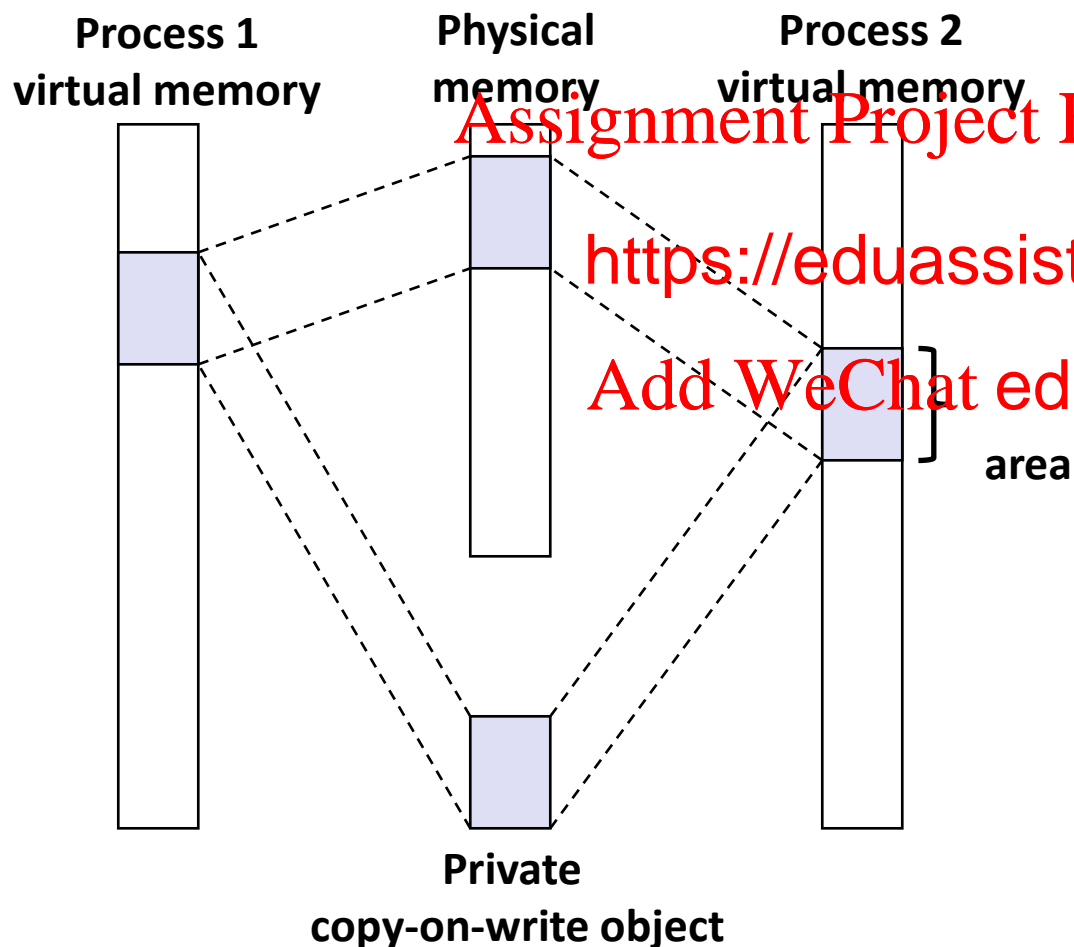


- Process 2 maps the same shared object.
- Notice how the virtual addresses can be different.
- But, difference must be multiple of page size.

Assignment Project Exam Help  
<https://eduassistpro.github.io/>  
 Add WeChat edu\_assist\_pro

# Sharing Revisited:

## Private Copy-on-write (COW) Objects



Assignment Project Exam Help

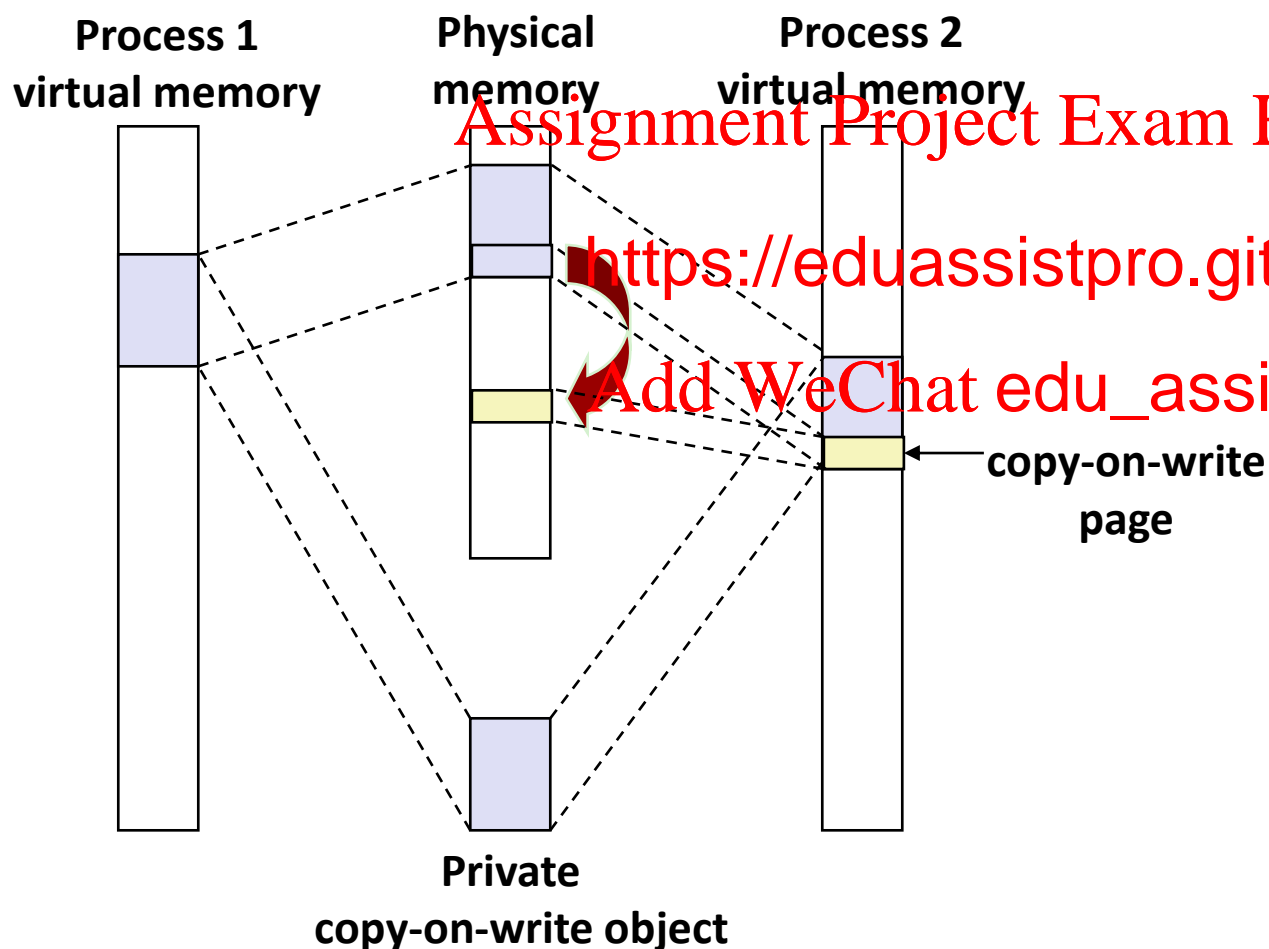
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

- Two processes mapping a *private copy-on-write (COW)* object
- Area flagged as private copy-on-write
- PTEs in private areas are flagged as read-only

# Sharing Revisited:

## Private Copy-on-write (COW) Objects



- Instruction writing to private page triggers protection fault.
- Handler creates new R/W page.
- Instruction restarts upon handler return.
- Copying deferred as long as possible!

# Finding Shareable Pages

## ■ Kernel Same-Page Merging

- OS scans through all of physical memory, looking for duplicate pages
- When found, merge into single copy, marked as copy-on-write
- Implemented in Linux kernel in 2009
- Limited to page
- Especially useful by virtual machines

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# User-Level Memory Mapping

```
void *mmap(void *start, int len,
           int prot, int flags, int fd, int offset)
```

- Map `len` bytes starting at offset `offset` of the file specified by file description `fd`, preferably at address `start`
  - `start`: may be `0`
  - `prot`: `PROT_R`, `PROT_W`, `PROT_RW`, ...
  - `flags`: `MAP_ANON`, `MAP_PRIVATE`, `MAP_SHARED`, ...
- Return a pointer to start of mapped area (may not be `start`)

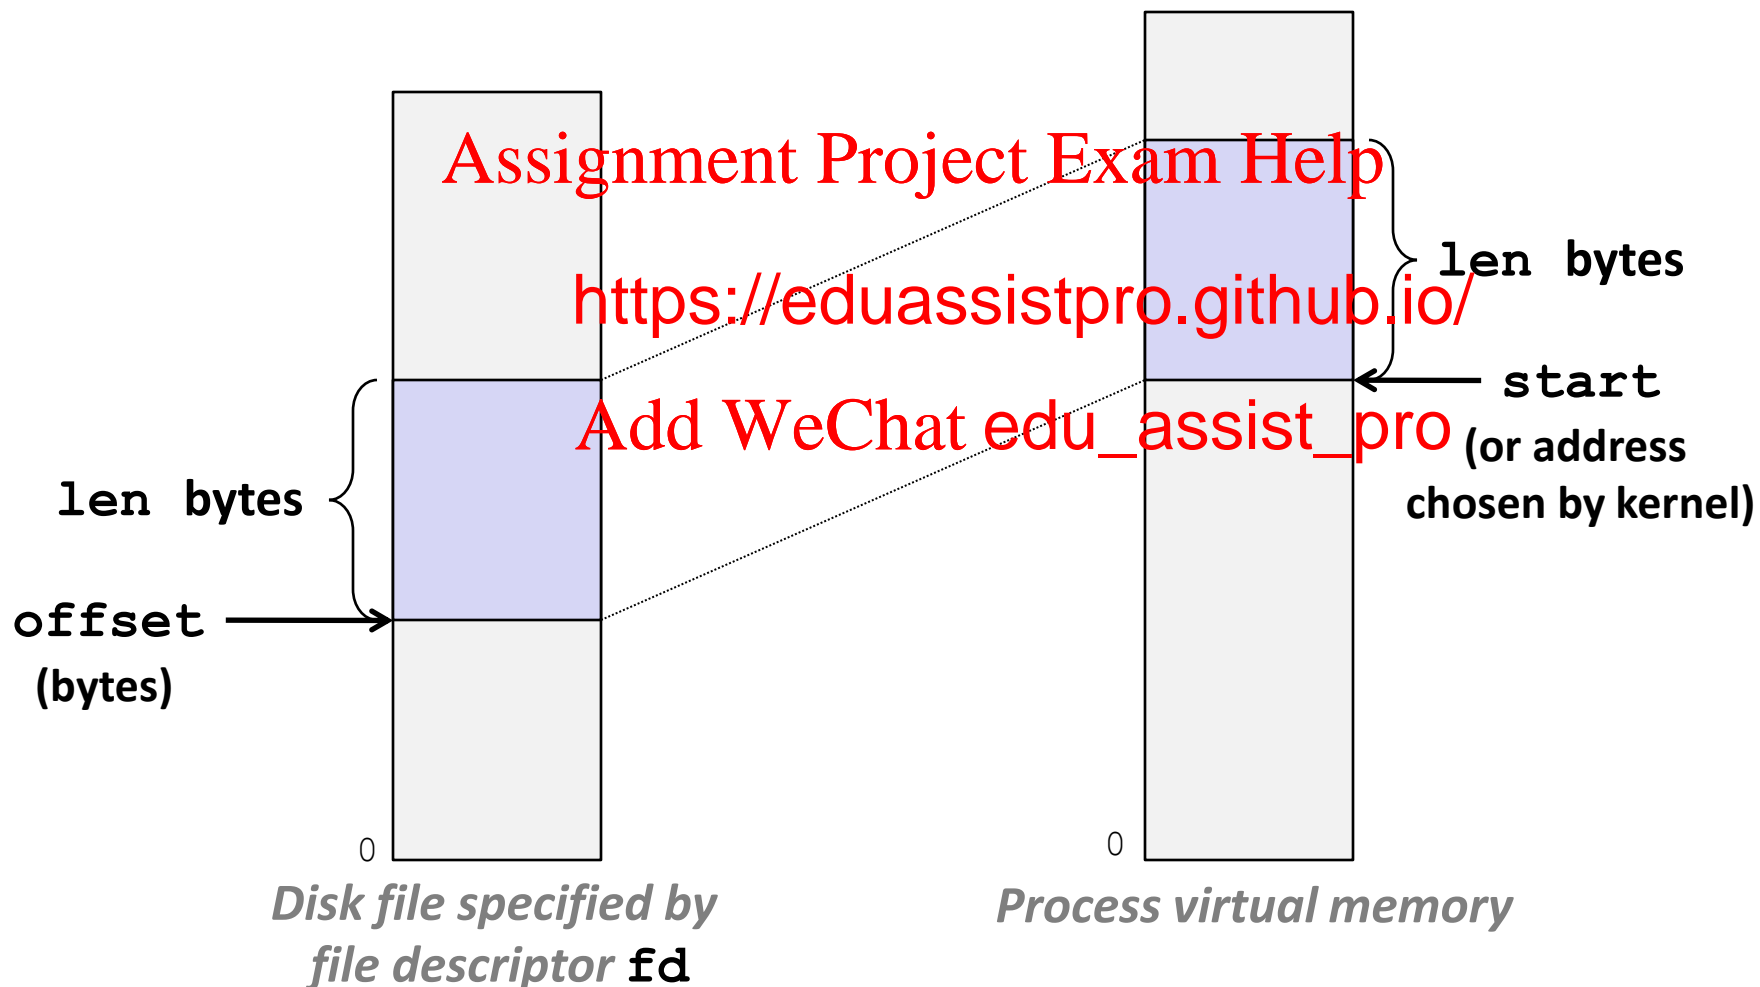
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu\_assist\_pro

# User-Level Memory Mapping

```
void *mmap(void *start, int len,
           int prot, int flags, int fd, int offset)
```



# Uses of mmap

## ■ Reading big files

- Uses paging mechanism to bring files into memory

## ■ Shared data structures

- When call with `MAP_SHARED` flag
  - Multiple processes can access same region of memory
  - Risky!

## ■ File-based data structures

- E.g., database
- Give `prot` argument `PROT_READ` | `PROT_WRITE`
- When unmap region, file will be updated via write-back
- Can implement load from file / update / write back to file

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Example: Using `mmap` to Support Attack Lab

## ■ Problem

- Want students to be able to perform code injection attacks
- Shark machine stacks are not executable

## ■ Solution

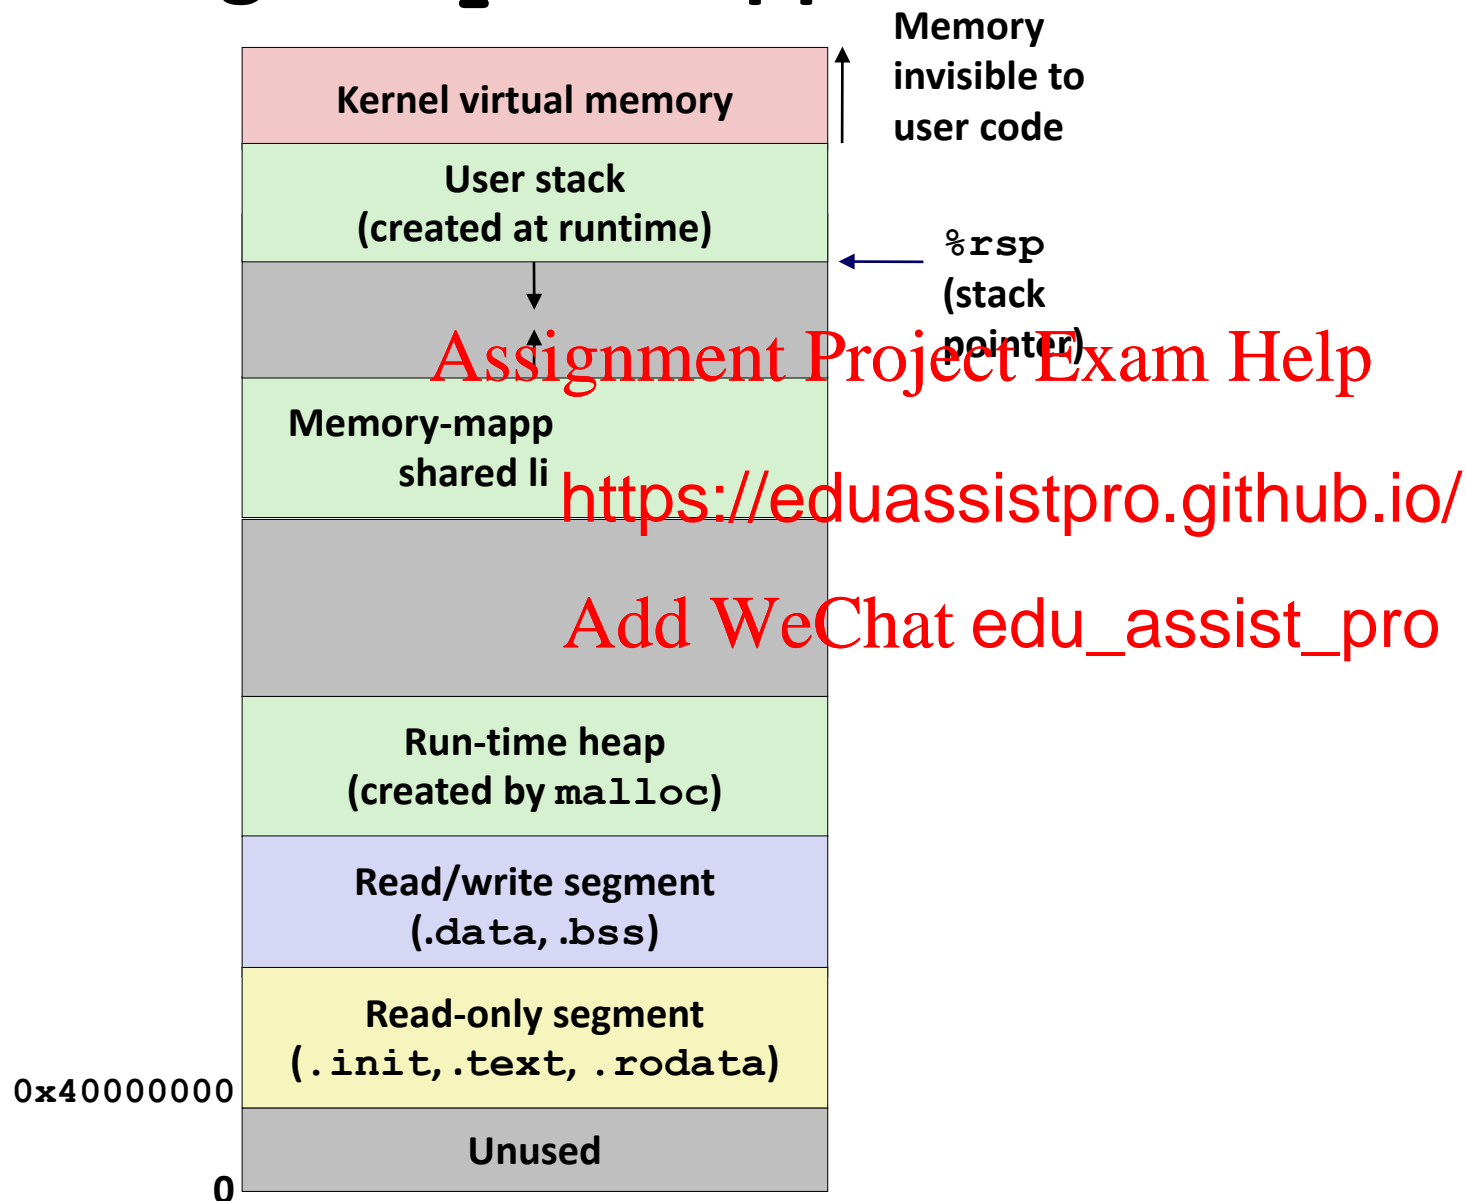
- Suggested by Sam King (now at UC Davis)
- Use `mmap` to <https://eduassistpro.github.io/> marked executable
- Divert stack to new region
- Execute student attack code
- Restore back to original stack
- Remove mapped region

Assignment Project Exam Help

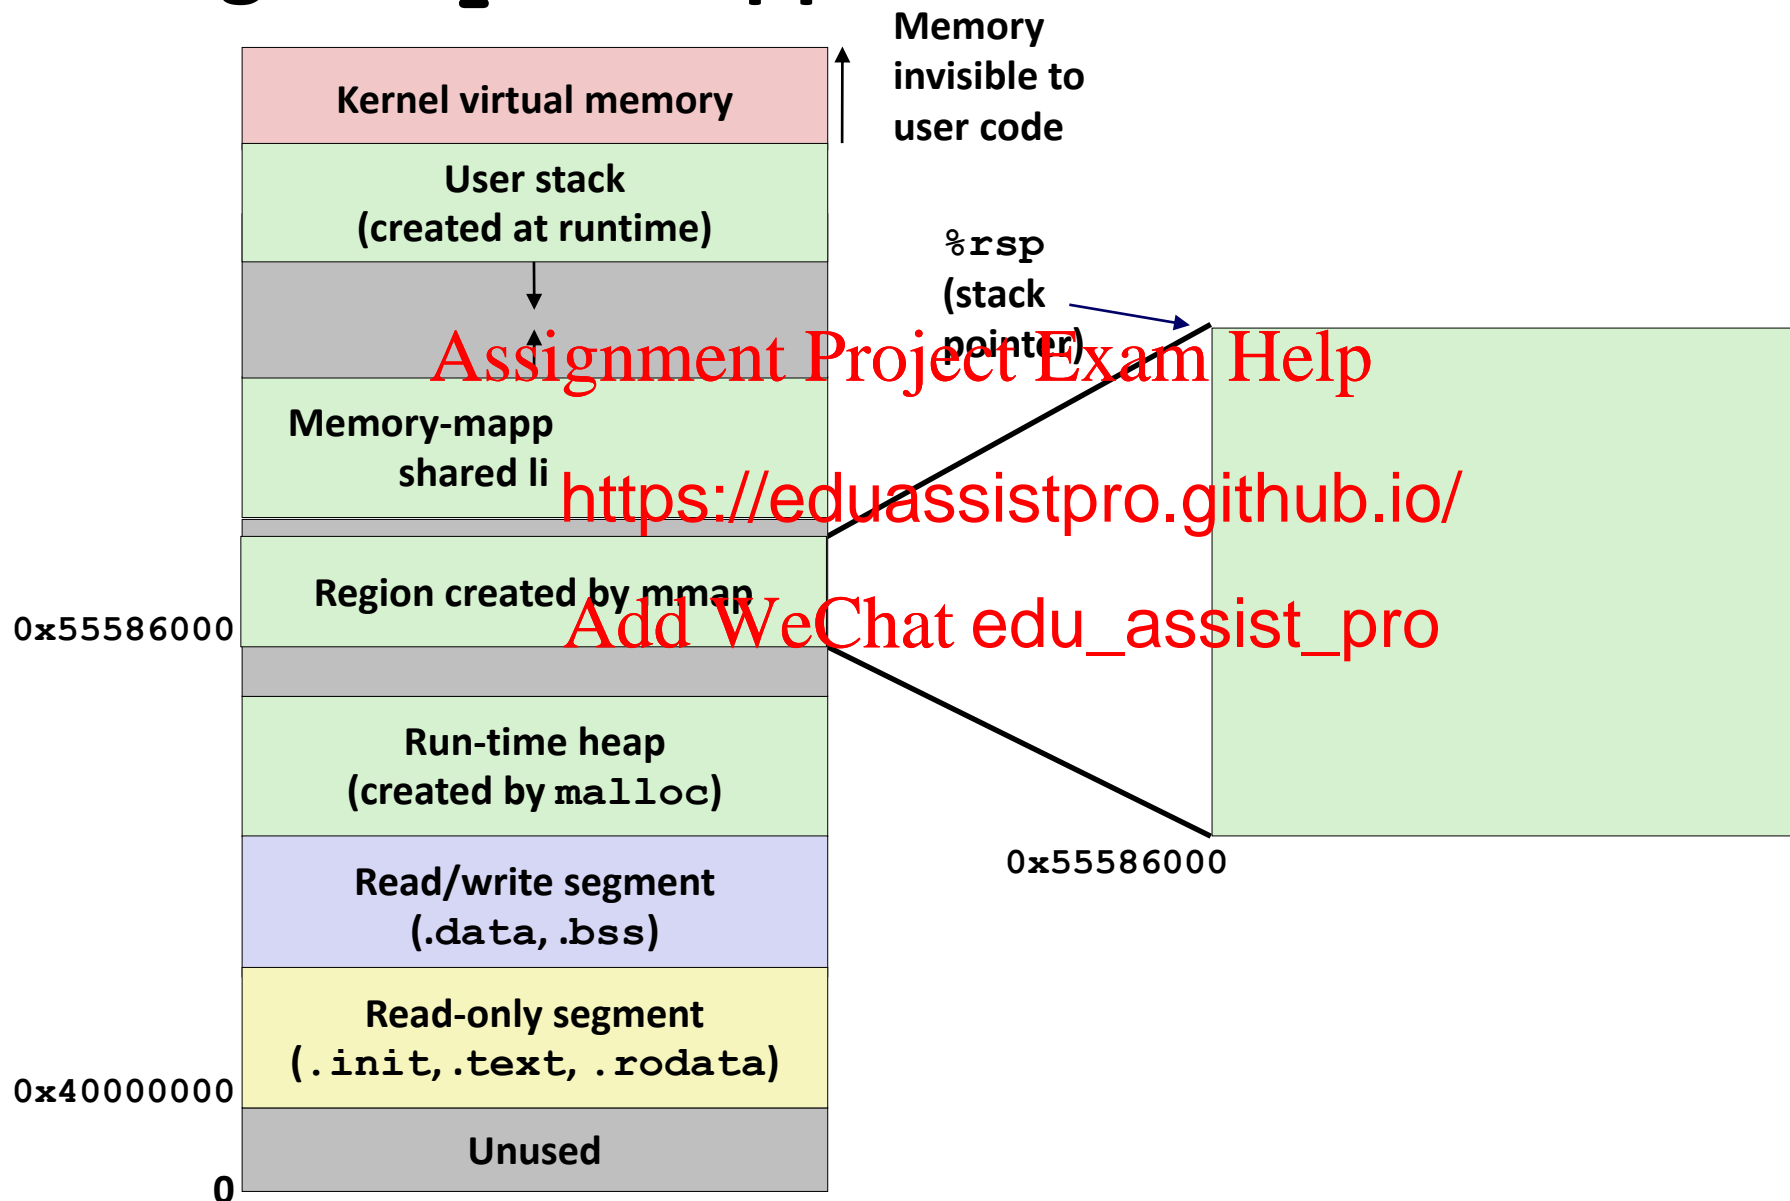
Add WeChat edu\_assist\_pro



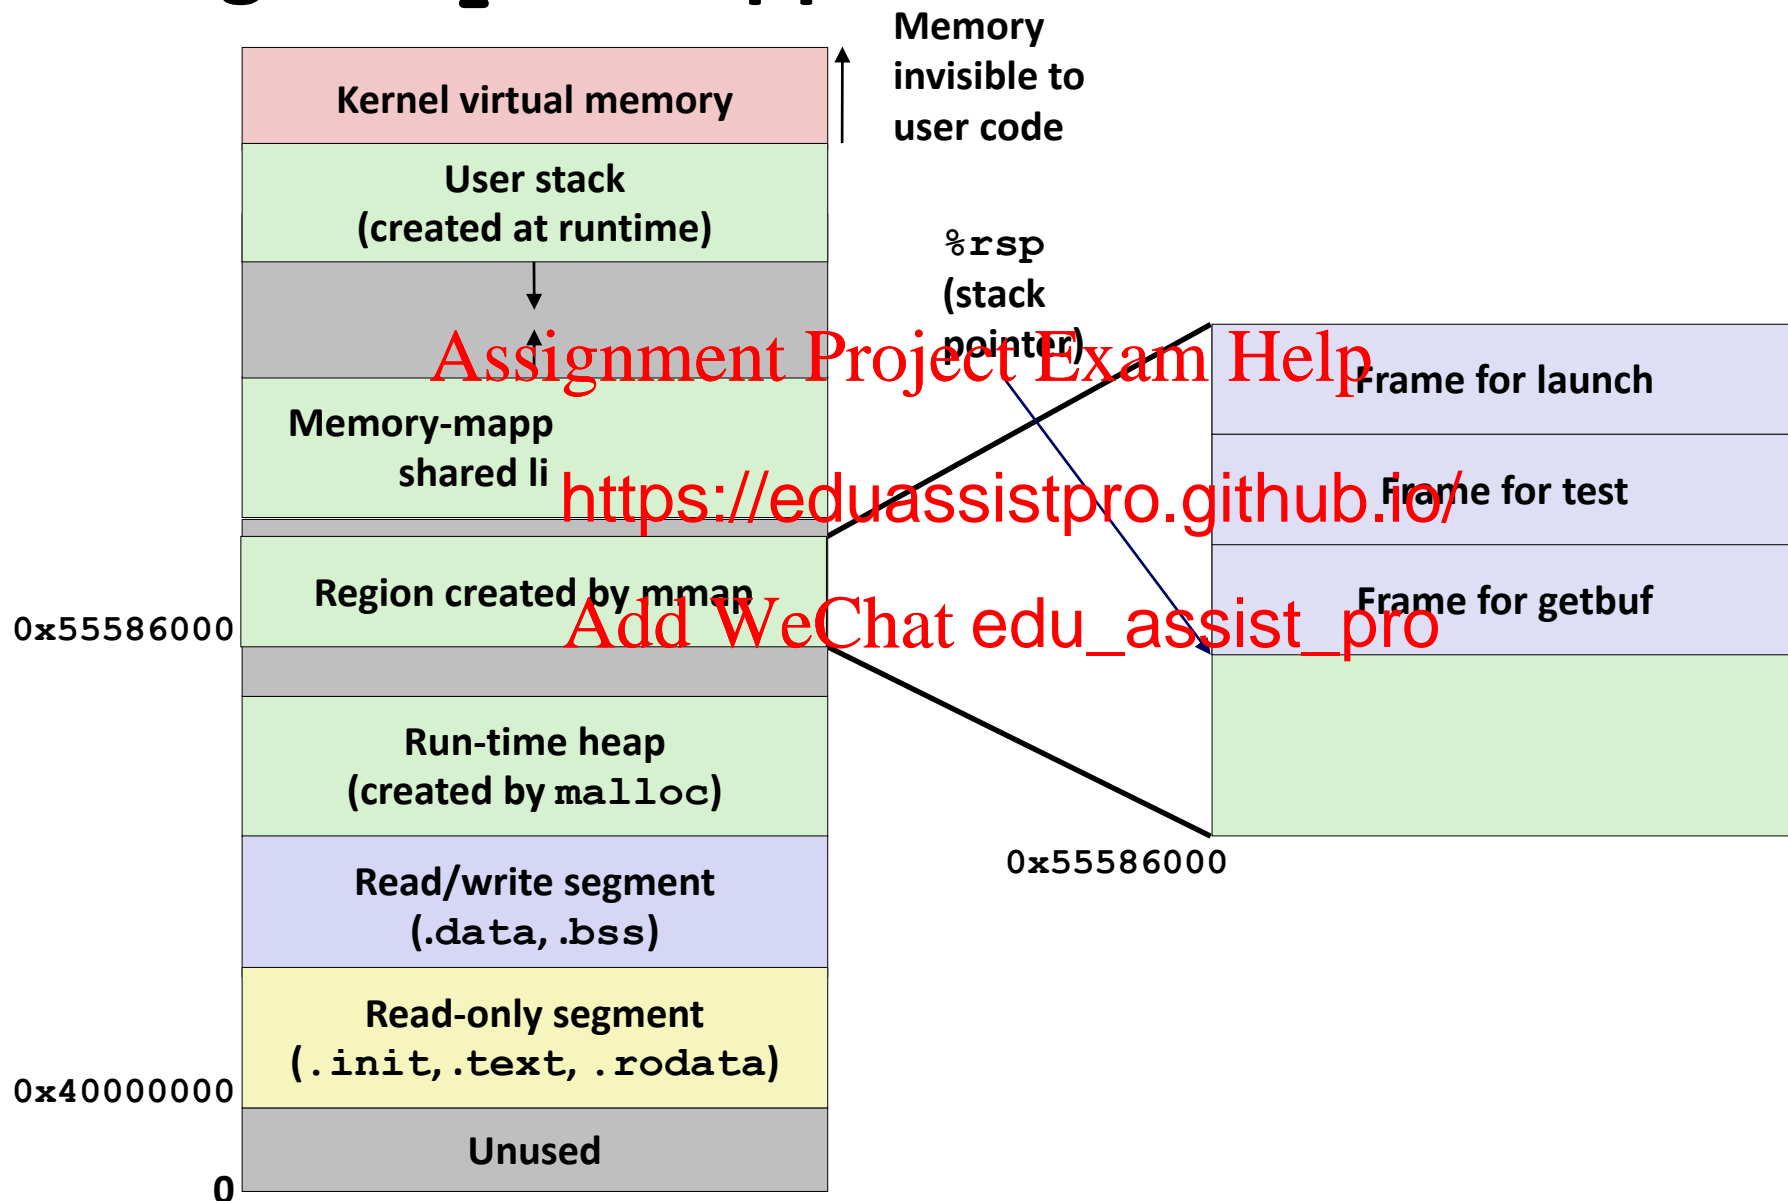
# Using mmap to Support Attack Lab



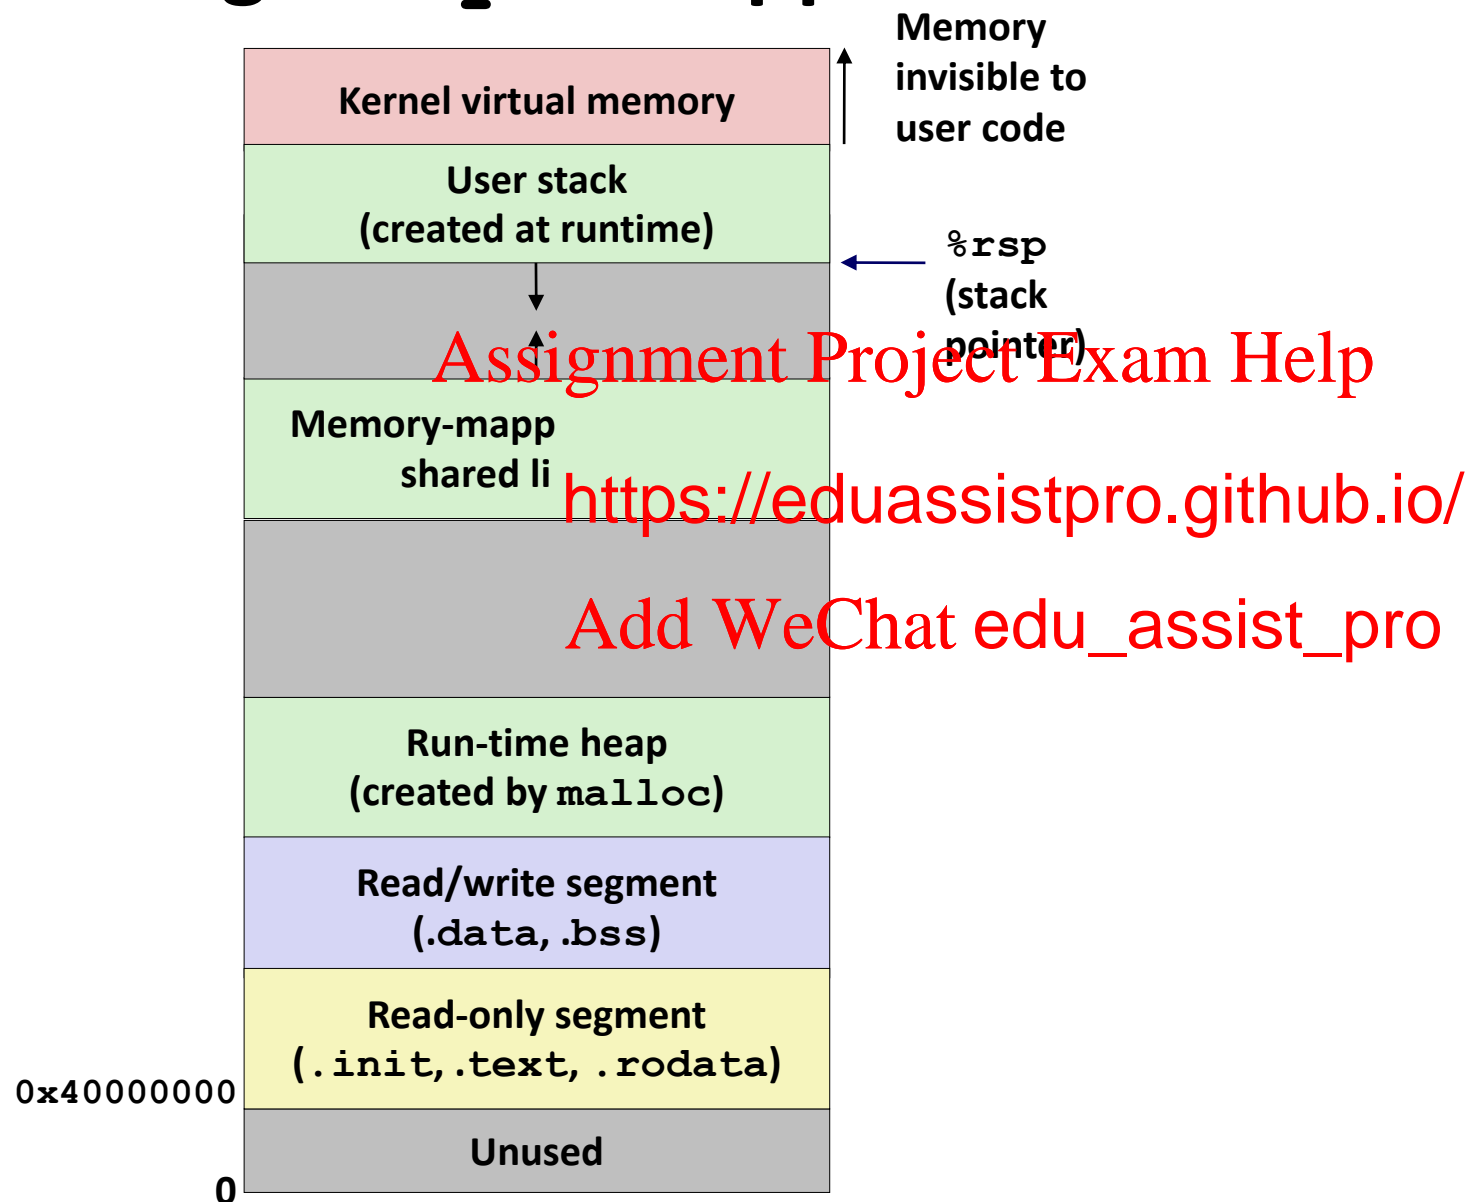
# Using mmap to Support Attack Lab



# Using mmap to Support Attack Lab



# Using mmap to Support Attack Lab



# Summary

## ■ VM requires hardware support

- Exception handling mechanism
- TLB
- Various control registers

## ■ VM requires OS

- Managing page
- Implementing page replacement
- Managing file system

## ■ VM enables many capabilities

- Loading programs from memory
- Providing memory protection

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Using mmap to Support Attack Lab

## Allocate new region

```
void *new_stack = mmap(START_ADDR, STACK_SIZE, PROT_EXEC|PROT_READ|PROT_WRITE,
                        MAP_PRIVATE | MAP_GROWSDOWN | MAP_ANONYMOUS | MAP_FIXED,
                        0, 0);
if (new_stack != START_ADDR) {
    munmap(new_stack, STACK_SIZE);
    exit(1);
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

## Divert stack to new region & execute attack co

```
stack_top = new_stack + STACK_SIZE - 8;
asm("movq %%rsp,%%rax ; movq %1,%%rsp ;
    movq %%rax,%0"
    : "=r" (global_save_stack) // %0
    : "r" (stack_top) // %1
    );
launch(global_offset);
```

## tack and remove region

```
0,%%rsp"
:
: "r" (global_save_stack) // %0
);
munmap(new_stack, STACK_SIZE);
```

Add WeChat edu\_assist\_pro