

# Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

14-513

18-613

# Network Programming: Part II

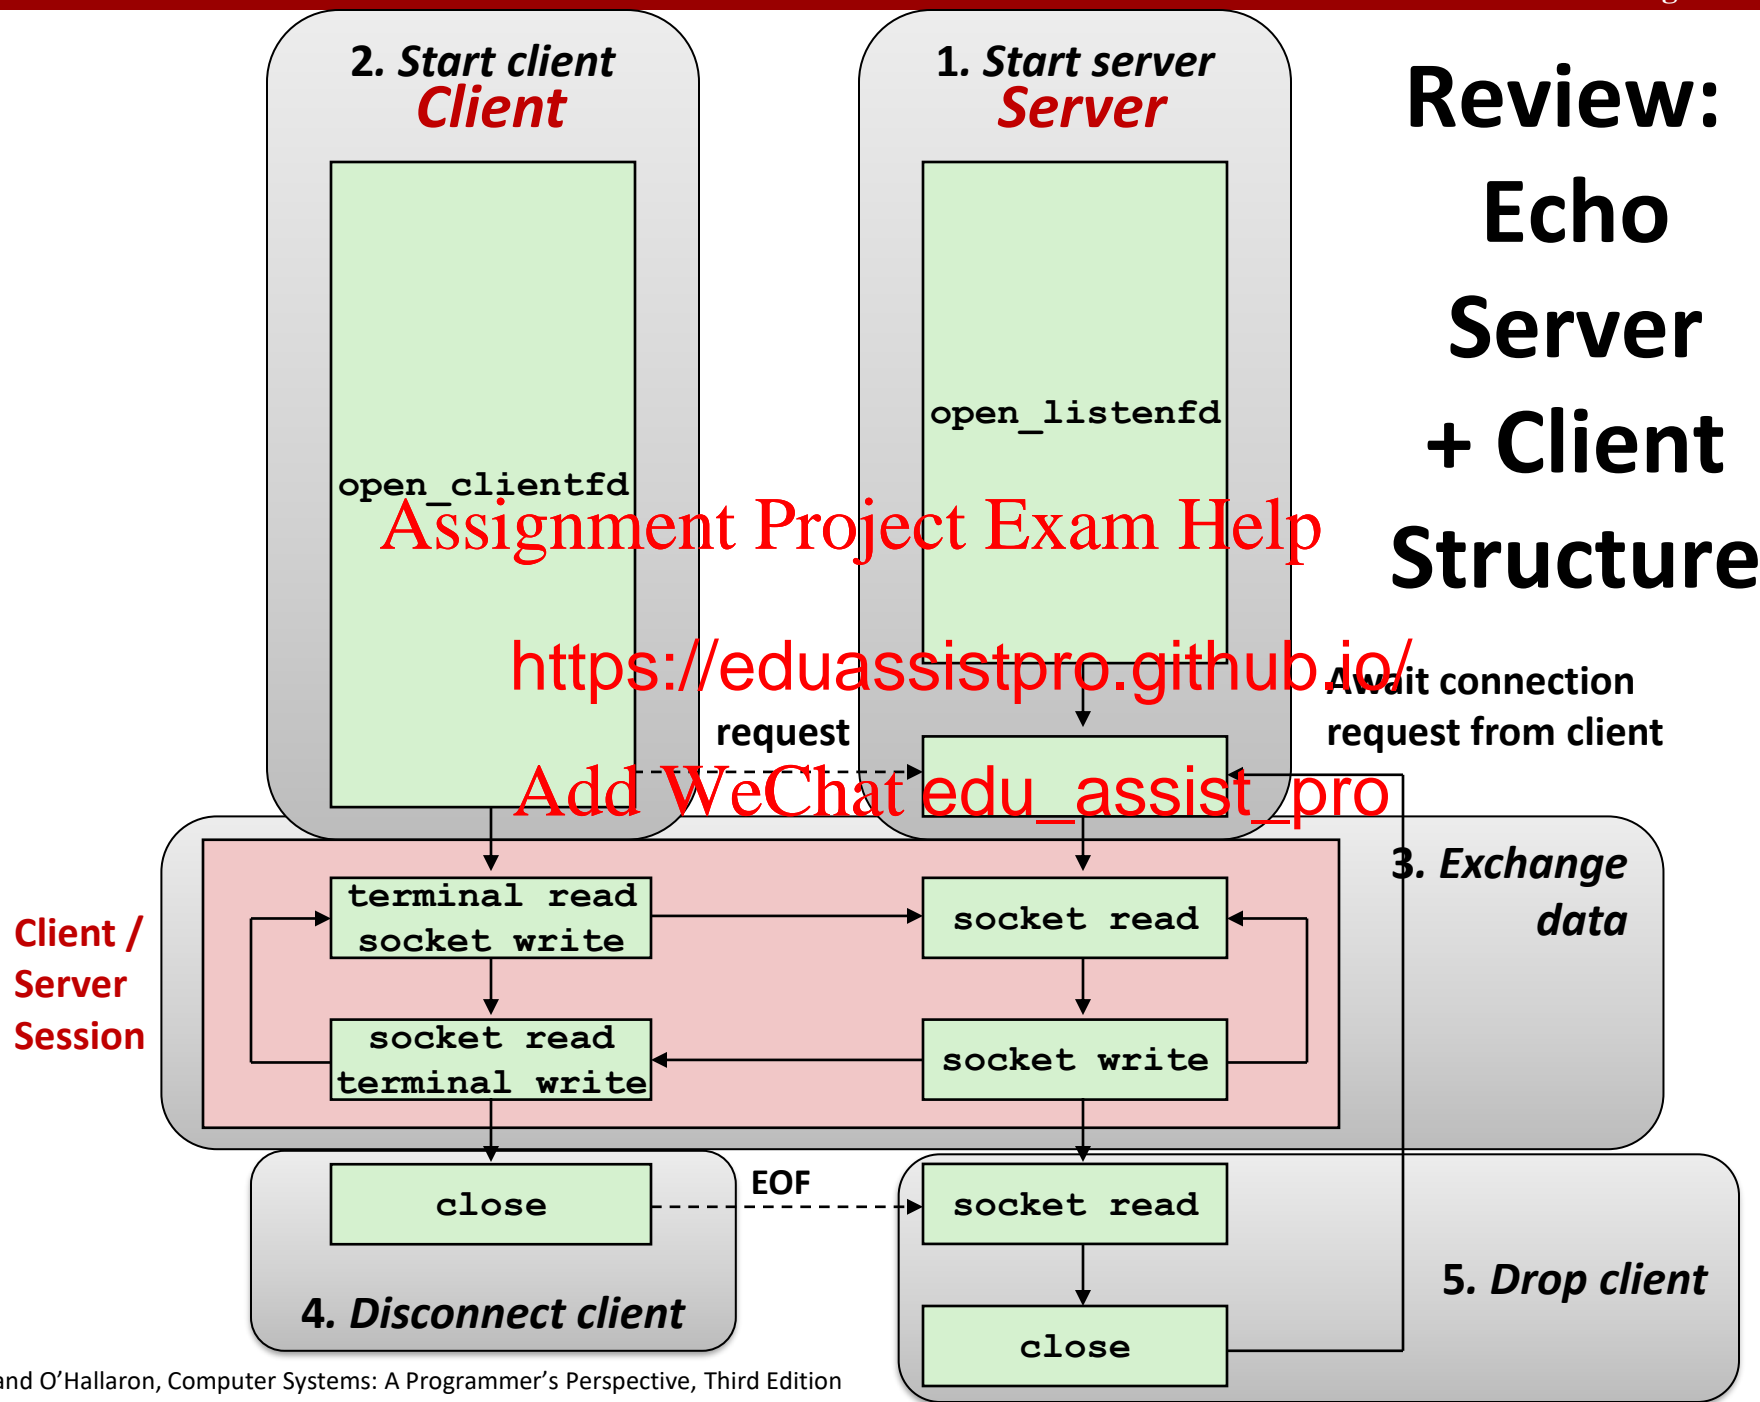
Assignment Project Exam Help

15-213/18-213/14-5

Introduction to Com  
23<sup>rd</sup> Lecture, Novem

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Today

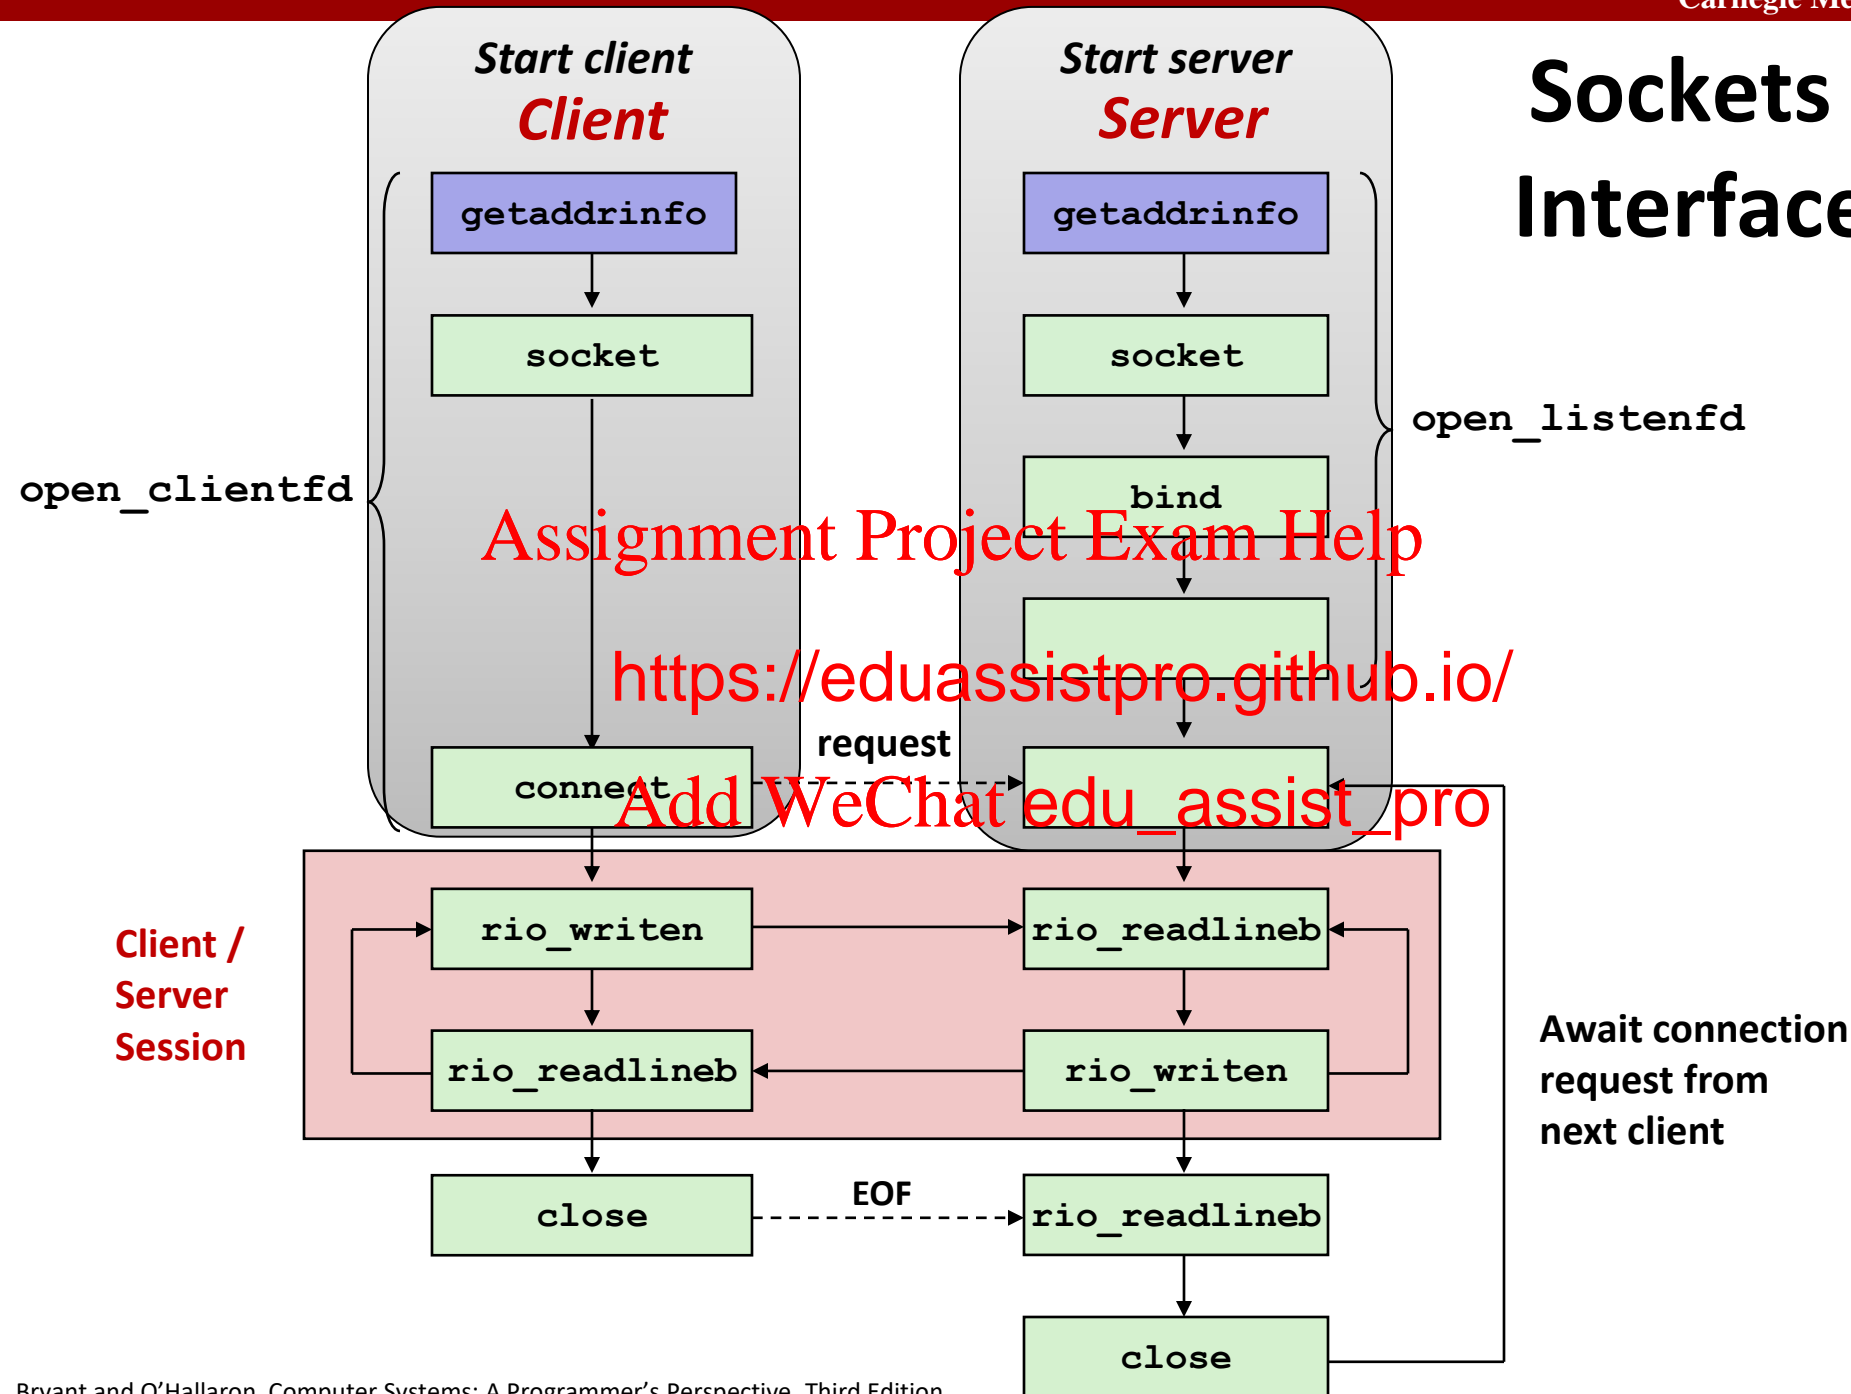
- The Sockets Interface CSAPP 11.4
- Web Servers CSAPP 11.5.1-11.5.3
- The Tiny Web Server CSAPP 11.6
- Serving Dynamic Content CSAPP 11.5.4

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Sockets Interface



# Review: Generic Socket Address

## ■ Generic socket address:


- For address arguments to **connect**, **bind**, and **accept**

```
struct sockaddr {
    uint16_t  s
    char      s
};
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



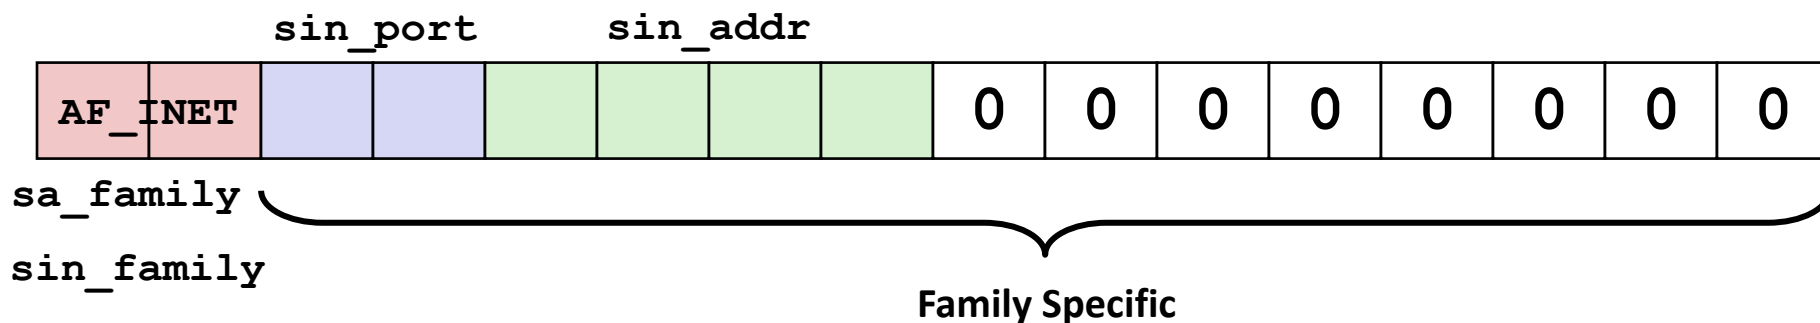
# Review: Socket Address Structures

## ■ Internet (IPv4) specific socket address:

- Must cast (`struct sockaddr_in *`) to (`struct sockaddr *`) for functions that take socket address arguments.

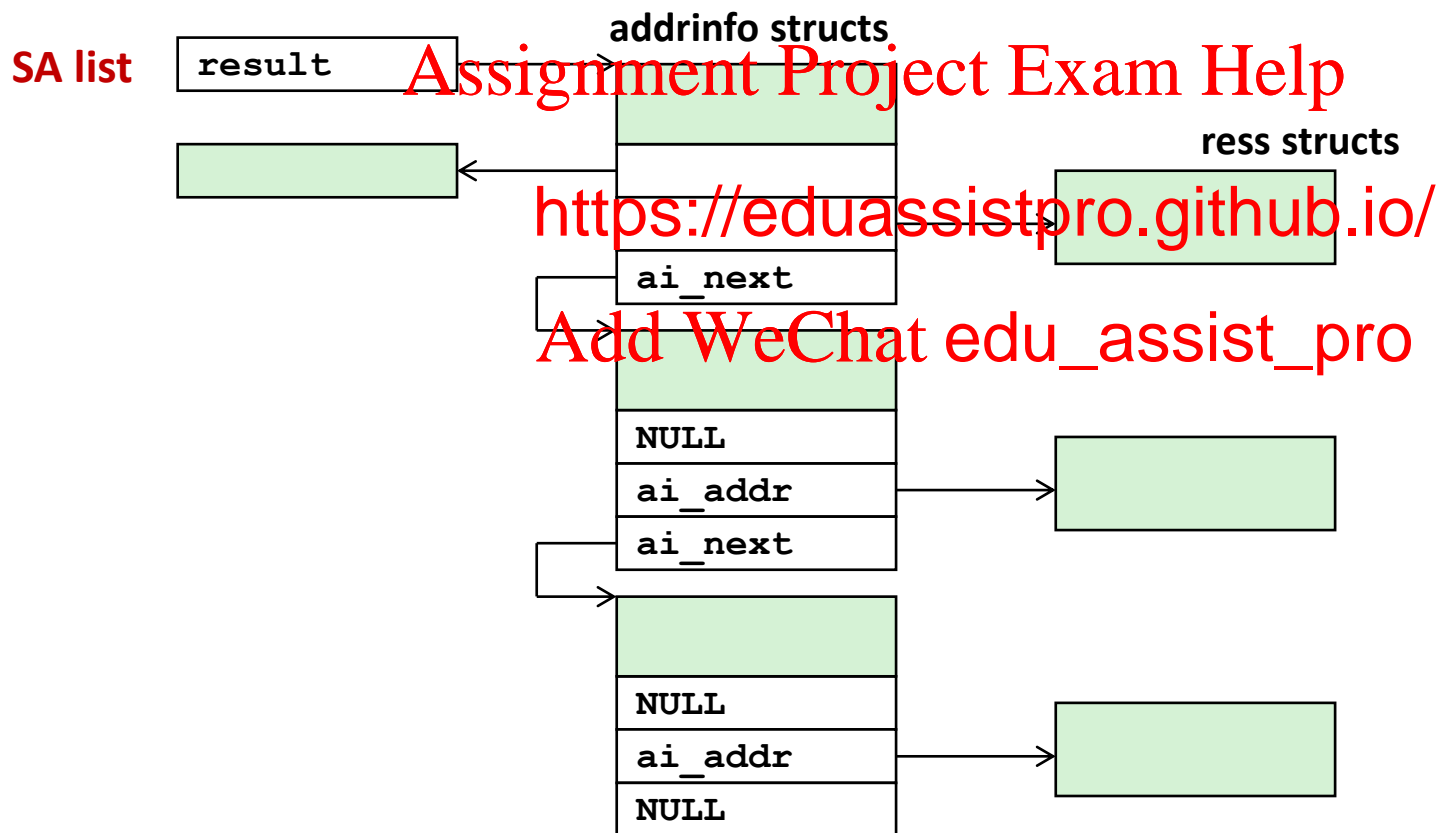
Assignment Project Exam Help

```
struct sockaddr_in {
    uint16_t      sin_port;           /* Port number, network byte order */
    uint16_t      sin_family;         /* Address family (always AF_INET) */
    struct in_addr sin_addr;          /* IP Address, network byte order */
    unsigned char sin_zero[8];        /* Pad to size of struct sockaddr */
};
```



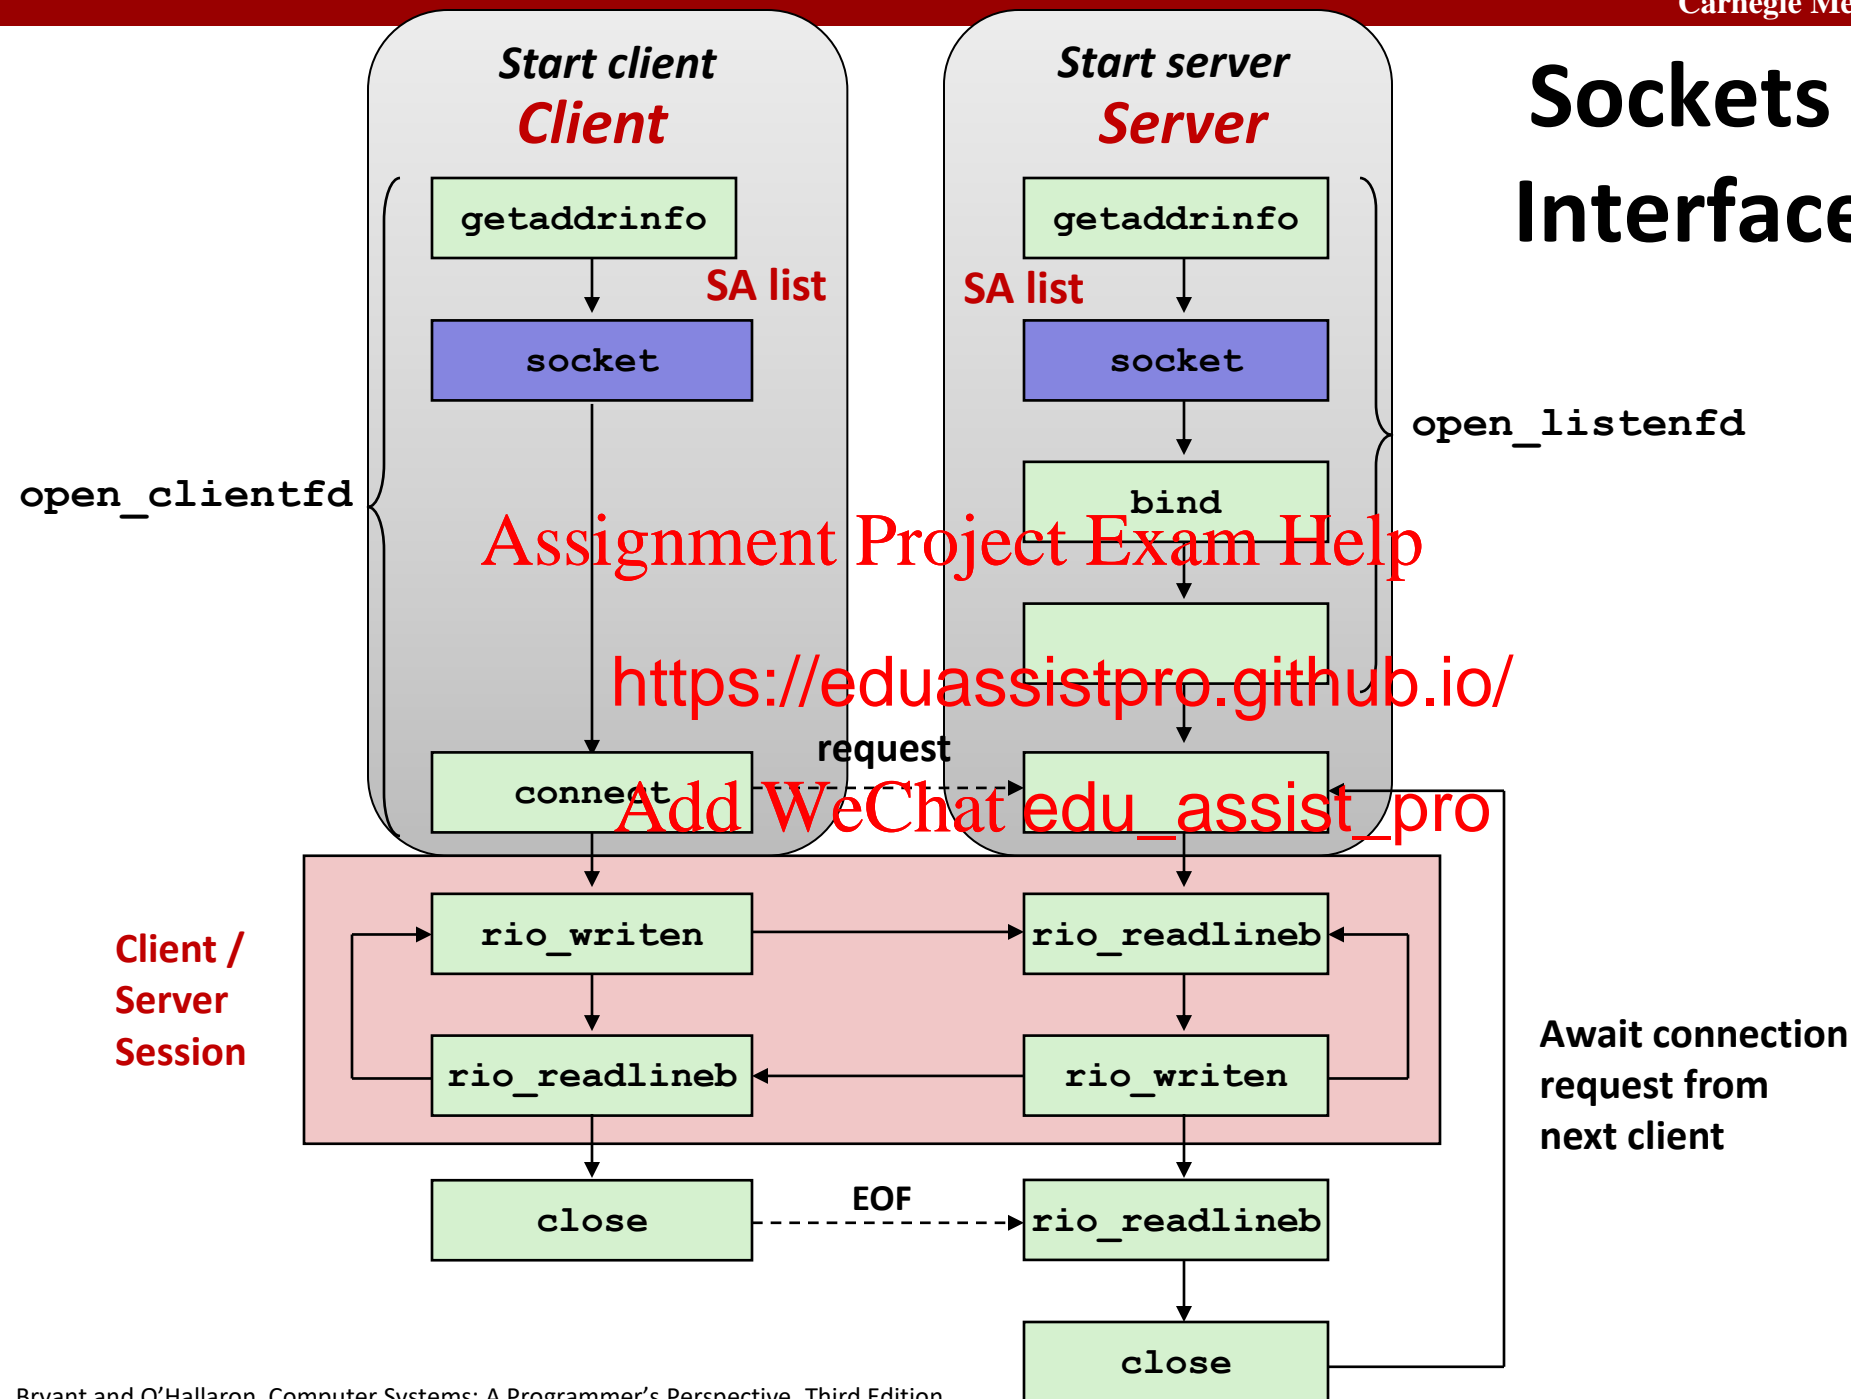
# Review: getaddrinfo

- `getaddrinfo` converts string representations of hostnames, host addresses, ports, service names to socket address structures





# Sockets Interface



# Sockets Interface: `socket`

- Clients and servers use the `socket` function to create a *socket descriptor*:

```
int socket(int domain, int type, int protocol)
```

- Example:

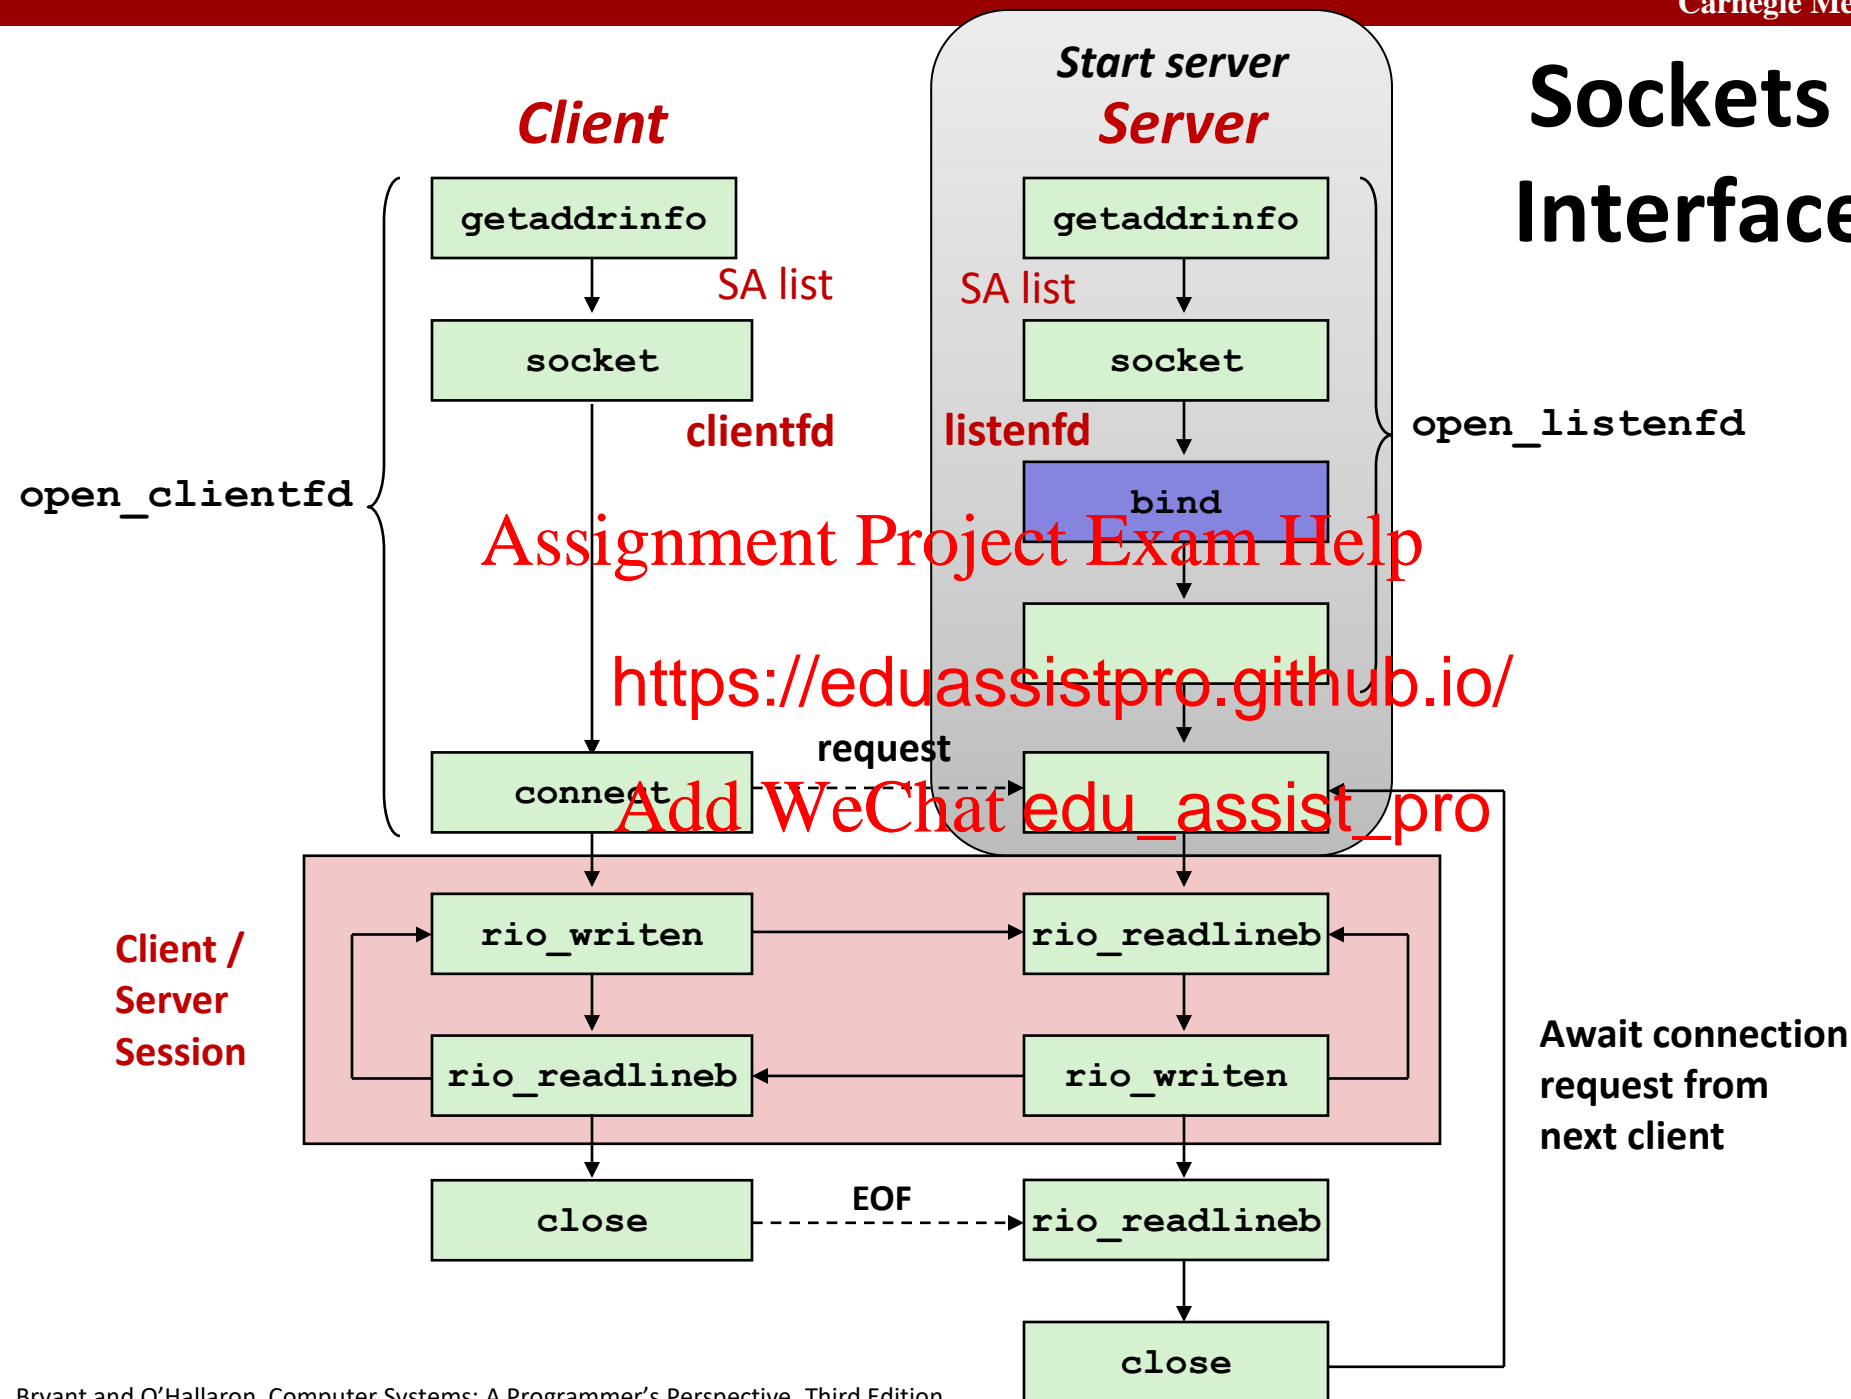
```
int clientfd = socket(AF_INET, SOCK_STREAM, 0);
```

Indicates that we are using  
32-bit IPV4 addresses

the socket  
will be the end point of a  
reliable (TCP) connection

**Protocol specific! Best practice is to use `getaddrinfo` to generate the parameters automatically, so that code is protocol independent.**

# Sockets Interface



# Sockets Interface: `bind`

- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:

```
int bind(int sockfd, SA *addr, socklen_t addrlen);
```

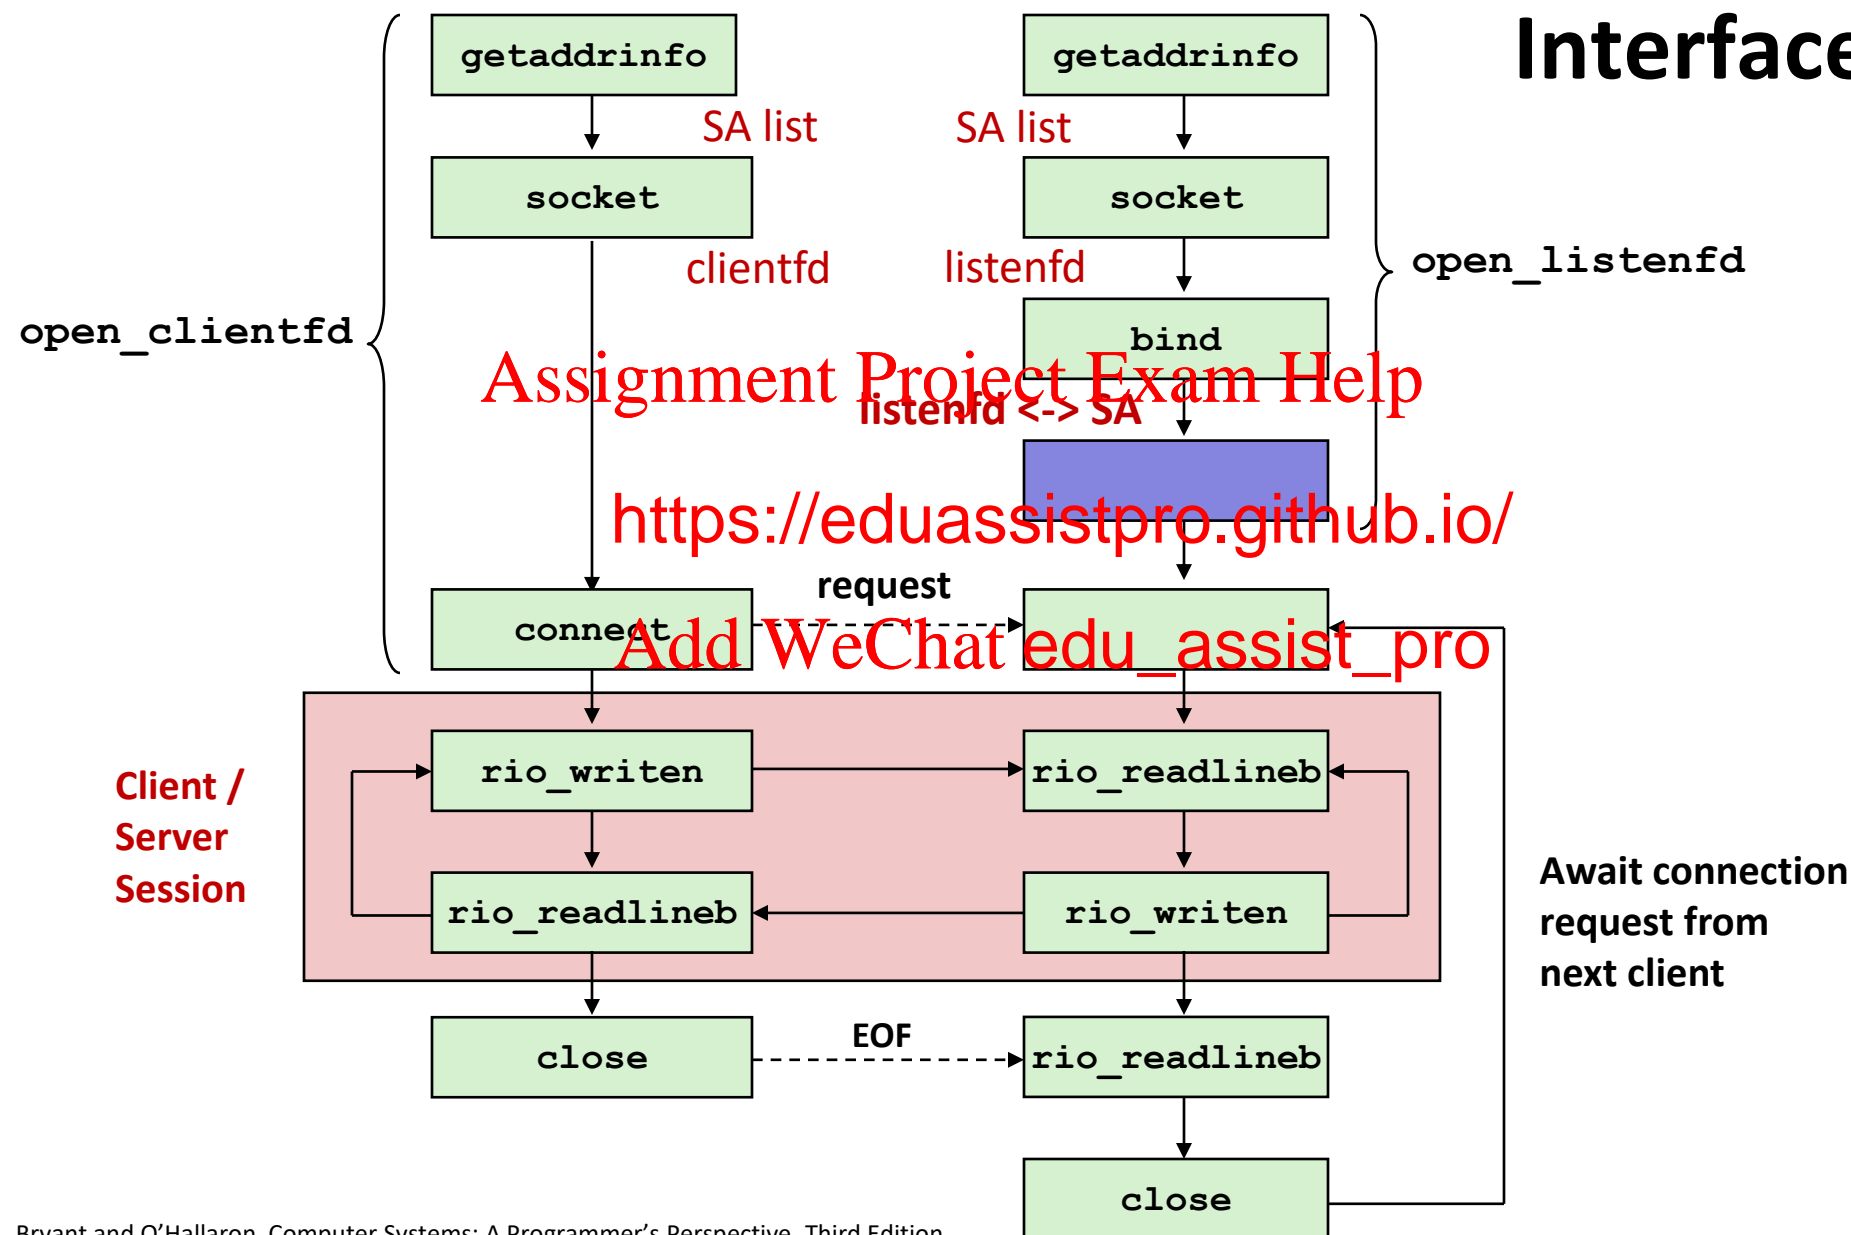
Assignment Project Exam Help  
Our convention: `typedef struct sockaddr SA;`

- Process can read connection whose endpoint is `addr` by reading from `sockfd`
- Similarly, writes to `sockfd` are along connection whose endpoint is `addr`

# Sockets Interface

*Client*

*Server*



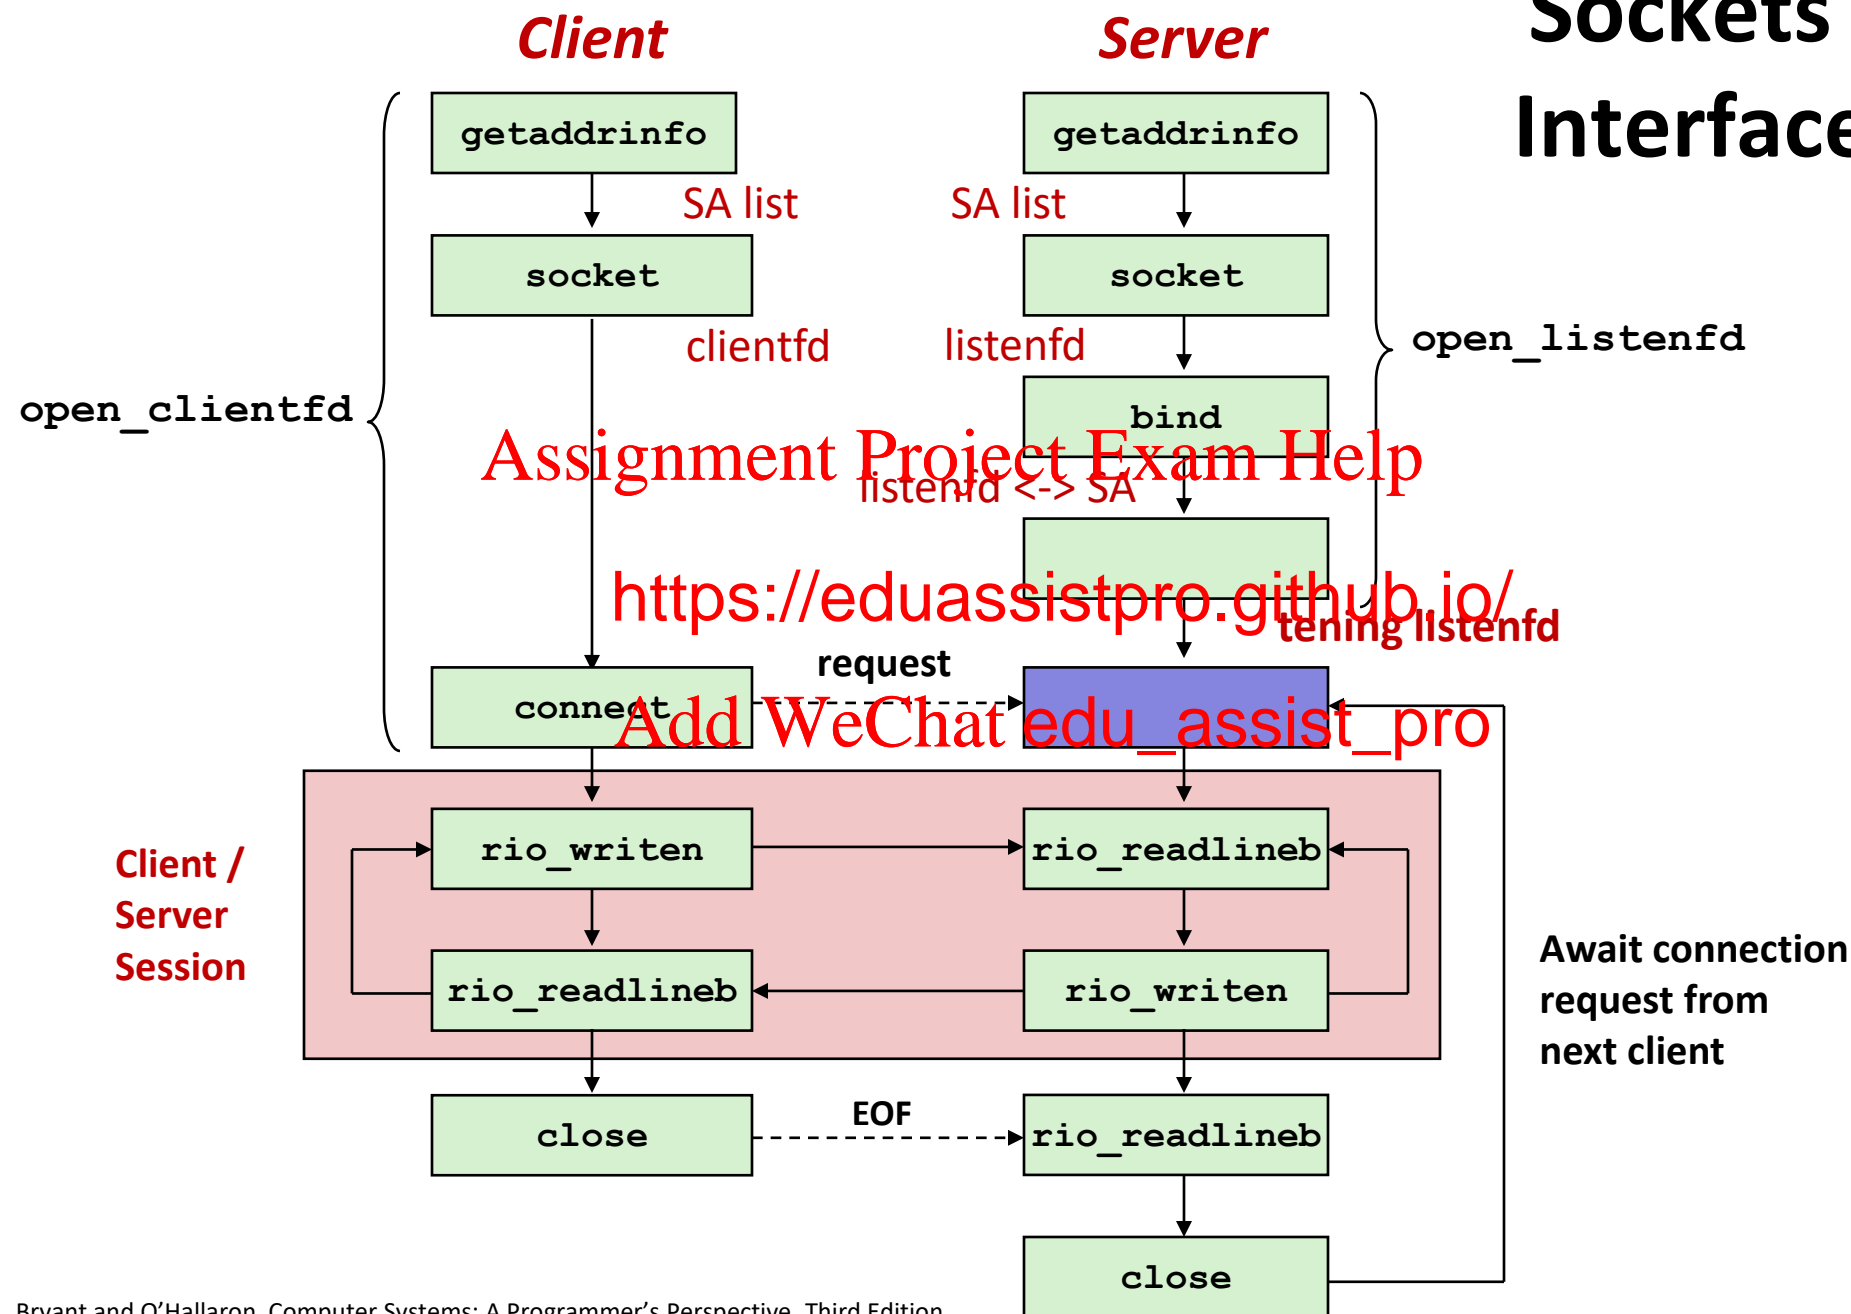
# Sockets Interface: `listen`

- Kernel assumes that descriptor from `socket` function is an **active socket** that will be on the client end
- A server calls the `listen` function to tell the kernel that a descriptor will be used by a server rather than a client:

```
int listen(int sockfd, int backlog);
```

- Converts `sockfd` from an active **socket** to a **listening socket** that can accept connection requests from clients.
- `backlog` is a hint about the number of outstanding connection requests that the kernel should queue up before starting to refuse requests (128-ish by default)

# Sockets Interface



# Sockets Interface: accept

- Servers wait for connection requests from clients by calling `accept`:

```
int accept(int listenfd, SA *addr, int *addrlen);
```

Assignment Project Exam Help

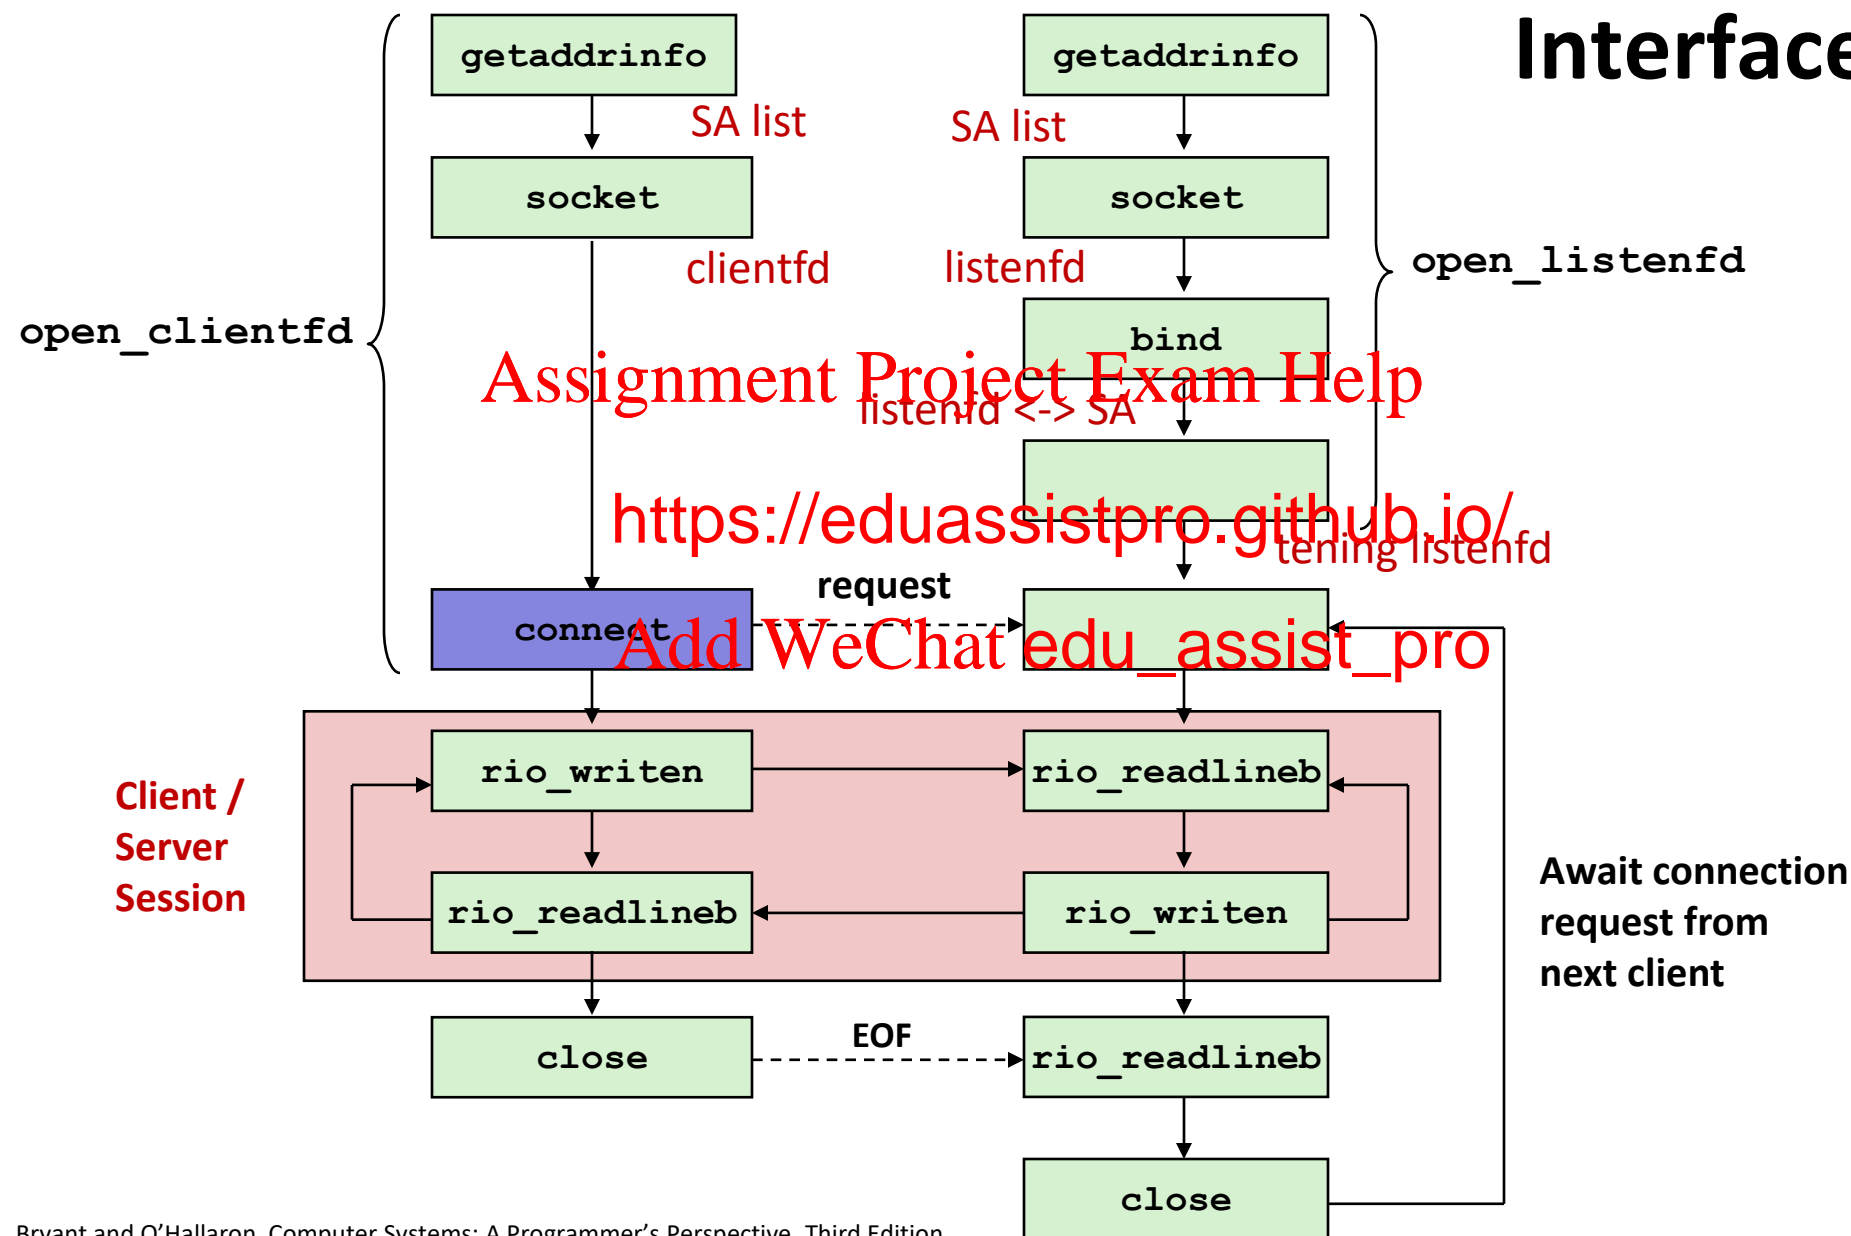
- Waits for connection requests from clients bound to `listenfd`, then fills `addr` with the connection's socket address and `addrlen` with the size of the socket address. <https://eduassistpro.github.io/>  
Add WeChat: edu\_assist\_pro
- Returns a ***connected descriptor*** `connfd` that can be used to communicate with the client via Unix I/O routines.



# Sockets Interface

**Client**

**Server**



# Sockets Interface: connect

- A client establishes a connection with a server by calling **connect**:

```
int connect(int clientfd, SA *addr, socklen_t addrlen);
```

- Attempts to establish a connection with server at socket address **addr**

- If successful, the connection is characterized by reading and writing
- Resulting connection is characterized by

(**x:y**, **addr.sin\_addr:addr.sin\_port**)

- **x** is client address
- **y** is ephemeral port that uniquely identifies client process on client host

Best practice is to use **getaddrinfo** to supply the arguments **addr** and **addrlen**.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu\_assist\_pro

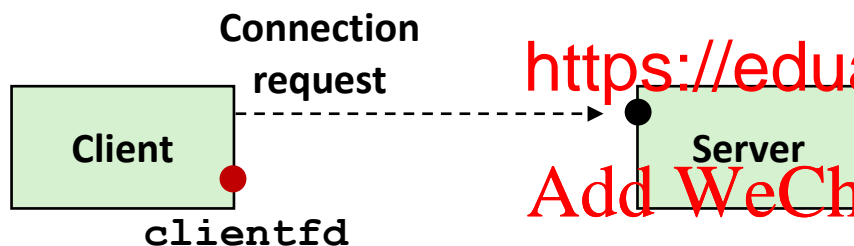
# connect/accept Illustrated



*1. Server blocks in `accept`, waiting for connection request on listening descriptor*

*listenfd*


**Assignment Project Exam Help**



<https://eduassistpro.github.io/>

Add WeChat `edu_assist_pro`

*takes connection request by locking in `connect`*



*3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`*

# Connected vs. Listening Descriptors

## ■ Listening descriptor

- End point for client connection requests
- Created once and exists for lifetime of the server

Assignment Project Exam Help

## ■ Connected descr

- End point of the <https://eduassistpro.github.io/> server
- A new descriptor is created each time the server accepts a connection request from a client
- Exists only as long as it takes to service client

Add WeChat edu\_assist\_pro

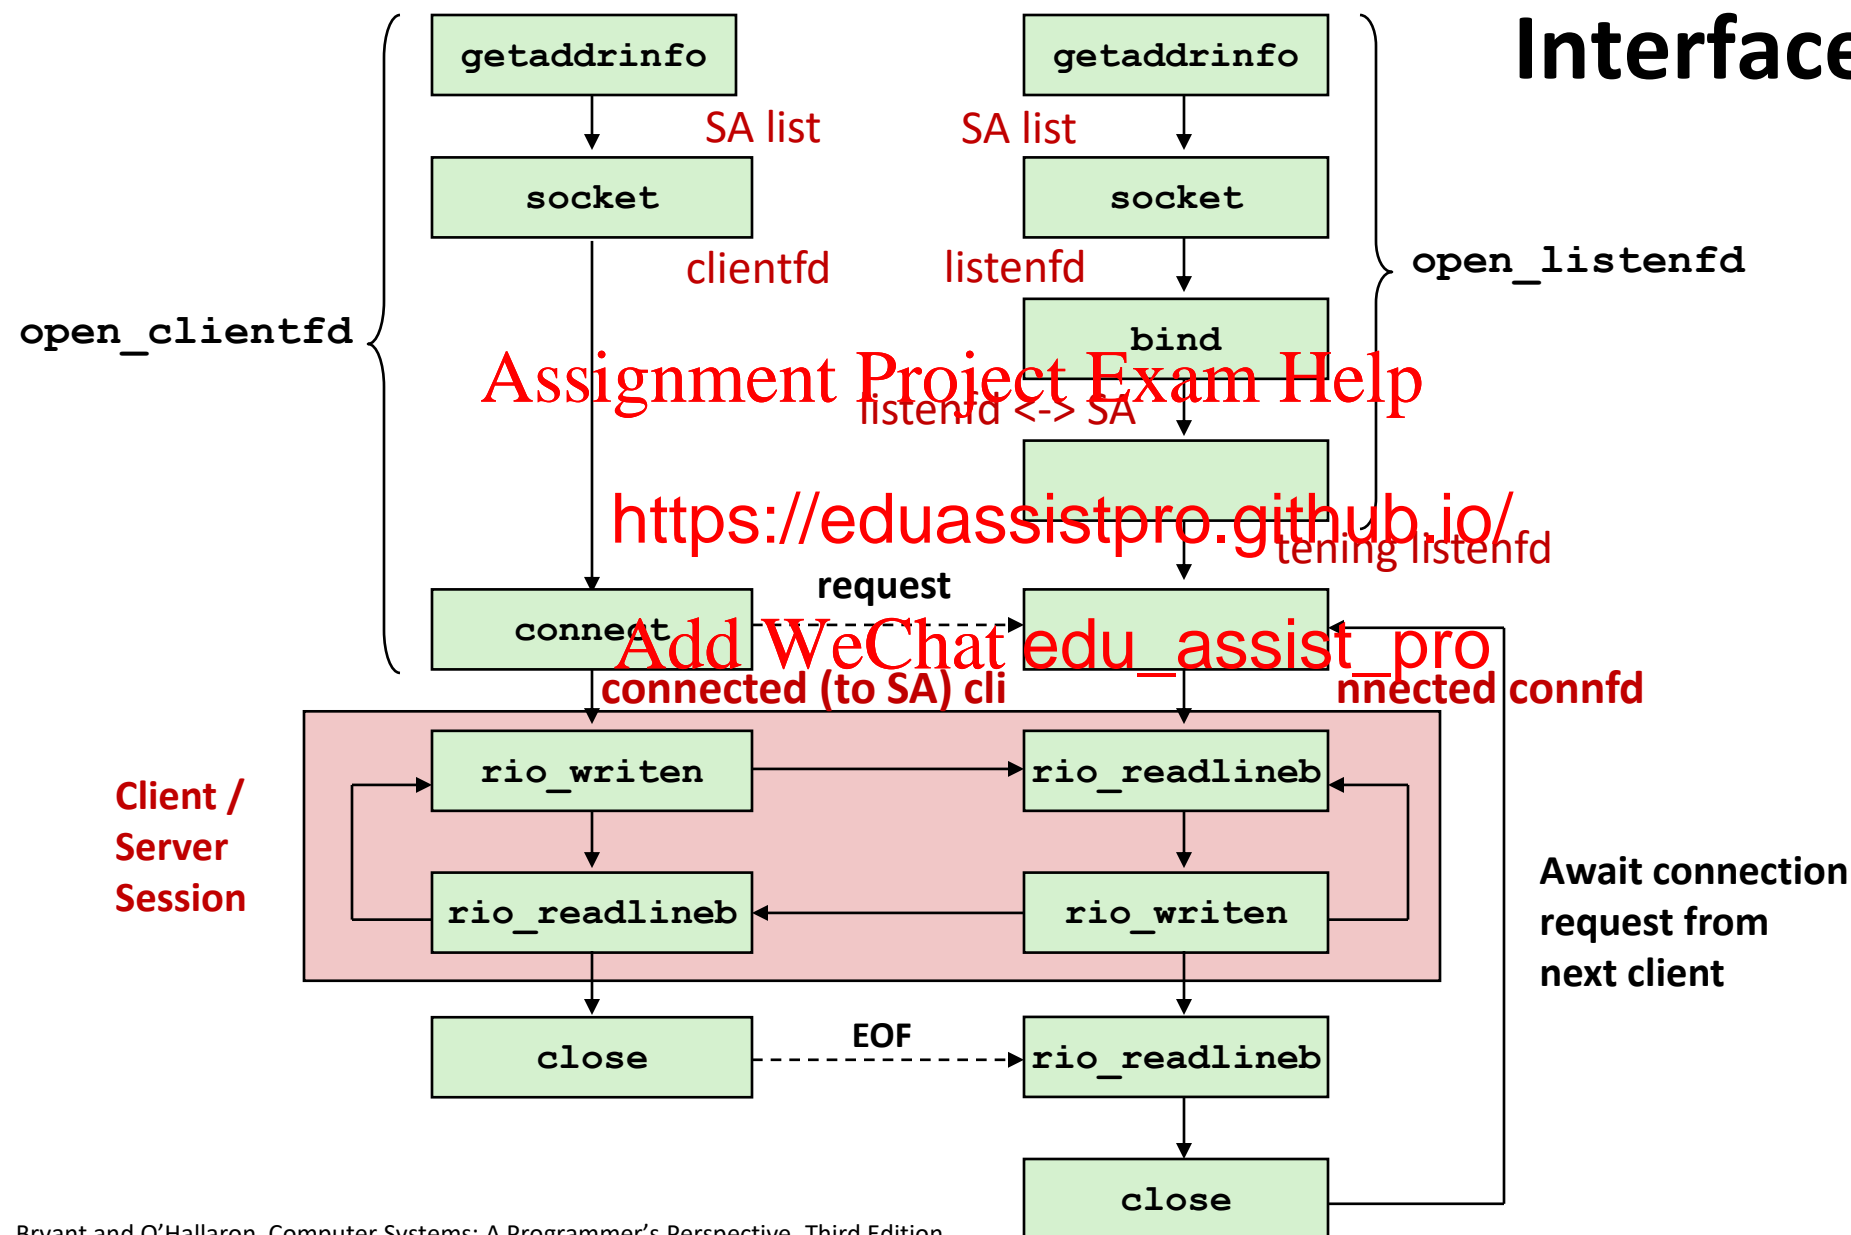
## ■ Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously
  - E.g., Each time we receive a new request, we fork a child to handle the request

# Sockets Interface

**Client**

**Server**



# Sockets Interface

*Client*

*Server*

getaddrinfo

getaddrinfo

socket

socket

open\_listenfd

bind

Assignment Project Exam Help

<https://eduassistpro.github.io/>

connect

request

Add WeChat edu\_assist\_pro

Client /  
Server  
Session

rio\_writen

rio\_readlineb

rio\_readlineb

rio\_writen

Await connection  
request from  
next client

close

EOF

rio\_readlineb

close

# Sockets Helper: `open_clientfd`

- Establish a connection with a server

```
int open_clientfd(char *hostname, char *port) {
    int clientfd;
    struct addrinfo hints, *listp, *p;

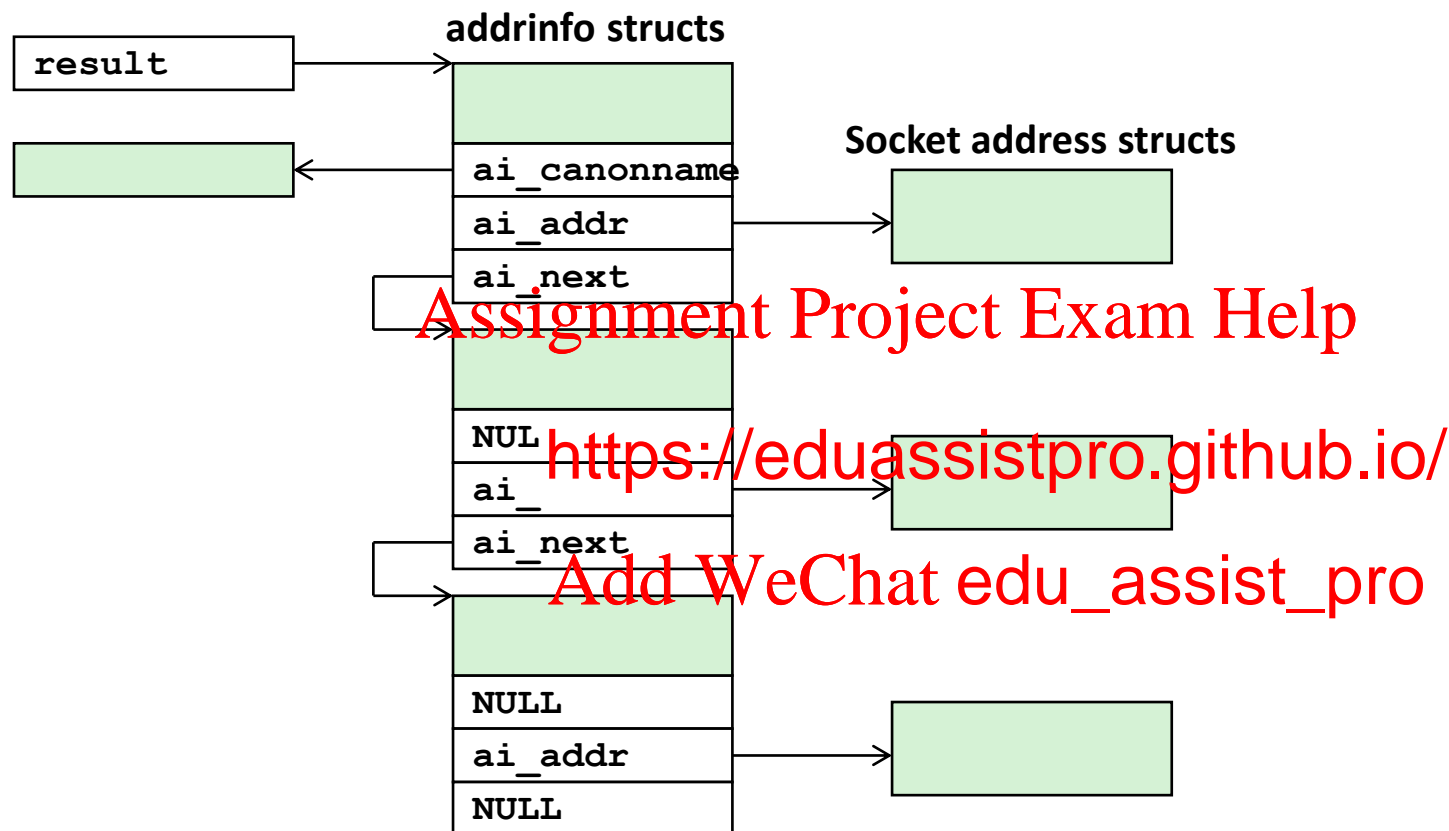
    /* Get a list of potential addresses */
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_socktype = SOCK_STREAM; /* connection */
    hints.ai_flags = AI_NUMERICSERV; /* numeric port arg. */
    hints.ai_flags |= AI_ADDRCONFIG; /* for connections */
    Getaddrinfo(hostname, port, &hints, &listp);
```

csapp.c

**AI\_ADDRCONFIG** – uses your system's address type.

You have at least one IPV4 iface? IPV4. At least one IPV6? IPV6.

# getaddrinfo



- **Clients:** walk this list, trying each socket address in turn, until the calls to `socket` and `connect` succeed.
- **Servers:** walk the list until calls to `socket` and `bind` succeed.



# Sockets Helper: open\_clientfd (cont)

```

/* Walk the list for one that we can successfully connect to */
for (p = listp; p; p = p->ai_next) {
    /* Create a socket descriptor */
    if ((clientfd = socket(p->ai_family, p->ai_socktype,
                          p->ai_protocol)) < 0)
        continue; /* Socket failed, try the next */

    /* Connect to the server */
    if (connect(clientfd, p->ai_addr, p->ai_addrlen) != -1)
        break; /* Success */

    Close(clientfd); /* Connect failed, try another */
}

/* Clean up */
Freeaddrinfo(listp);
if (!p) /* All connects failed */
    return -1;
else /* The last connect succeeded */
    return clientfd;
}

```

csapp.c

# Sockets Interface

*Client*

*Server*

getaddrinfo

getaddrinfo

socket

socket

open\_listenfd

bind

Assignment Project Exam Help

<https://eduassistpro.github.io/>

connect

request

Add WeChat edu\_assist\_pro

rio\_writen

rio\_readlineb

rio\_readlineb

rio\_writen

Await connection  
request from  
next client

close

EOF

rio\_readlineb

close

Client /  
Server  
Session

# Sockets Helper: open\_listenfd

- Create a listening descriptor that can be used to accept connection requests from clients.

```

int open_listenfd(char *port)
{
    struct addrinfo h
    int listenfd, opt

    /* Get a list of potential server
    memset(&hints, 0, sizeof(struct a
    hints.ai_socktype = SOCK_STREAM;
    hints.ai_flags = AI_PASSIVE | AI_ADDRCONFIG;
    hints.ai_flags |= AI_NUMERICSERV;
    Getaddrinfo(NULL, port, &hints, &listp);
  
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

\* Accept connect. \*/  
 /\* ...on any IP addr \*/  
 /\* ...using port no. \*/

csapp.c

# Sockets Helper: open\_listenfd (cont)

```

/* Walk the list for one that we can bind to */
for (p = listp; p; p = p->ai_next) {
    /* Create a socket descriptor */
    if ((listenfd = socket(p->ai_family, p->ai_socktype,
                          p->ai_protocol)) < 0)
        continue; /* Socket failed, try the next */

    /* Eliminates from bind */
    Setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
               (const void *)&opt_int);

    /* Bind the descriptor to the address */
    if (bind(listenfd, p->ai_addr, p->ai_addrlen) == 0)
        break; /* Success */
    Close(listenfd); /* Bind failed, try the next */
}

```

csapp.c

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Sockets Helper: open\_listenfd (cont)

```
/* Clean up */
Freeaddrinfo(listp);
if (!p) /* No address worked */
    return -1;

/* Make it a listening socket ready to accept conn. requests */
if (listen(listenfd,
           Close(listenfd)
           return -1;
}
return listenfd;
}
```

Assignment Project Exam Help  
<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro

csapp.c

- **Key point:** open\_clientfd and open\_listenfd are both independent of any particular version of IP.

# Testing Servers Using `telnet`

- The `telnet` program is invaluable for testing servers that transmit ASCII strings over Internet connections

- Our simple echo server
- Web servers
- Mail servers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Usage: Add WeChat `edu_assist_pro`

- `linux> telnet <host> <portnumber>`
- Creates a connection with a server running on `<host>` and listening on port `<portnumber>`

# Testing the Echo Server With telnet

```
whaleshark> ./echoserveri 15213
Connected to (MAKOSHARK.ICS.CS.CMU.EDU, 50280)
server received 11 bytes
server received 8 bytes
```

## Assignment Project Exam Help

```
makoshark> telnet w 5213
Trying 128.2.210.17
Connected to whales .210.175) .
Escape character is '^]'.
Hi there!
Hi there!
Howdy!
Howdy!
^]
telnet> quit
Connection closed.
makoshark>
```

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Today

- The Sockets Interface
- **Web Servers**
- The Tiny Web Server
- Serving Dynamic

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



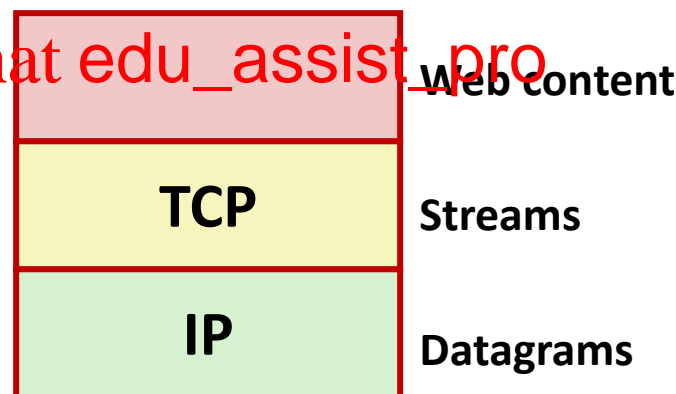
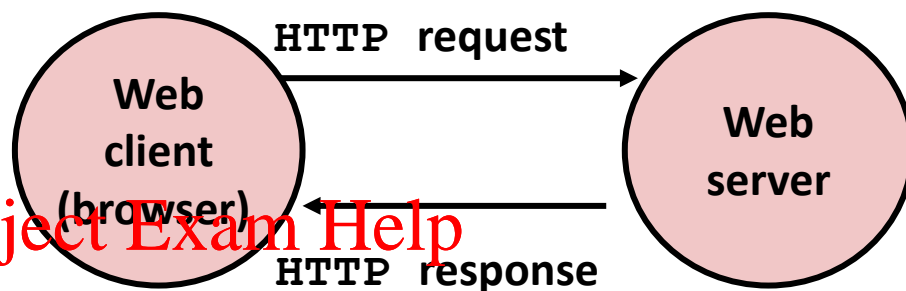
# Web Server Basics

## ■ Clients and servers communicate using the HyperText Transfer Protocol (HTTP)

- Client and server establish TCP connection
- Client requests content
- Server responds with requested content
- Client and server close connection (eventually)

## ■ Current version is HTTP/1.1

- RFC 2616, June, 1999.



<http://www.w3.org/Protocols/rfc2616/rfc2616.html>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Web Content

## ■ Web servers return *content* to clients

- *content*: a sequence of bytes with an associated MIME (Multipurpose Internet Mail Extensions) type

Assignment Project Exam Help

## ■ Example MIME types

- `text/html` <https://eduassistpro.github.io/>
- `text/plain` [Add WeChat edu\\_assist\\_pro](#)
- `image/gif` encoded in GIF format
- `image/png` Binary image encoded in PNG format
- `image/jpeg` Binary image encoded in JPEG format

You can find the complete list of MIME types at:

<http://www.iana.org/assignments/media-types/media-types.xhtml>

# Static and Dynamic Content

- The content returned in HTTP responses can be either *static* or *dynamic*
  - *Static content*: content stored in files and retrieved in response to an HTTP request
    - Examples: HTML, CSS, JavaScript, and other ascript programs
    - Request identifies file containing content
  - *Dynamic content*: content produced in response to an HTTP request
    - Example: content produced by a program executed by the server on behalf of the client
    - Request identifies file containing executable code
- ***Web content associated with a file that is managed by the server***

# URLs and how clients and servers use them

- Unique name for a file: URL (Universal Resource Locator)
- Example URL: `http://www.cmu.edu:80/index.html`
- Clients use *prefix* (`http://www.cmu.edu:80`) to infer:
  - What kind (protocol) of server to contact (HTTP)
  - Where the server is
  - What port it is listening on
- Servers use *suffix* (`/index.html`) to:
  - Determine if request is for static or dynamic content.
    - No hard and fast rules for this
    - One convention: executables reside in `cgi-bin` directory
  - Find file on file system
    - Initial “/” in suffix denotes home directory for requested content.
    - Minimal suffix is “/”, which server expands to configured default filename (usually, `index.html`)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# HTTP Requests

- HTTP request is a *request line*, followed by zero or more *request headers*
- Request line: `<method> <uri> <version>`
  - `<method>` is `GET`, `POST`, `PUT`, `DELETE`, or `HEAD`
  - `<uri>` is typically a URL for proxies, servers
    - A URL is a type of URI (Uniform Resource Identifier)
    - See <http://www.ietf.org/rfc/rfc2396.txt>
  - `<version>` is HTTP version of request (`HTTP/1.0` or `HTTP/1.1`)
- Request headers: `<header name>: <header data>`
  - Provide additional information to the server

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# HTTP Responses

- HTTP response is a **response line** followed by zero or more **response headers**, possibly followed by **content**, with blank line (“\r\n”) separating headers from content.

- Response line: **<version> status msg>**

- <version> is HTTP version
- <status code> is numeric status
- <status msg> is corresponding English message
  - 200 OK Request was handled without error
  - 301 Moved Provide alternate URL
  - 404 Not found Server couldn't find the file

- Response headers: **<header name>: <header data>**

- Provide additional information about response
- **Content-Type:** MIME type of content in response body
- **Content-Length:** Length of content in response body

# Example HTTP Transaction

whaleshark> telnet www.cmu.edu 80

Trying 128.2.42.52...

Connected to WWW-CMU-PROD-VIP.ANDREW.cmu.edu.

Escape character is '^['.

**GET / HTTP/1.1**

Host: www.cmu.edu

HTTP/1.1 **301 Moved Permanently**

Date: Wed, 05 Nov 2014 17:05:11 GMT

Server: Apache/1.3.42 (U

Location: **http://www.cmu**

Transfer-Encoding: chunk

Content-Type: text/html; charset=...

15c

<HTML><HEAD>

...

</BODY></HTML>

0

Connection closed by foreign host.

Client: open connection to server

Telnet prints 3 lines to terminal

Client: request line

Client: required HTTP/1.1 header

Client: blank line terminates headers

Server: response line

Server: followed by 5 response headers

is an Apache server

e has moved here

onse body will be chunked

HTML in response body

line terminates headers

ine in response body

Server: start of HTML content

Server: end of HTML content

Server: last line in response body

Server: closes connection

Assignment Project Exam Help  
<https://eduassistpro.github.io/>  
 Add WeChat edu\_assist\_pro

- HTTP standard requires that each text line end with “\r\n”
- Blank line (“\r\n”) terminates request and response headers

# Example HTTP Transaction, Take 2

whaleshark> telnet www.cmu.edu 80	Client: open connection to server
Trying 128.2.42.52...	Telnet prints 3 lines to terminal
Connected to WWW-CMU-PROD-VIP.ANDREW.cmu.edu.	
Escape character is '^['.	
GET /index.shtml HTTP/1.1	Client: request line
Host: www.cmu.edu	Client: required HTTP/1.1 header
	Client: blank line terminates headers
HTTP/1.1 200 OK	Server: response line
Date: Wed, 05 Nov 2014 1	Server: followed by 4 response headers
Server: Apache/1.3.42 (U	
Transfer-Encoding: chunked	
Content-Type: text/html; charset=.	Server: blank line terminates headers
1000	Server: begin response body
<html ..>	Server: first line of HTML content
...	
</html>	
0	Server: end response body
Connection closed by foreign host.	Server: close connection

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Example HTTP(S) Transaction, Take 3

```

whaleshark> openssl s_client www.cs.cmu.edu:443
CONNECTED(00000005)
...
Certificate chain
...
-
Server certificate
-----BEGIN CERTIFICATE-----
MIIGDjCCBPagAwIBAgIRAMiF70BPPdySiInNoU+mp+gwDQYJKoZIhvcNAQELBQAw
djELMAkGA1UEBhMCVVMxCzAJBgNVBAGTAk1JMRIWEAYDVQQHEw1Bbm4gQXJib3Ix
EjAQBgNVBAoTCUluGvybmV0HzAdBgNVBAMT
wkWkvDVBBCwKXrShVxQNs6J
-----END CERTIFICATE-----
subject=/C=US/postalCode=15213/ST=PA/L=Pittsburgh t=5000 Forbes
Ave/O=Carnegie Mellon University/OU=School of Computer
Science/CN=www.cs.cmu.edu issuer=/C=US/ST=MI/L=Ann
Arbor/O=Internet2/OU=InCommon/CN=InCommon RSA Server CA
SSL handshake has read 6274 bytes and written 483 bytes
...
>GET / HTTP/1.0

HTTP/1.1 200 OK
Date: Tue, 12 Nov 2019 04:22:15 GMT
Server: Apache/2.4.10 (Ubuntu)
Set-Cookie: SHIBLOCATION=scsweb; path=/; domain=.cs.cmu.edu
... HTML Content Continues Below ...

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Quiz Time!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Check out:

Add WeChat edu\_assist\_pro

<https://canvas.cmu.edu/courses/17808>

# Today

- The Sockets Interface
- Web Servers
- The Tiny Web Server
- Serving Dynamic Content

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Tiny Web Server

## ■ Tiny Web server described in text

- Tiny is a sequential Web server
- Serves static and dynamic content to real browsers
  - text files, HTML files, GIF, PNG, and JPEG images
- 239 lines of code
- Not as complete server
  - You can break it with poorly formatted requests (e.g., terminate lines with “\n” instead of “\r\n”)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Tiny Operation

- Accept connection from client
- Read request from client (via connected socket)
- Split into `<method> <uri> <version>`
  - If method not GET, then return error
- If URI contains “`https://eduassistpro.github.io/`”
  - (Would do wrong thing if had file “`o.html`”)
  - Fork process to execute program
- Otherwise serve static content
  - Copy file to output

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Tiny Serving Static Content

```

void serve_static(int fd, char *filename, int filesize)
{
    int srcfd;
    char *srcp, filetype[MAXLINE], buf[MAXBUF];

    /* Send response headers to client */
    get_filetype(filename, filetype);
    sprintf(buf, "H
    sprintf(buf, "%s", buf);
    sprintf(buf, "%s
    sprintf(buf, "Content-length: %d", filesize);
    sprintf(buf, "Content-type: %s", filetype);
    Rio_writen(fd, buf, strlen(buf)

    /* Send response body to client */
    srcfd = Open(filename, O_RDONLY, 0);
    srcp = Mmap(0, filesize, PROT_READ, MAP_PRIVATE, srcfd, 0);
    Close(srcfd);
    Rio_writen(fd, srcp, filesize);
    Munmap(srcp, filesize);
}

```

tiny.c

# Today

- The Sockets Interface
- Web Servers
- The Tiny Web Server
- **Serving Dynamic**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Serving Dynamic Content

- Client sends request to server

GET /cgi-bin/env.pl HTTP/1.1


Assignment Project Exam Help

- If request URI co

string "/cgi-bin

server assumes that the

request is for dynamic content




<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# Serving Dynamic Content (cont)

- The server creates a child process and runs the program identified by the URI in that process



Assignment Project Exam Help

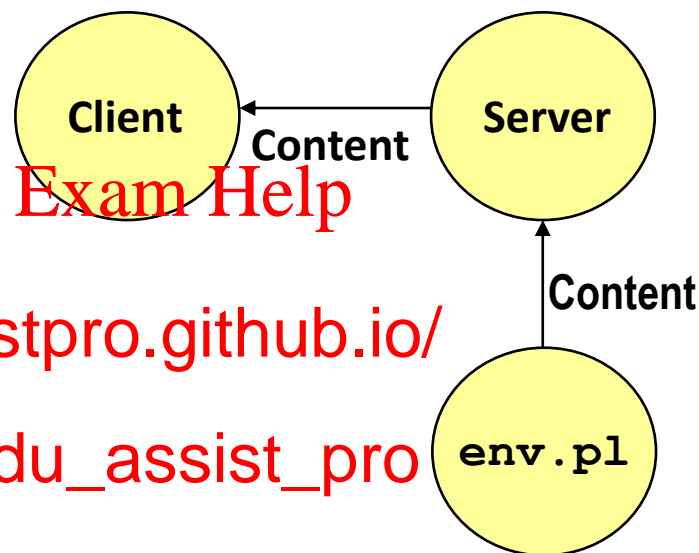
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Serving Dynamic Content (cont)

- The child runs and generates the dynamic content

- The server captures content of the child and forwards it without modification to the client



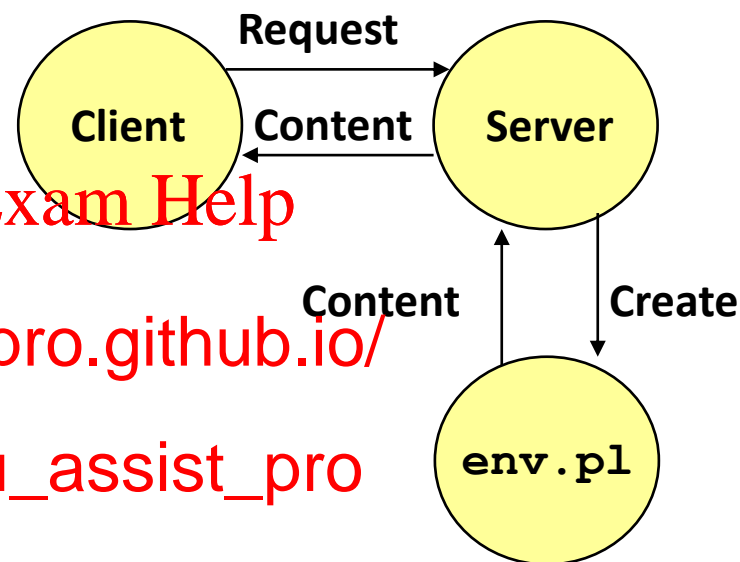
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Issues in Serving Dynamic Content

- How does the client pass program arguments to the server?
- How does the server pass these arguments to the child?
- How does the server pass the request to the child?
- How does the server capture the content produced by the child?
- These issues are addressed by the **Common Gateway Interface (CGI)** specification.



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# CGI

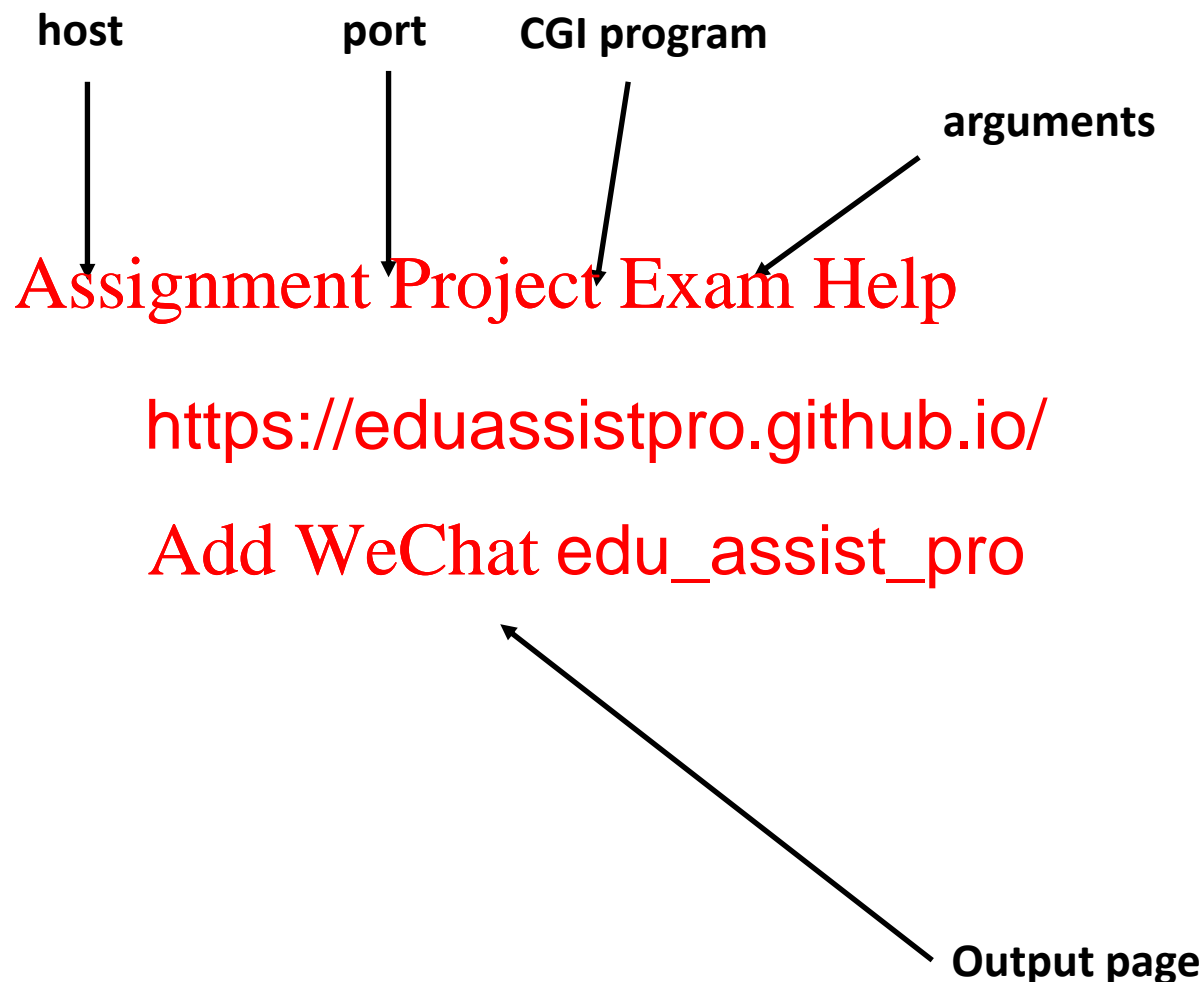
- Because the children are written according to the CGI spec, they are often called **CGI programs**.
- However, CGI really defines a simple standard for transferring information (browser), the server, and t
- CGI is the original standard for generating dynamic content. Has been largely replaced by other, faster techniques:
  - E.g., fastCGI, Apache modules, Java servlets, Rails controllers
  - Avoid having to create process on the fly (expensive and slow).

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# The add.com Experience



# Serving Dynamic Content With GET

- Question: How does the client pass arguments to the server?
- Answer: The arguments are appended to the URI
- Can be encoded directly in a URL typed to a browser or a URL in an HTML link
  - `http://adder` <https://eduassistpro.github.io/213818213>
  - `adder` is the CGI program on the server. I do the addition.
  - argument list starts with “?”
  - arguments separated by “&”
  - spaces represented by “+” or “%20”

# Serving Dynamic Content With GET

- URL suffix:
  - `cgi-bin/adder?15213&18213`

- Result displayed on browser:

Welcome to <https://eduassistpro.github.io/> ternet  
addition portal.  
Add WeChat edu\_assist\_pro  
The answer is:  $15213 + 18213 = 33426$   
Thanks for visiting!

# Serving Dynamic Content With GET

■ Question: How does the server pass these arguments to the child?

■ Answer: In environment variable QUERY\_STRING

- A single string containing everything after the “?”
- For add: QUE 3”

<https://eduassistpro.github.io/>

```

/* Extract the two arguments
if ((buf = getenv("QUERY_STRING")) != NULL) {
    p = strchr(buf, '&');
    *p = '\0';
    strcpy(arg1, buf);
    strcpy(arg2, p+1);
    n1 = atoi(arg1);
    n2 = atoi(arg2);
}
    
```

adder.c



# Serving Dynamic Content with GET

- Question: How does the server capture the content produced by the child?
- Answer: The child generates its output on `stdout`. Server uses `dup2` to redirect `stdout` to its connected socket.

```
void serve_dynamic(int fd, char *filename, char *cgiargs)
```

```
{
```

```
    char buf[MAXLINE],
```

```
    /* Return first pa
```

```
    sprintf(buf, "HTTP/1.0 200 OK\r\n"
```

```
    Rio_writen(fd, buf, strlen(buf));
```

```
    sprintf(buf, "Server: Tiny Web Ser
```

```
    Rio_writen(fd, buf, strlen(buf));
```

```
    if (Fork() == 0) { /* Child */
```

```
        /* Real server would set all CGI vars here */
```

```
        setenv("QUERY_STRING", cgiargs, 1);
```

```
        Dup2(fd, STDOUT_FILENO);
```

```
        /* Redirect stdout to client */
```

```
        Execve(filename, emptylist, environ); /* Run CGI program */
```

```
    }
```

```
    Wait(NULL); /* Parent waits for and reaps child */
```

```
}
```

# Serving Dynamic Content with GET

- Notice that only the CGI child process knows the content type and length, so it must generate those headers.

```

/* Make the response body */
sprintf(content, "Welcome to add.com: ");
sprintf(content, "                                rtal.\r\n<p>", content);
sprintf(content, "https://eduassistpro.github.io/\r\n<p>",
          content, n1, n2, n1 + n2);
sprintf(content, "%sThanks for visiting\r\n", content);

/* Generate the HTTP response */
printf("Content-length: %d\r\n", (int)strlen(content));
printf("Content-type: text/html\r\n\r\n");
printf("%s", content);
fflush(stdout);

exit(0);

```

adder.c

# Serving Dynamic Content With GET

```
bash:makoshark> telnet whaleshark.ics.cs.cmu.edu 15213
Trying 128.2.210.175...
Connected to whaleshark.ics.cs.cmu.edu (128.2.210.175).
Escape character is '^]'.
GET /cgi-bin/adder?15213&18213 HTTP/1.0
```

*HTTP request sent by client*

```
HTTP/1.0 200 OK
Server: Tiny Web Ser
Connection: close
Content-length: 117
Content-type: text/html
```

*HTTP response generated  
by the server*

```
Welcome to add.com: THE Internet addition portal.
<p>The answer is: 15213 + 18213 = 33426
<p>Thanks for visiting!
Connection closed by foreign host.
bash:makoshark>
```

*HTTP response generated  
by the CGI program*

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# For More Information

- **W. Richard Stevens et. al. “Unix Network Programming: The Sockets Networking API”, Volume 1, Third Edition, Prentice Hall, 2003**
  - THE network programming bible
- **Michael Kerrisk, “Linux in a Nutshell”, No Starch Press, 2010**
  - THE Linux programming bible.
- **Complete versions of all code in the 213 schedule page.**
  - `http://www.cs.cmu.edu/~213/schedule.html`
  - `csapp.{c,h}`, `hostinfo.c`, `echoclient.c`, `echoserveri.c`, `tiny.c`, `adder.c`
  - You can use any of this code in your assignments.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

is available

# Additional slides

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Web History

## ■ 1989:

- Tim Berners-Lee (CERN) writes internal proposal to develop a distributed hypertext system
  - Connects “a web of notes with links”
  - Intended to manage info

Assignment Project Exam Help

<https://eduassistpro.github.io/>

## ■ 1990:

- Tim BL writes a graphical brows

Add WeChat edu\_assist\_pro  
achines

# Web History (cont)

## ■ 1992

- NCSA server released
- 26 WWW servers worldwide

## ■ 1993

Assignment Project Exam Help

- Marc Andreess
- Mosaic version
- Web (port 80) traffic at 1% of NS
- Over 200 WWW servers worldwide

CSA Mosaic browser  
(Macintosh, Windows, Unix)

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

## ■ 1994

- Andreessen and colleagues leave NCSA to form “Mosaic Communications Corp” (predecessor to Netscape)

# HTTP Versions

## ■ Major differences between HTTP/1.1 and HTTP/1.0

- HTTP/1.0 uses a new connection for each transaction
- HTTP/1.1 also supports *persistent connections*
  - multiple transactions over the same connection
  - Connecti
- HTTP/1.1 requir
  - Host: `www.cmu.edu`
  - Makes it possible to host multi single Internet host
- HTTP/1.1 supports *chunked encoding*
  - Transfer-Encoding: `chunked`
- HTTP/1.1 adds additional support for caching

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro



# GET Request to Apache Server From Firefox Browser

URI is just the suffix, not the entire URL

```
GET /~bryant/test.html HTTP/1.1
Host: www.cs.cmu.edu
User-Agent: Mozilla/5.0 (Windows NT 6.0; en-US;
rv:1.9.2.11) Gecko/20110607 Firefox/3.5.1
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-us,en;q=0.5
Accept-Encoding: gzip,deflate
Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
Keep-Alive: 115
Connection: keep-alive
CRLF (\r\n)
```

Assignment Project Exam Help  
<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro

# GET Response From Apache Server

```
HTTP/1.1 200 OK
Date: Fri, 29 Oct 2010 19:48:32 GMT
Server: Apache/2.2.14 (Unix) mod_ssl/2.2.14 OpenSSL/0.9.7m
mod_pubcookie/3.3.2b PHP/5.3.1
Accept-Ranges: bytes
Content-Length: 47
Keep-Alive: timeout=5
Connection: Keep-Alive
Content-Type: text/html

<html>
<head><title>Some Tests</title></head>

<body>
<h1>Some Tests</h1>
. . .
</body>
</html>
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Data Transfer Mechanisms

## ■ Standard

- Specify total length with content-length
- Requires that program buffer entire message

## ■ Chunked **Assignment Project Exam Help**

- Break into block
- Prefix each block

<https://eduassistpro.github.io/>  
ded)

Add WeChat edu\_assist\_pro

# Chunked Encoding Example

```
HTTP/1.1 200 OK\n
Date: Sun, 31 Oct 2010 20:47:48 GMT\n
Server: Apache/1.3.41 (Unix)\n
Keep-Alive: timeout=15, max=100\n
Connection: Keep-Alive\n
Transfer-Encoding: chunked\n
Content-Type: text/html\n
\r\n
```

```
d75\r\n
```

**First Ch**

```
<html>
<head>
.<link href="http://.css" rel="stylesheet"
type="text/css">
</head>
<body id="calendar_body">

<div id='calendar'><table width='100%' border='0' cellpadding='0'
cellspacing='1' id='cal'>

. . .
</body>
</html>
\r\n
```

```
0\r\n
```

**Second Chunk: 0 bytes (indicates last chunk)**

```
\r\n
```

Assignment Project Exam Help


<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Proxies

- A **proxy** is an intermediary between a client and an **origin server**
  - To the client, the proxy acts like a server
  - To the server, the proxy acts like a client

## Assignment Project Exam Help



# Why Proxies?

- Can perform useful functions as requests and responses pass by
  - Examples: Caching, logging, anonymization, filtering, transcoding

Assignment Project Exam Help

<https://eduassistpro.github.io/>  
Add WeChat edu\_assist\_pro

