

Synchronization: Basics

Assignment Project Exam Help

15-213: Introduction

25th Lecture, April 1 <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Today

- **Threads review**
- **Sharing**
- **Mutual exclusion**
- **Semaphores**

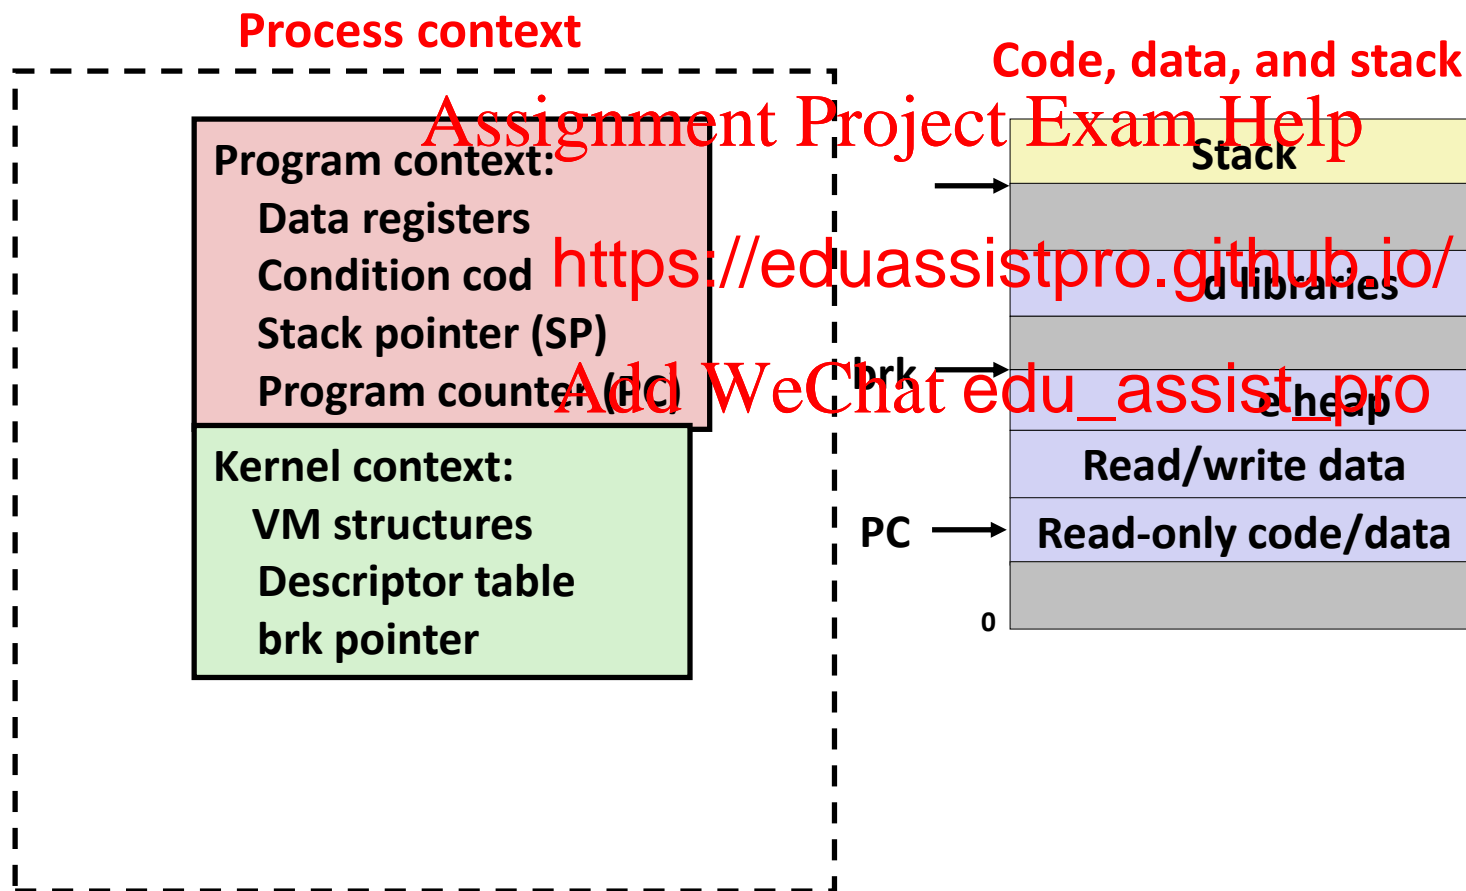
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

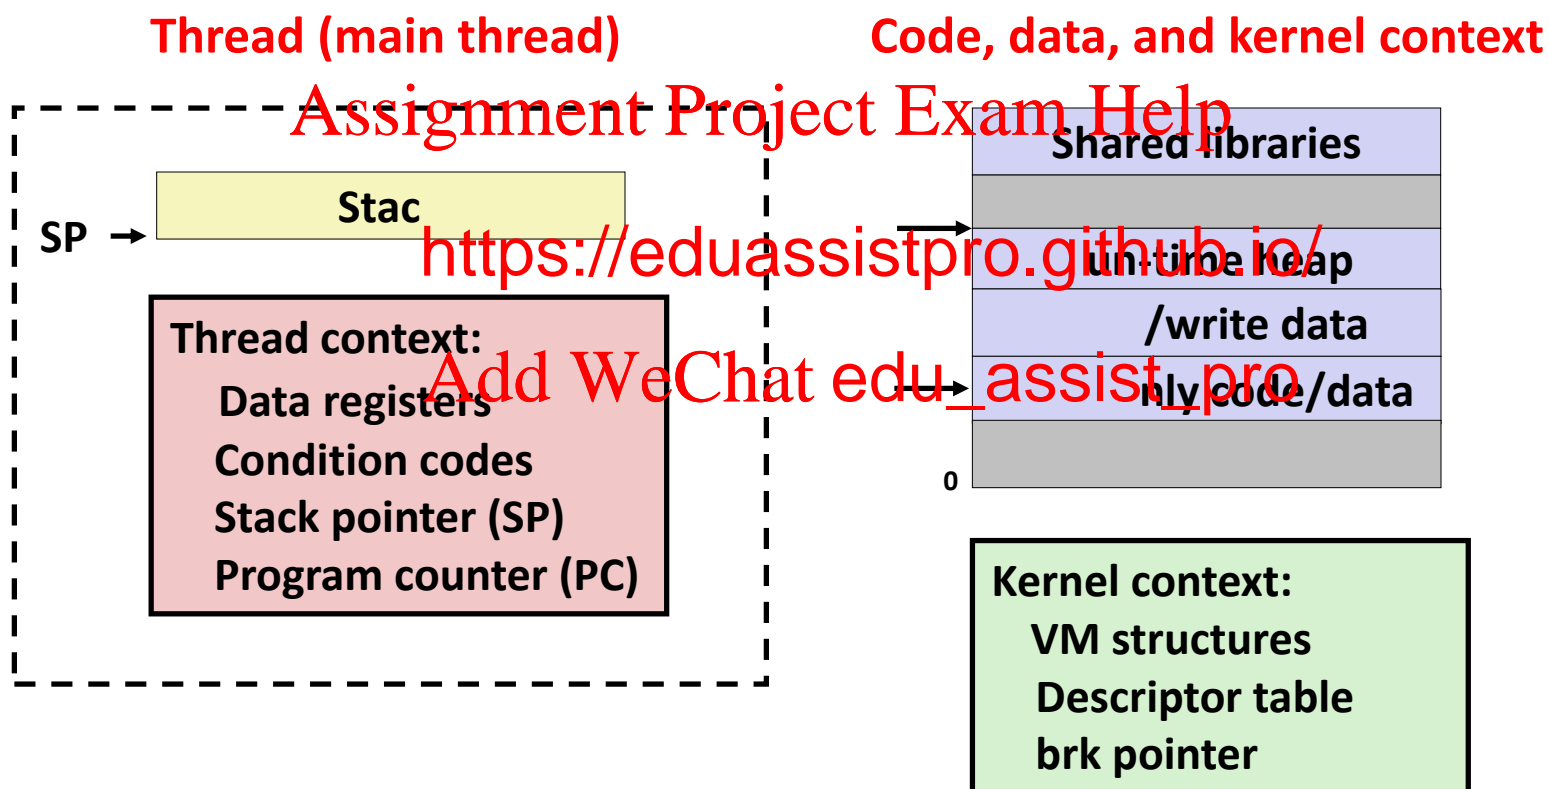
Traditional View of a Process

- Process = process context + code, data, and stack



Alternate View of a Process

- Process = thread + (code, data, and kernel context)



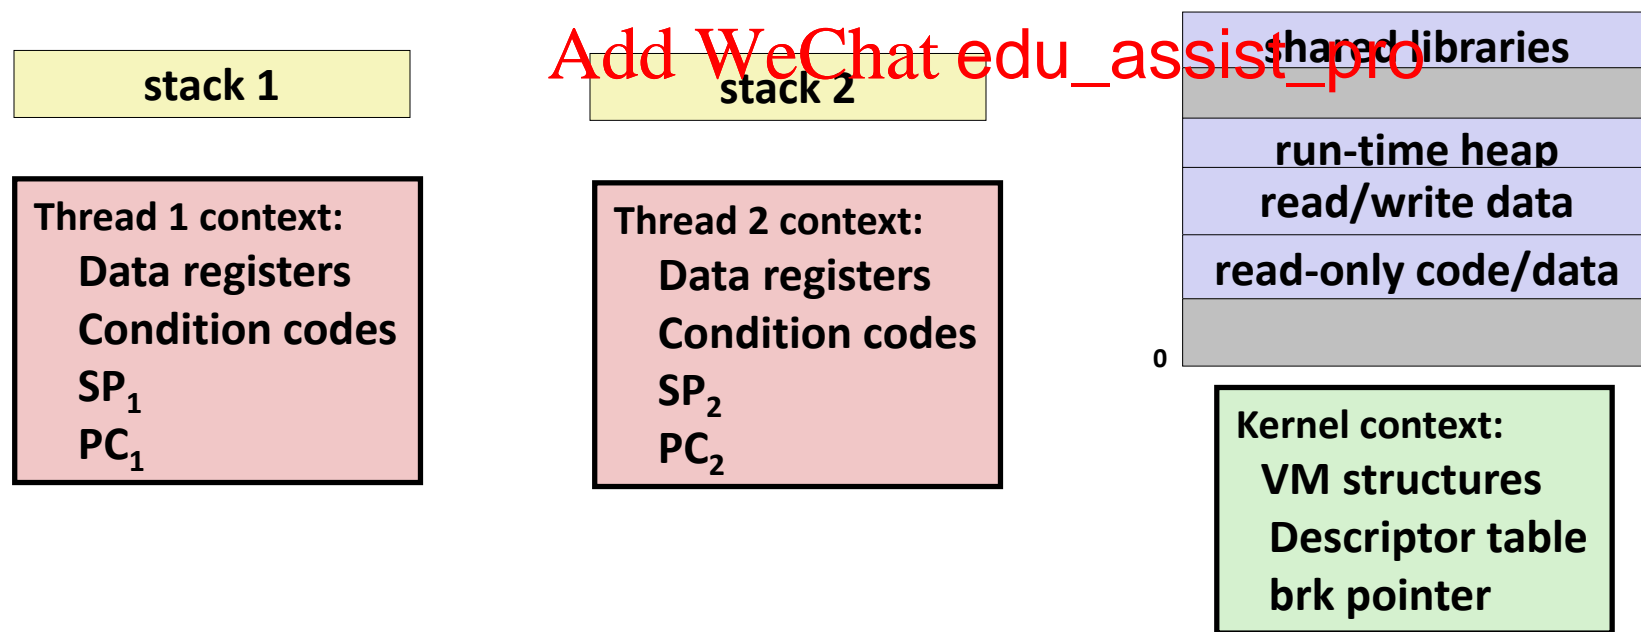
A Process With Multiple Threads

- **Multiple threads can be associated with a process**
 - Each thread has its own logical control flow
 - Each thread shares the same code, data, and kernel context
 - Each thread has its own stack for local variables
 - but not protected from other threads
 - Each thread has its own thread id (TID)

Assignment Project Exam Help

Thread 1 (main thread) <https://eduassistpro.github.io/> and data

Add WeChat edu_assist_pro



Don't let picture confuse you!

Assignment Project Exam Help

Thread 1 (main thread) <https://eduassistpro.github.io/> and data

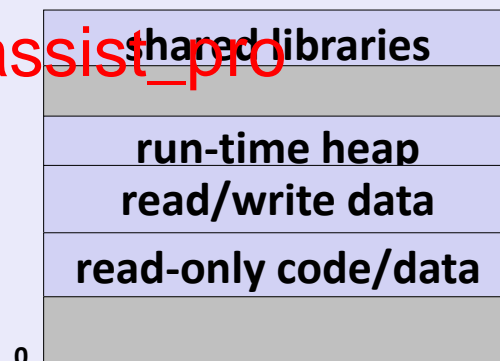
Add WeChat edu_assist_pro

stack 1

stack 2

Thread 1 context:
Data registers
Condition codes
 SP_1
 PC_1

Thread 2 context:
Data registers
Condition codes
 SP_2
 PC_2



Kernel context:
VM structures
Descriptor table
brk pointer

Memory is shared between all threads

Today

- Threads review

- **Sharing**

- Mutual exclusion

- Semaphores

- Producer-Consumer <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu_assist_pro

Shared Variables in Threaded C Programs

■ Question: Which variables in a threaded C program are shared?

- The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”

Assignment Project Exam Help

■ Def: A variable x <https://eduassistpro.github.io/> multiple threads reference some instance of x .

Add WeChat edu_assist_pro

■ Requires answers to the following questions:

- What is the memory model for threads?
- How are instances of variables mapped to memory?
- How many threads might reference each of these instances?

Threads Memory Model: Conceptual

- Multiple threads run within the context of a single process
- Each thread has its own separate thread context
 - Thread ID, stack, stack pointer, PC, condition codes, and GP registers
- All threads share the remaining process context
 - Code, data, heap, and shared virtual address space
 - Open files and installed device drivers

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Shared code and data

Thread 1
(private)

stack 1

Thread 1 context:
Data registers
Condition codes
 SP_1
 PC_1

Thread 2
(private)

stack 2

Thread 2 context:
Data registers
Condition codes
 SP_2
 PC_2

shared libraries

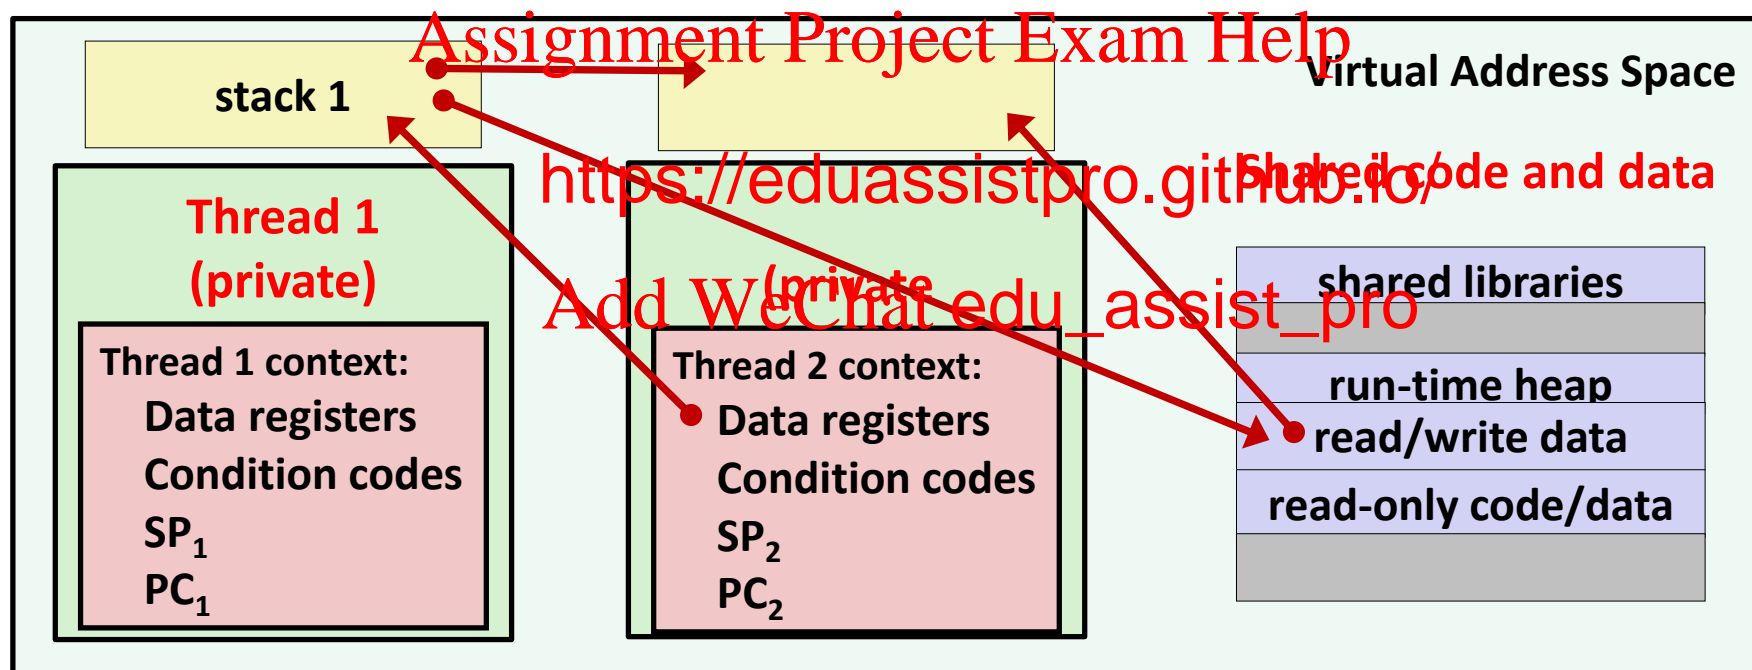
run-time heap
read/write data

read-only code/data

Threads Memory Model: Actual

■ Separation of data is not strictly enforced:

- Register values are truly separate and protected, but...
- Any thread can read and write the stack of any other thread



*The mismatch between the conceptual and operation model
is a source of confusion and errors*

Passing an argument to a thread - Pedantic

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++) {
        long* p = Malloc
        *p = i;
        Pthread_create(&tids[i],
            NULL,
            thread,
            (void *)p);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[*(long *)vargp] += 1;
    Free(vargp);
    return NULL;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
void main() {
    for (int i=0; i<N; i++) {
        if (hist[i] != 1) {
            printf("Failed at %d\n", i);
            exit(-1);
        }
    }
    printf("OK\n");
}
```

Passing an argument to a thread - Pedantic

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++) {
        long* p = Malloc
        *p = i;
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *)p);
    }
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[* (long *)vargp] += 1;
    Free(vargp);
    return NULL;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

alloc to create a per
d heap allocated

place in memory for the
argument

- Remember to free in thread!
- Producer-consumer pattern

Passing an argument to a thread – Also OK!

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&
                        NULL,
                        thread,
                        (void *)i);

    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[(long)vargp] += 1;
    return NULL;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

se cast since
ong) <= sizeof(void*)

- Cast does NOT change bits

Passing an argument to a thread – **WRONG!**

```

int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&
                        NULL,
                        thread,
                        (void *) &i);

    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}

```

```

void *thread(void *vargp)
{
    hist[*(long*)vargp] += 1;
    return NULL;
}

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

ints to same location
threads!

- Creates a data race!

Three Ways to Pass Thread Arg

■ Malloc/free

- Producer malloc's space, passes pointer to `pthread_create`
- Consumer dereferences pointer

■ Ptr to stack slot

- Producer passes `pthread_create`
- Consumer deref

■ Cast of int

- Producer casts an int/long to address in `pthread_create`
- Consumer casts `void*` argument back to int/long

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Example Program to Illustrate Sharing

```

char **ptr;  /* global var */

int main(int argc, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from fo",
        "Hello from ba"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}

```

sharing.c

```

void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;
    printf("[%ld]:  %s (cnt=%d)\n",
        id, ptr[myid], ++cnt);
    return NULL;
}

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

*reference main thread's stack
indirectly through global ptr variable*

*A common way to pass a single
argument to a thread routine*

Shared Variables in Threaded C Programs

■ Question: Which variables in a threaded C program are shared?

- The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”

Assignment Project Exam Help

■ Def: A variable x <https://eduassistpro.github.io/> multiple threads reference some instance of x .

Add WeChat edu_assist_pro

■ Requires answers to the following questions:

- What is the memory model for threads?
- How are instances of variables mapped to memory?
- How many threads might reference each of these instances?

Mapping Variable Instances to Memory

■ Global variables

- *Def*: Variable declared outside of a function
- **Virtual memory contains exactly one instance of any global variable**

Assignment Project Exam Help

■ Local variables

- *Def*: Variable de <https://eduassistpro.github.io/> `static` attribute
- Each thread stack contains one ins `local` variable

Add WeChat edu_assist_pro

■ Local static variables

- *Def*: Variable declared inside function with the `static` attribute
- **Virtual memory contains exactly one instance of any local static variable.**

Mapping Variable Instances to Memory

```
char **ptr;    /* global var */

int main(int main, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
void thread(void *vargp)
{
    int myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

Mapping Variable Instances to Memory

Global var: 1 instance (ptr [data])

Local vars: 1 instance (i.m, msgs.m, tid.m)

```
char **ptr; /* global var */

int main(int main, char *argv[])
{
    long i;
    pthread_t tid;
    char *msgs[2] = {
        "Hello from foo",
        "Hello from bar"
    };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
            NULL,
            thread,
            (void *)i);
    Pthread_exit(NULL);
}
```

sharing.c

Local var: 2 instances (
myid.p0 [peer thread 0's stack],
myid.p1 [peer thread 1's stack]
)

```
void *thread(void *vargp)
{
    = (long)vargp;

    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
        myid, ptr[myid], ++cnt);
    return NULL;
}
```

Local static var: 1 instance (cnt [data])

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Shared Variable Analysis

■ Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>
<code>ptr</code>	yes	yes	yes
<code>cnt</code>	no	yes	yes
<code>i.m</code>	yes	no	no
<code>msgs.m</code>	yes		es
<code>myid.p0</code>	no		o
<code>myid.p1</code>	no	no	

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
char **ptr; /* global var */
int main(int main, char *argv[]) {
    long i; pthread_t tid;
    char *msgs[2] = {"Hello from foo",
                     "Hello from bar" };

    ptr = msgs;
    for (i = 0; i < 2; i++)
        Pthread_create(&tid,
                       NULL, thread, (void *)i);
    Pthread_exit(NULL);}
```

```
void *thread(void *vargp)
{
    long myid = (long)vargp;
    static int cnt = 0;

    printf("[%ld]: %s (cnt=%d)\n",
           myid, ptr[myid], ++cnt);
    return NULL;
}
```

Shared Variable Analysis

■ Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>
<code>ptr</code>	yes	yes	yes
<code>cnt</code>	no	yes	yes
<code>i.m</code>	yes		no
<code>msgs.m</code>	yes		yes
<code>myid.p0</code>	no	yes	
<code>myid.p1</code>	no	no	

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

■ Answer: A variable **x** is shared iff multiple threads reference at least one instance of **x**. Thus:

- `ptr`, `cnt`, and `msgs` are shared
- `i` and `myid` are **not** shared

Synchronizing Threads

- Shared variables are handy...
- ...but introduce the possibility of nasty *synchronization* errors.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

badcnt.c: Improper Synchronization

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1,
        thread, &niters);
    Pthread_create(&tid2,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);
    for (i = 0; i < niters; i++)
        cnt++;
    return NULL;
}

```

```

linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>

```

cnt should equal 20,000.

What went wrong?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

Assembly Code for Counter Loop

C code for counter loop in thread i

```
for (i = 0; i < niters; i++)
    cnt++;
```

Assignment Project Exam Help

Asm code for thread i

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

<code>jle .L2</code>	} : Head
<code>movl \$0, %eax</code>	

<code>.L3:</code>	} L_i : Load cnt U_i : Update cnt S_i : Store cnt
<code>movq cnt(%rip), %rdx</code>	
<code>addq \$1, %rdx</code>	
<code>movq %rdx, cnt(%rip)</code>	} T_i : Tail
<code>addq \$1, %rax</code>	
<code>cmpq %rcx, %rax</code>	
<code>jne .L3</code>	
<code>.L2:</code>	

Concurrent Execution

- **Key idea:** In general, any **sequentially consistent*** interleaving is possible, but some give an unexpected result!

- I_i denotes that thread i executes instruction I
- $\%rdx_i$ is the content of $\%rdx$ in thread i 's context

Assignment Project Exam Help

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

i (thread)	$instr_i$			
1	H_1			
1	L_1	0	-	
1	U_1	1		
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2

OK

**For now. In reality, on x86 even non-sequentially consistent interleavings are possible*

Concurrent Execution

- **Key idea:** In general, any sequentially consistent interleaving is possible, but some give an unexpected result!

- I_i denotes that thread i executes instruction I
- $\%rdx_i$ is the content of $\%rdx$ in thread i 's context

Assignment Project Exam Help

i (thread) $instr_i$

1	H_1			
1	L_1	0	-	
1	U_1	1	-	
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Thread 1

critical section



Thread 2

critical section

OK

Concurrent Execution (cont)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2

i (thread)	instr _i	%rdx ₁	%rdx ₂	cnt
1	H ₁	-	-	0
1	L ₁			
1	U ₁			
2	H ₂			
2	L ₂		0	
1	S ₁	1	-	
1	T ₁	1	-	1
2	U ₂	-	1	1
2	S ₂	-	1	1
2	T ₂	-	1	1

Oops!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Concurrent Execution (cont)

■ How about this ordering?

i (thread)	instr _i	%rdx ₁	%rdx ₂	cnt
1	H ₁			0
1	L ₁	0		
2	H ₂			
2	L ₂			
2	U ₂			
2	S ₂		1	
1	U ₁	1		
1	S ₁	1		
1	T ₁			1
2	T ₂			1

Assignment Project Exam Help

<https://eduassistpro.github.io/>

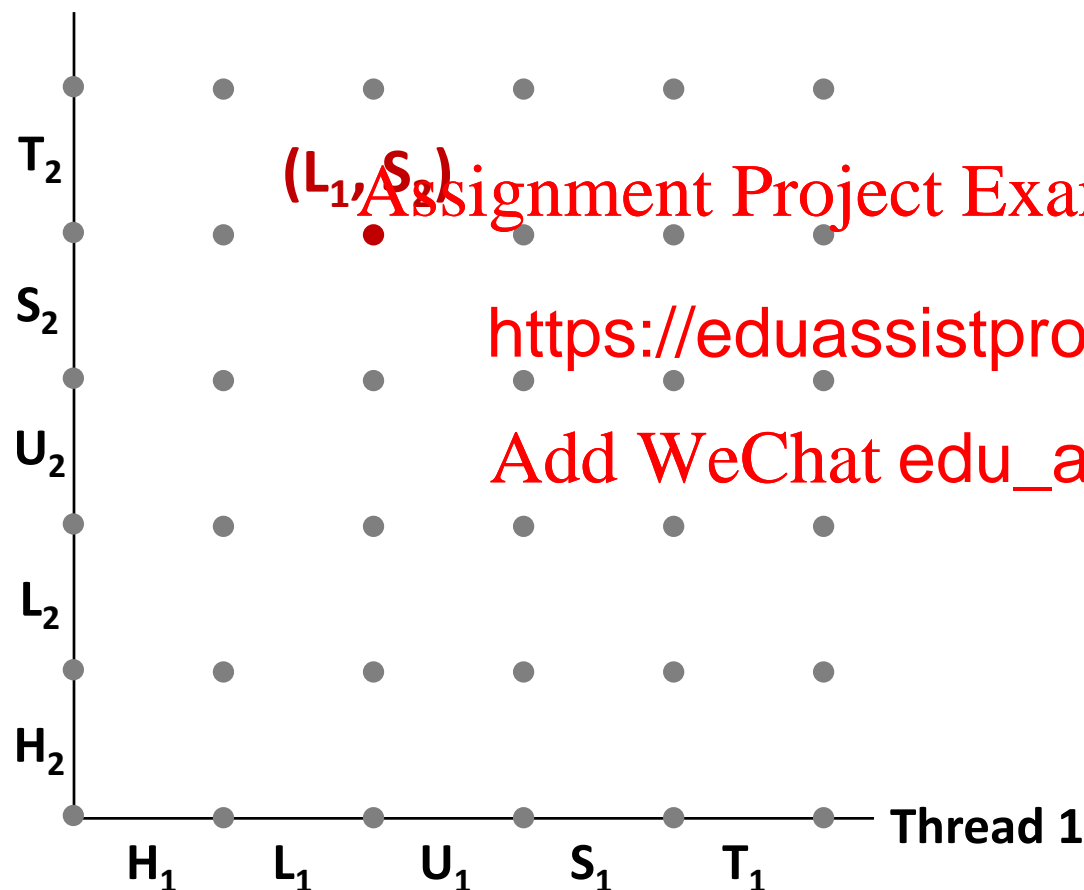
Add WeChat edu_assist_pro

Oops!

■ We can analyze the behavior using a *progress graph*

Progress Graphs

Thread 2



A **progress graph** depicts the discrete **execution state space** of concurrent threads.

Each axis corresponds to the sequential order of instructions in a thread.

Each point corresponds to a possible **execution state** $(Inst_1, Inst_2)$.

E.g., (L_1, S_2) denotes state where thread 1 has completed L_1 and thread 2 has completed S_2 .

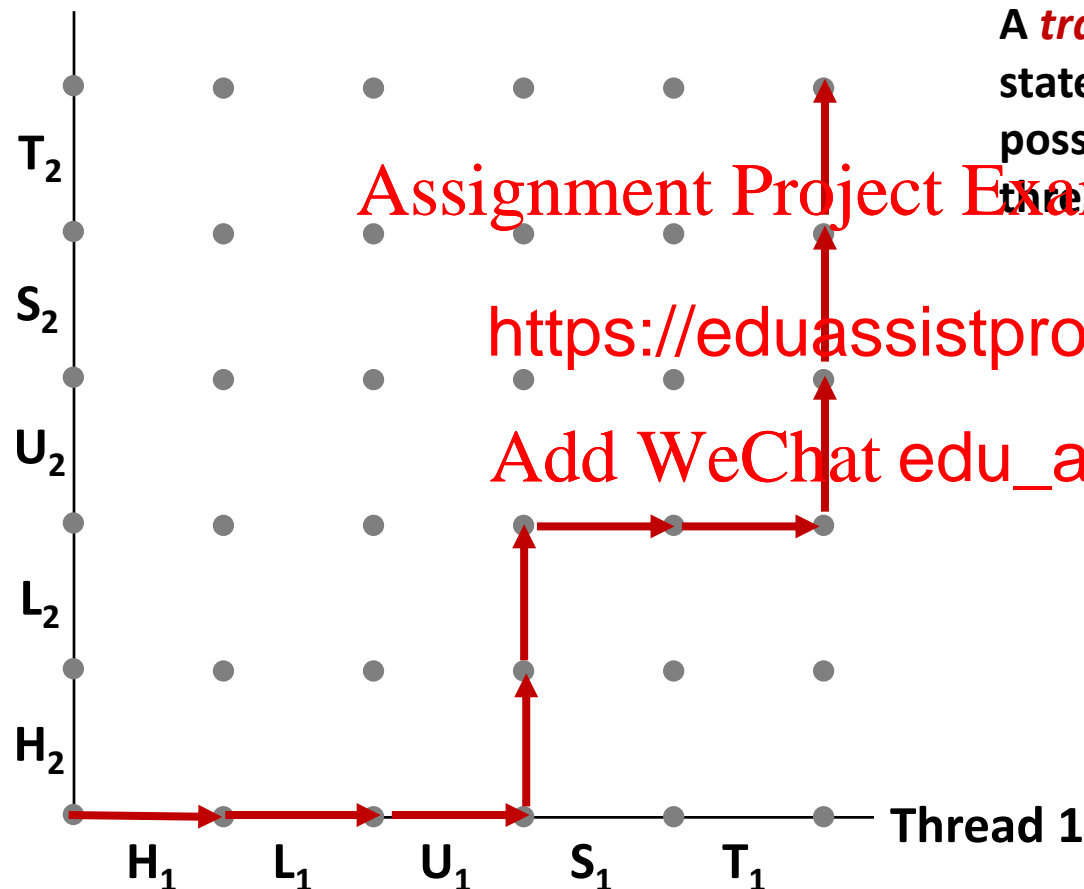
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Trajectories in Progress Graphs

Thread 2



A **trajectory** is a sequence of legal state transitions that describes one possible concurrent execution of the threads.

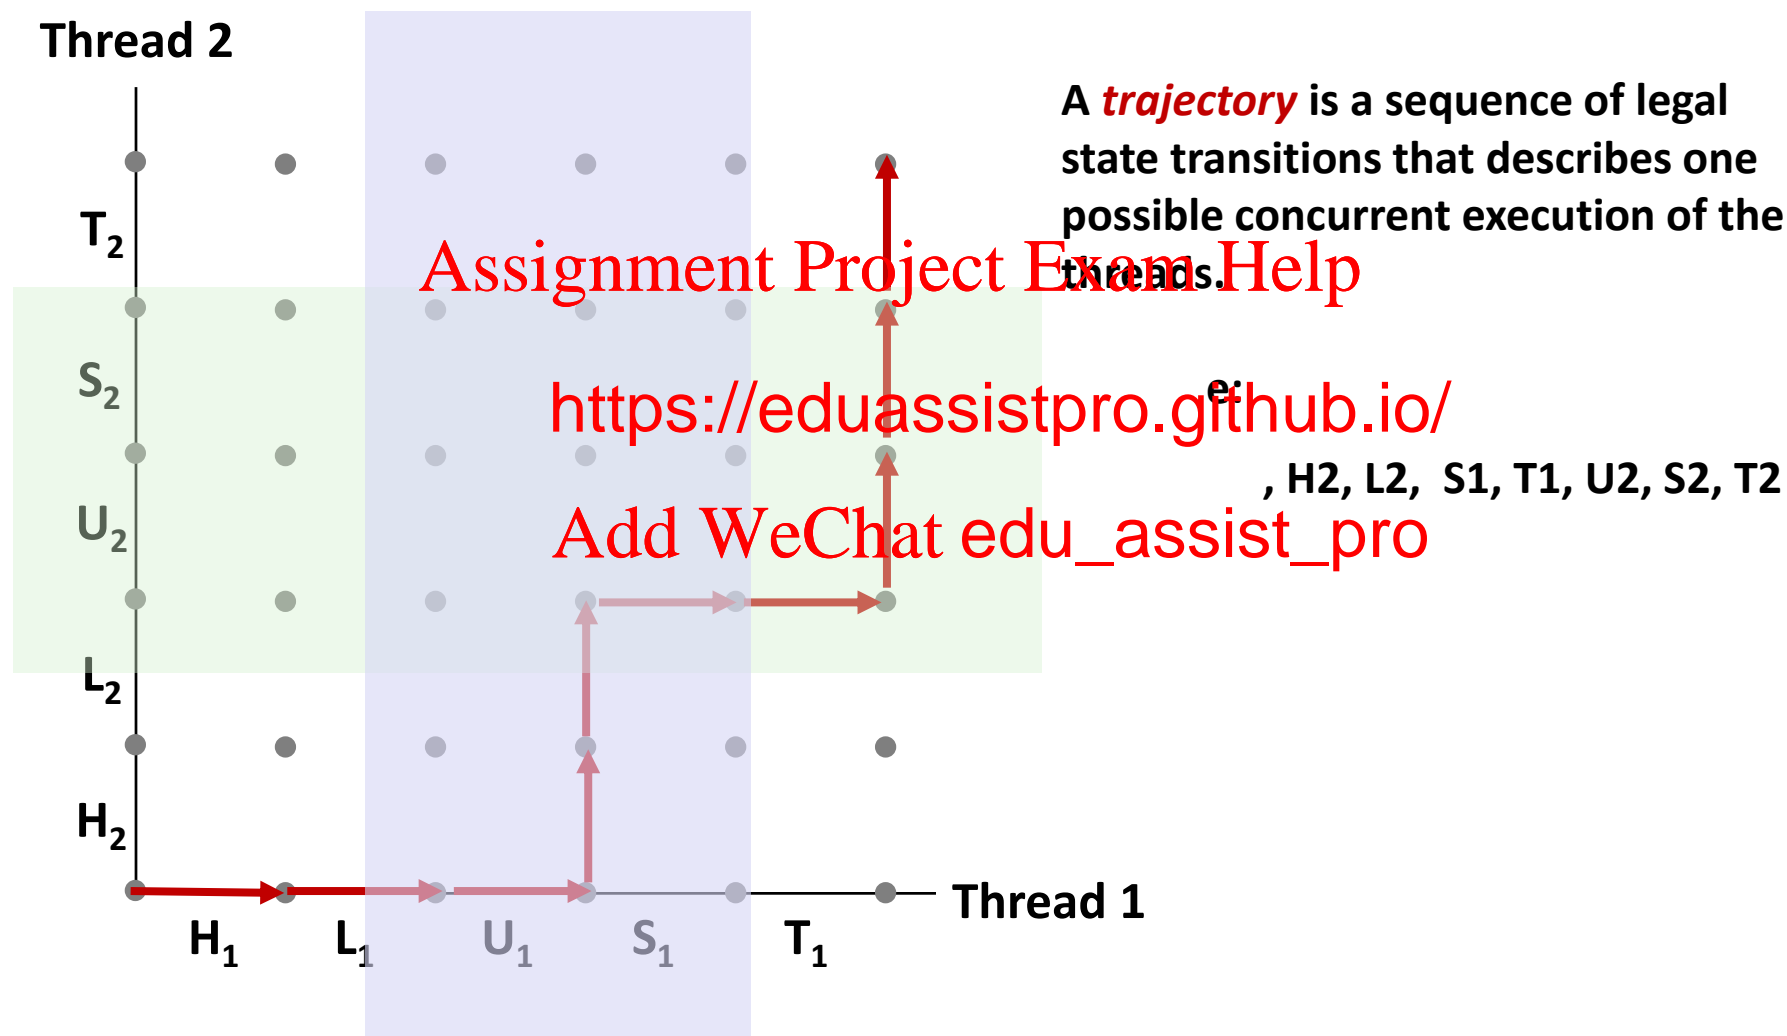
Assignment Project Exam Help

<https://eduassistpro.github.io/>

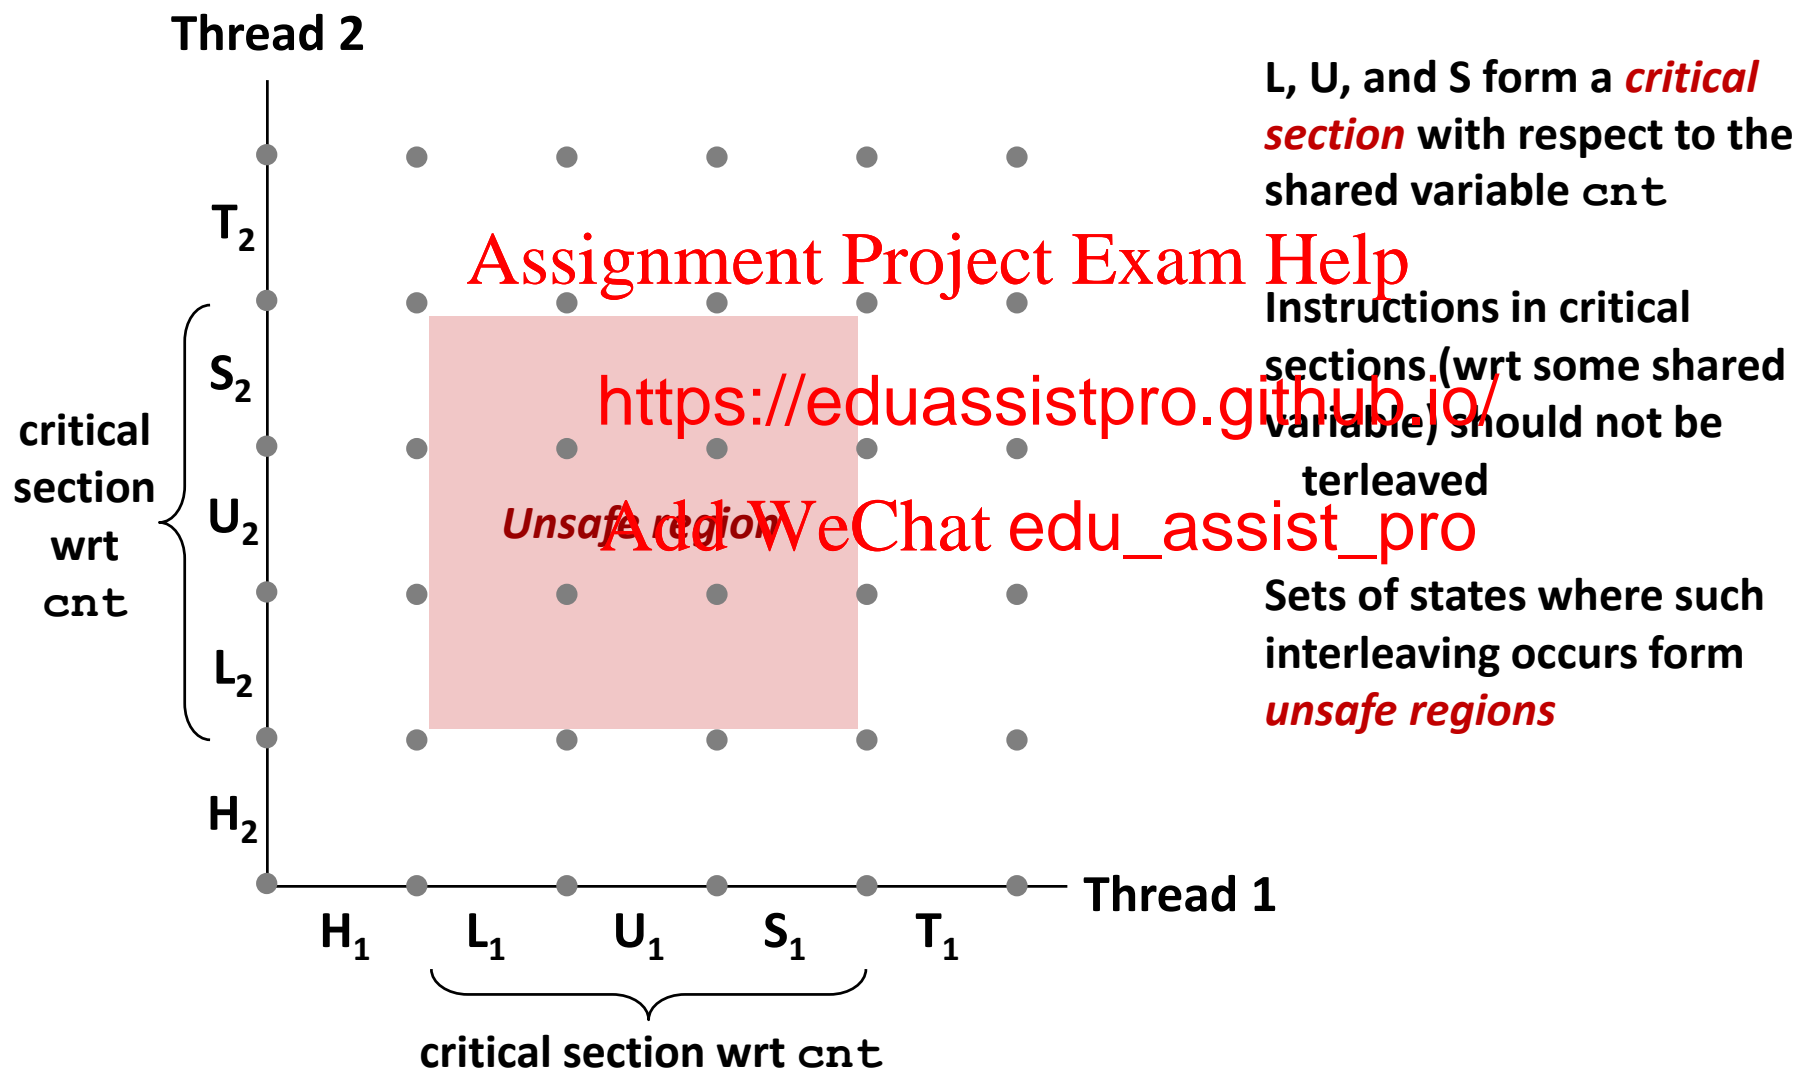
Add WeChat edu_assist_pro

, H2, L2, S1, T1, U2, S2, T2

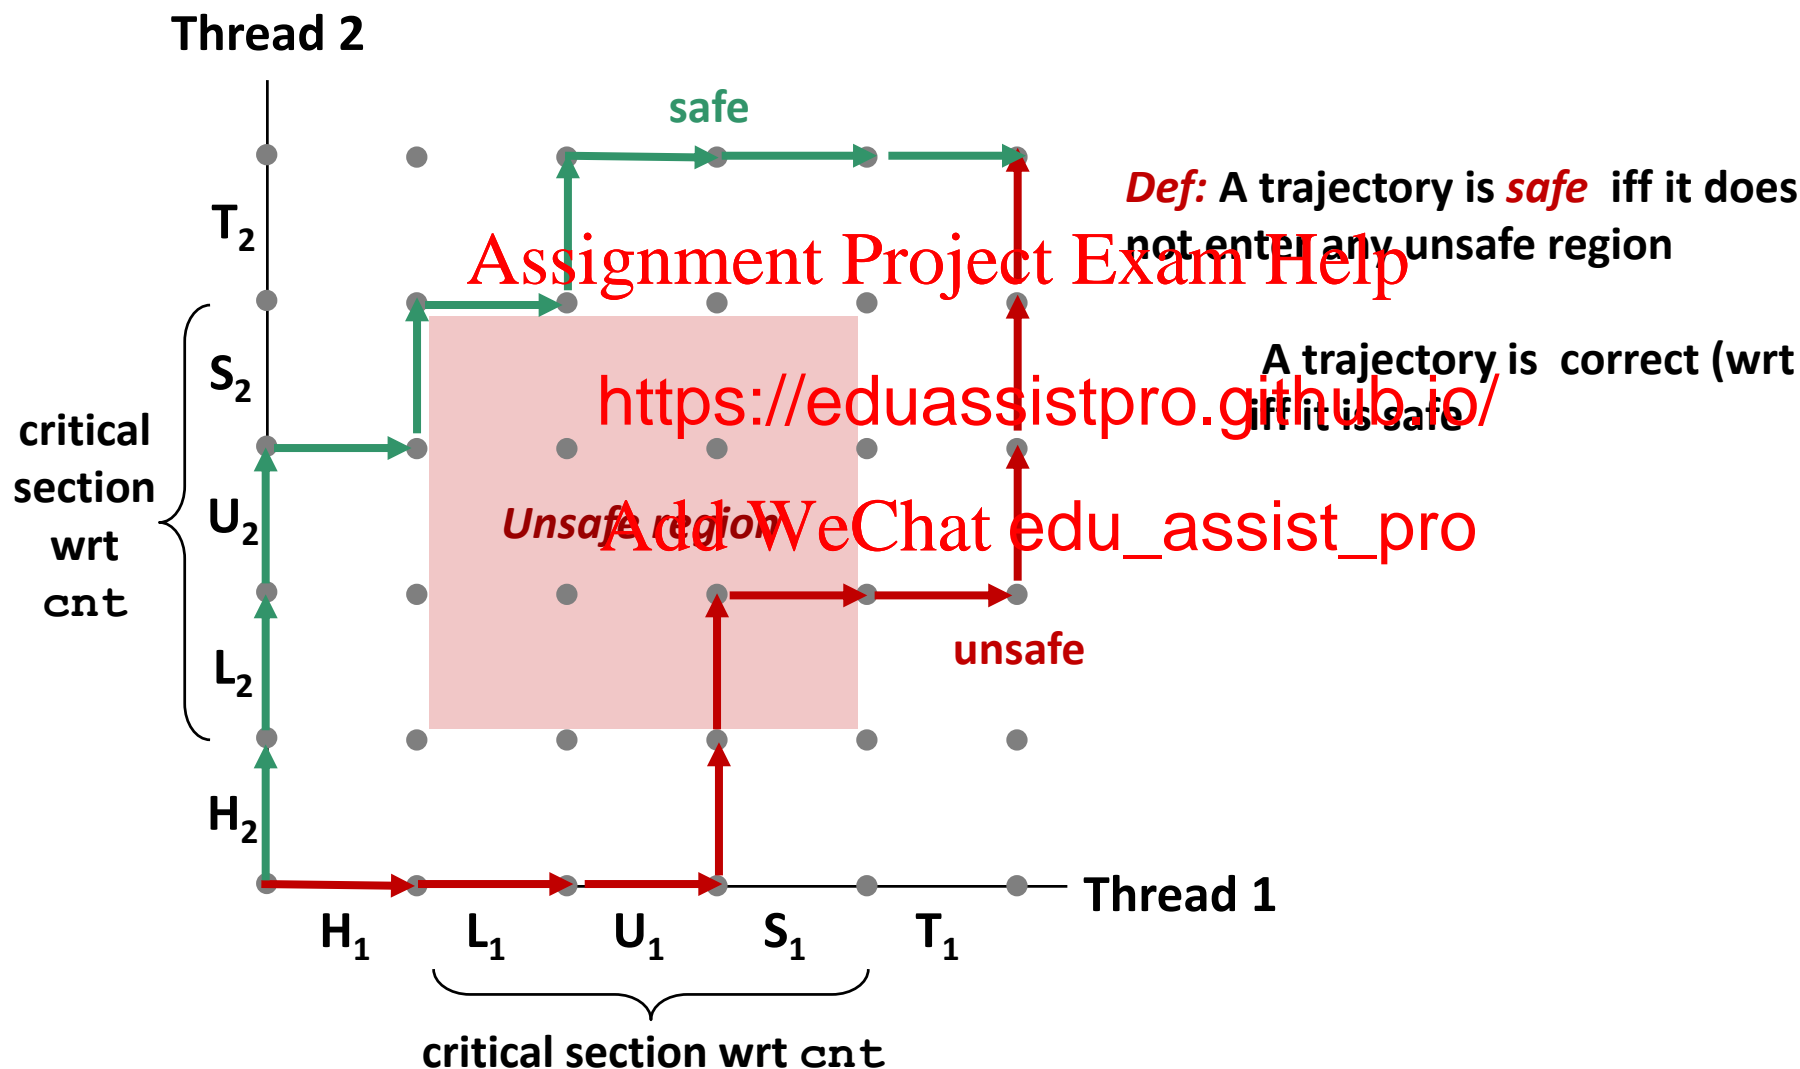
Trajectories in Progress Graphs



Critical Sections and Unsafe Regions



Critical Sections and Unsafe Regions



badcnt.c: Improper Synchronization

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
        thread, &niters);
    Pthread_create(&tid2, NULL,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);
    for (i = 0; i < niters; i++)
        cnt++;
    return NULL;
}

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

	main	thread1	thread2
cnt			
niters.m			
tid1.m			
i.1			
i.2			
niters.1			
niters.2			

badcnt.c: Improper Synchronization

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
        thread, &niters);
    Pthread_create(&tid2, NULL,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);
    for (i = 0; i < niters; i++)
        cnt++;
    return
        NULL;
}

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

	main	thread1	thread2
cnt	yes*	yes	yes
niters.m	yes	no	no
tid1.m	yes	no	no
i.1	no	yes	no
i.2	no	no	yes
niters.1	no	yes	no
niters.2	no	no	yes

Break Time!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Check out:

Add WeChat edu_assist_pro

Quiz: day 25: Synchronization Basic

<https://canvas.cmu.edu/courses/31656>

Bonus Quiz Question 6:

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
#include "csapp.h"
#define N 2
void *thread(void *vargp);
long *pointers[N];

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&tids[i], NULL, thread, (void *) i);
    sleep(1);    // Sleep-#1
    for (i = 0; i < N; i++)
        printf("Thread id %u has loca
               (int) tids[i], *pointers[i]);
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    return 0;
}

void *thread(void *vargp) {
    long myid = (long) vargp;
    pointers[myid] = &myid;
    sleep(2);    // Sleep-2
    return NULL;
}
```

If the statement labeled "Sleep #1" is kept, the main thread might have a segmentation fault when referencing "pointers"?

- True?
- False?

Today

- Threads review

- Sharing

- Mutual exclusion

- Semaphores

- Producer-Consumer <https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Assignment Project Exam Help

Enforcing Mutual Exclusion

- **Question:** How can we guarantee a safe trajectory?
- **Answer:** We must *synchronize* the execution of the threads so that they can never have an unsafe trajectory.
 - i.e., need to guarantee *mutually exclusive access* for each critical section.
- **Classic solution:**
 - Semaphores (Edsger Dijkstra)
- **Other approaches (out of our scope)**
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Semaphores

- ***Semaphore***: non-negative global integer synchronization variable. Manipulated by *P* and *V* operations.
- ***P(s)***
 - If *s* is nonzero, then decrement *s* by 1 and return immediately.
 - Test and decrement operations occur atomically (indivisibly)
 - If *s* is zero, then suspend thread until *s* becomes nonzero and the thread is restarted by a *V*
 - After restarting, thread returns control to the caller.
- ***V(s)***:
 - Increment *s* by 1.
 - Increment operation occurs atomically
 - If there are any threads blocked in a *P* operation waiting for *s* to become non-zero, then restart exactly one of those threads, which then completes its *P* operation by decrementing *s*.
- **Semaphore invariant: ($s \geq 0$)**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Semaphores

- ***Semaphore***: non-negative global integer synchronization variable
- **Manipulated by *P* and *V* operations:**
 - $P(s)$: [**while** ($s == 0$) **wait()**; $s--$;]
 - Dutch for "P"
 - $V(s)$: [$s++$;]
 - Dutch for "Verhogen" (increase)
- **OS kernel guarantees that operations between brackets [] are executed indivisibly**
 - Only one P or V operation at a time can modify s .
 - When **while** loop in P terminates, only that P can decrement s
- **Semaphore invariant: ($s \geq 0$)**

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat: edu_assist_pro

C Semaphore Operations

Pthreads functions:

```
#include <semaphore.h>

int sem_init(sem_t *s, unsigned int val); /* s = val */

int sem_wait(sem_t *s);
int sem_post(sem_t *s);
```

Assignment Project Exam Help
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

CS:APP wrapper functions:

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```

badcnt.c: Improper Synchronization

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1,
        thread, &niters);
    Pthread_create(&tid2,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i, niters =
        *((long *)vargp);
    for (i = 0; i < niters; i++)
        cnt++;
    return NULL;
}

```

How can we fix this using semaphores?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Using Semaphores for Mutual Exclusion

■ Basic idea:

- Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables).
- Surround corresponding critical sections with $P(mutex)$ and $V(mutex)$ operations

Assignment Project Exam Help

<https://eduassistpro.github.io/>

■ Terminology:

Add WeChat edu_assist_pro

- *Binary semaphore*: semaphore whose value is always 0 or 1
- *Mutex*: binary semaphore used for mutual exclusion
 - P operation: “locking” the mutex
 - V operation: “unlocking” or “releasing” the mutex
 - “Holding” a mutex: locked and not yet unlocked.
- *Counting semaphore*: used as a counter for set of available resources.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */

sem_init(&mutex, 0, 1); /* mutex = 1 */
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Surround critical

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

goodcnt.c

```
> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

Warning: It's orders of magnitude slower than badcnt.c.

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */

sem_init(&mutex, 0, 1); /* mutex = 1 */
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Surround critical

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```

Add WeChat edu_assist_pro

Func	cnt	goodcnt
Time (ms) niters = 10 ⁶	12.0	450.0
Slowdown	1.0	37.5

Warning: It's orders of magnitude slower than badcnt.c.

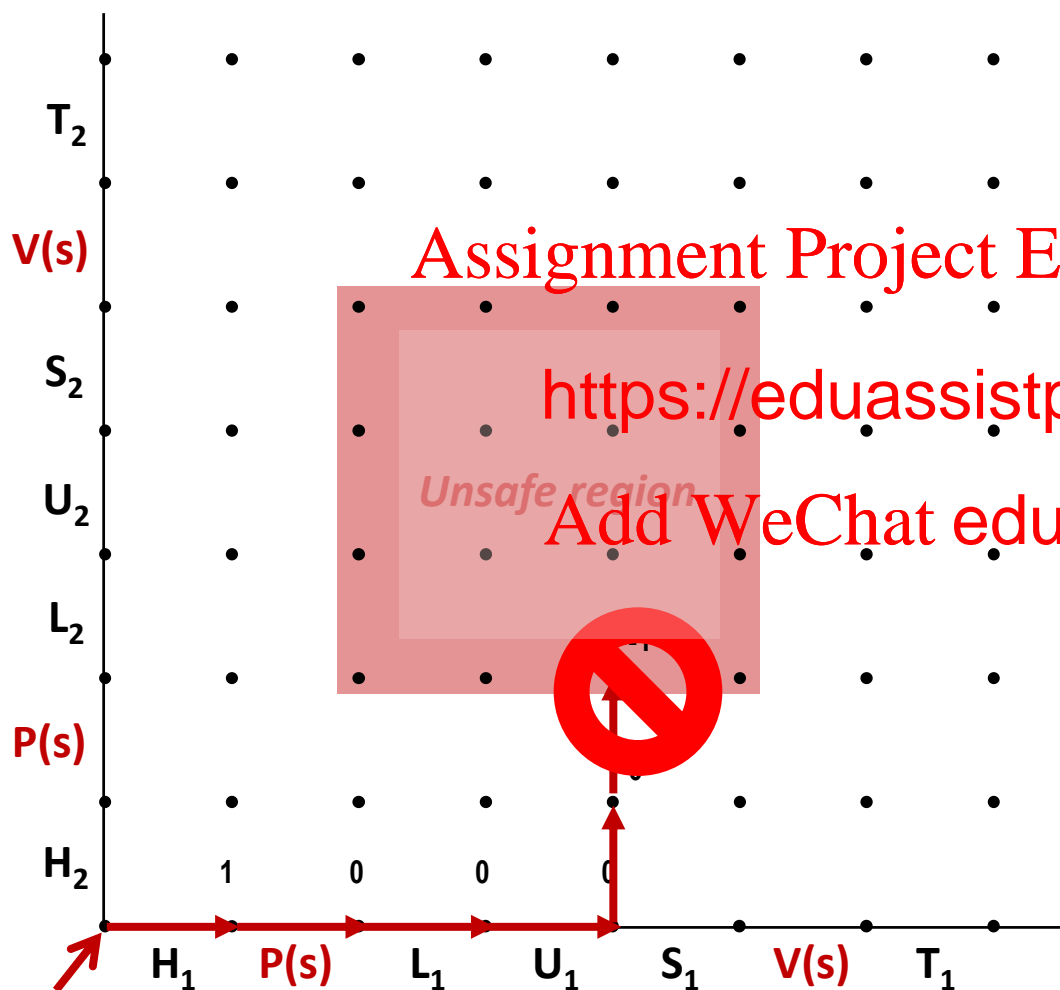
Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)



$s = 1$

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

semaphore invariant

creates a *forbidden region*

encloses unsafe region

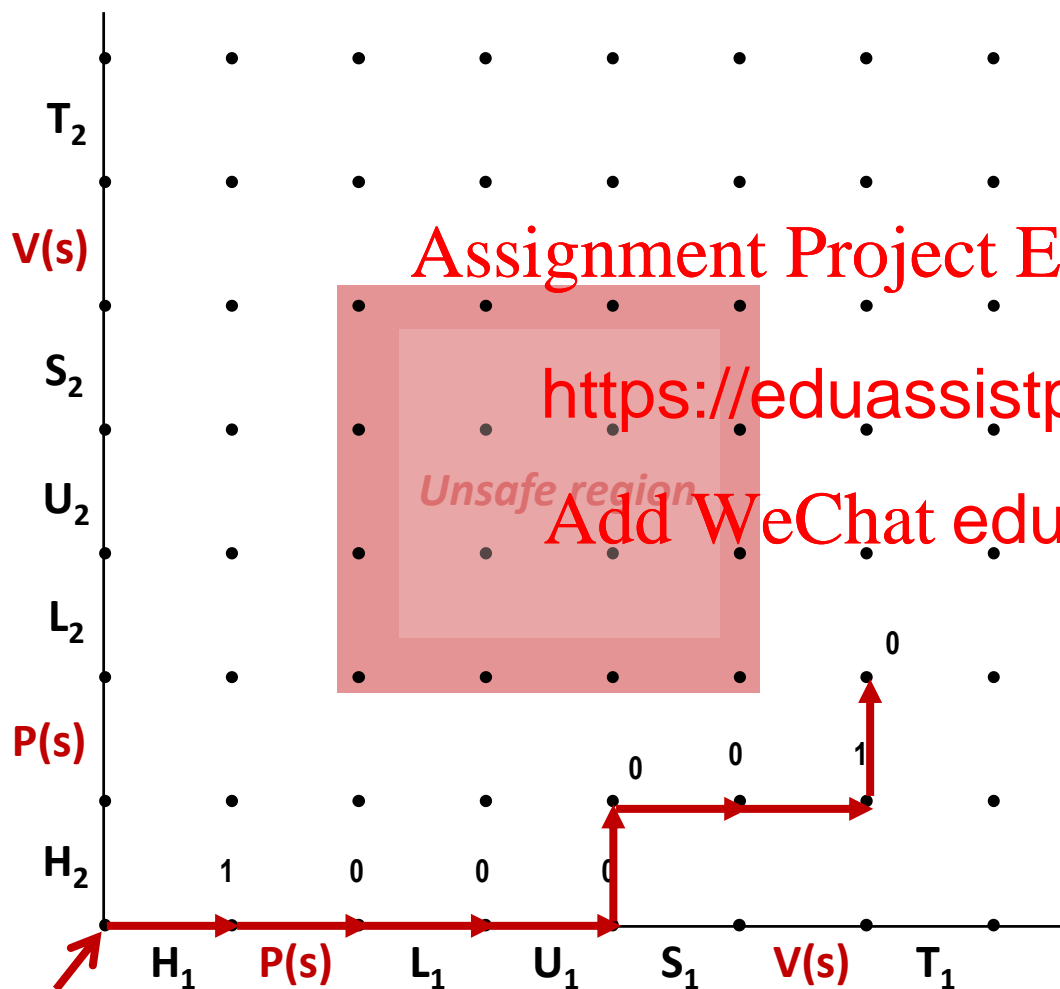
that cannot be entered by any trajectory.

Initially

$s = 1$

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

semaphore invariant

creates a *forbidden region*

encloses unsafe region

that cannot be entered by any trajectory.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

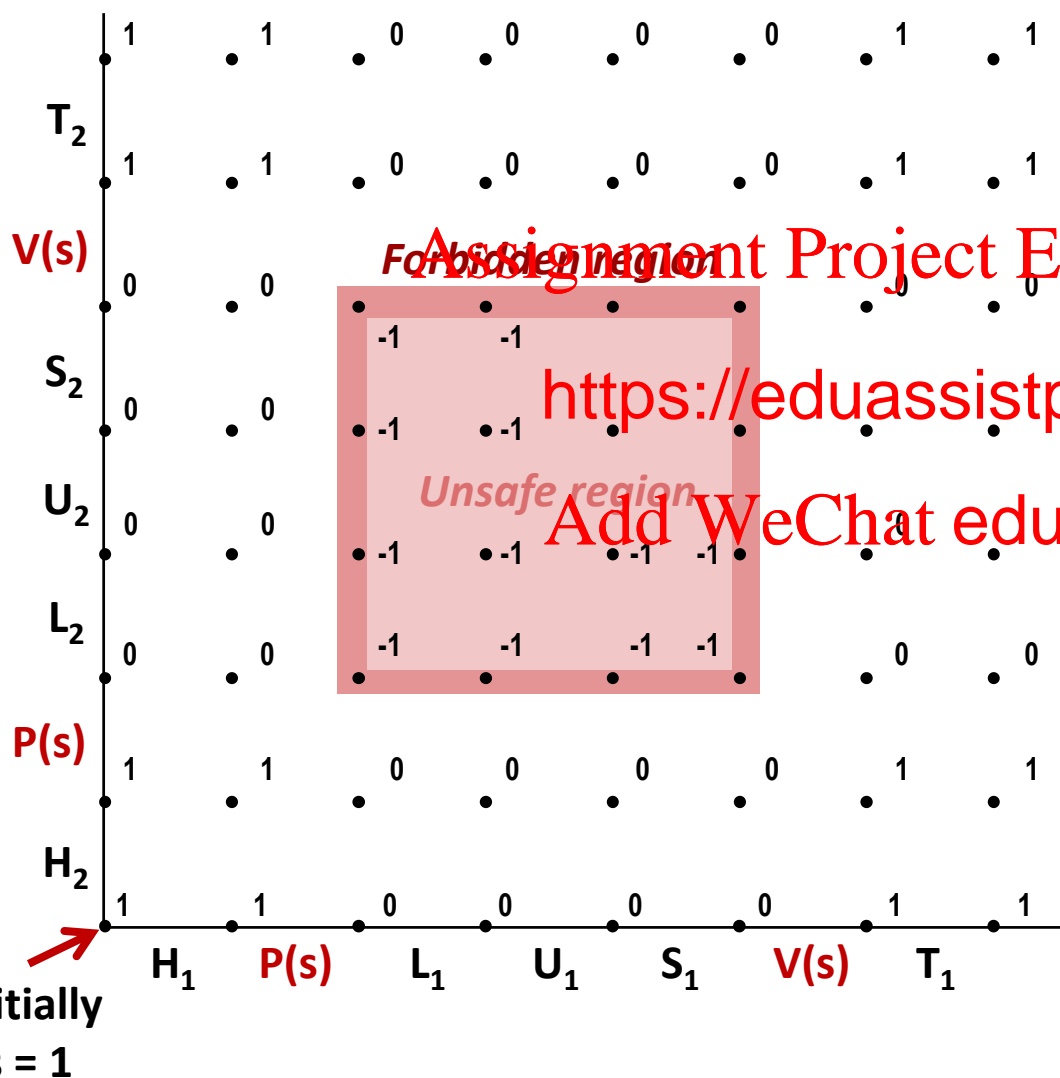
Unsafe region

Initially

$s = 1$

Why Mutexes Work

Thread 2



Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)

semaphore invariant

creates a *forbidden region*

encloses unsafe region

that cannot be entered by any trajectory.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Binary Semaphores – For Mutual Exclusion

- **Mutex is special case of semaphore**

- Value either 0 or 1

- **Pthreads provides `pthread_mutex_t`**

- Operations: lock, unlock

- **Recommended
appropriate**

Assignment Project Exam Help

<https://eduassistpro.github.io/> when

Add WeChat edu_assist_pro

goodmcent.c: Mutex Synchronization

- Define and initialize a mutex for the shared variable cnt:

```
volatile long cnt = 0; /* Counter */
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL); // No special attributes
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

- Surround critical

```
for (i = 0; i < niters; i++) {
    pthread_mutex_lock(&mutex);
    cnt++;
    pthread_mutex_unlock(&mutex);
}
```

```
> ./goodmcent 10000
OK cnt=20000
linux> ./goodmcent 10000
OK cnt=20000
```

Function	badcnt	goodcnt	goodmcent
Time (ms) niters = 10 ⁶	12.0	450.0	214.0
Slowdown	1.0	37.5	17.8

Today

- Threads review

- Sharing

- Mutual exclusion

- Semaphores

- **Producer-Consumer** <https://eduassistpro.github.io/>

Assignment Project Exam Help

Add WeChat edu_assist_pro

Using Semaphores to Coordinate Access to Shared Resources

- **Basic idea:** Thread uses a semaphore operation to notify another thread that some condition has become true

- Use counting semaphores to keep track of resource state.
- Use binary semaphore.

Assignment Project Exam Help


<https://eduassistpro.github.io/>

- **The Producer-Consumer Problem**

Add WeChat edu_assist_pro

- Mediating interactions between processes that generate information and that then make use of that information

Producer-Consumer Problem



Assignment Project Exam Help

■ Common synchronization pattern:

- Producer waits for buffer, and notifies consumer
- Consumer waits for producer, and notifies producer


Add WeChat edu_assist_pro

■ Examples

- Multimedia processing:
 - Producer creates video frames, consumer renders them
- Event-driven graphical user interfaces
 - Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - Consumer retrieves events from buffer and paints the display

Producer-Consumer on 1-element Buffer




- Maintain two semaphores: `full` + `empty`



Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



Producer-Consumer on 1-element Buffer

```
#include "csapp.h"

#define NITERS 5

void *producer(void *arg),
void *consumer(void

struct {
    int buf; /* shared var */
    sem_t full; /* sems */
    sem_t empty;
} shared;
```

```
int main(int argc, char** argv) {
    pthread_t tid_producer;
    pthread_t tid_consumer;

    /* Initialize the semaphores */
    red.empty, 0, 1);
    red.full, 0, 0);

    /* ... and wait */
    Pt &tid_producer, NULL,
    producer, NULL);
    Pthread_create(&tid_consumer, NULL,
    consumer, NULL);
    Pthread_join(tid_producer, NULL);
    Pthread_join(tid_consumer, NULL);

    return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Producer-Consumer on 1-element Buffer

Initially: `empty==1, full==0`

Producer Thread

```
void *producer(void *arg) {
    int i, item;

    for (i=0; i<NITER; i++) {
        /* Produce item */
        item = i;
        printf("produced %d\n",
              item);

        /* Write item to buf */
        P(&shared.empty);
        shared.buf = item;
        V(&shared.full);
    }
    return NULL;
}
```

Consumer Thread

```
void *consumer(void *arg) {
    int i, item;

    for (i=0; i<NITER; i++) {
        /* Consume item */
        item = shared.buf;
        V(&shared.empty);
        printf("consumed %d\n", item);
    }
    return NULL;
}
```

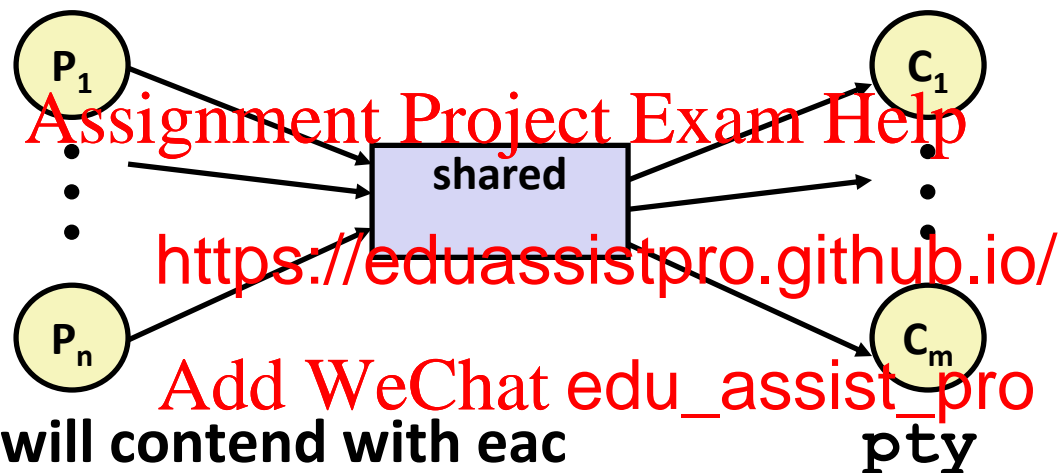
Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Why 2 Semaphores for 1-Entry Buffer?

- Consider multiple producers & multiple consumers



- Producers will contend with each other
- Consumers will contend with each other to get full


Producers

```
P(&shared.empty);
shared.buf = item;
V(&shared.full);
```

empty




full



Consumers

```
P(&shared.full);
item = shared.buf;
V(&shared.empty);
```

Producer-Consumer on an n -element Buffer




Add WeChat edu_assist_pro

- Implemented using a shared buffer package called `sbuf`.

Circular Buffer (n = 10)

- Store elements in array of size n
- items: number of elements in buffer
- Empty buffer:
 - front = rear
- Nonempty buffer
 - rear: index of most recently inserted
 - front: (index of next element to remove)
- Initially:




Assignment Project Exam Help

<https://eduassistpro.github.io/>


Add WeChat edu_assist_pro

Circular Buffer Operation (n = 10)


■ Insert 7 elements




■ Remove 5 elements



■ Insert 6 elements



■ Remove 8 elements



Sequential Circular Buffer Code

```
init(int v)
{
    items = front = rear = 0;
}
```

```
insert(int v)
{
    if (items >=
        error();
    if (++rear >=
        buf[rear] = v;
    items++;
}
```


```
int remove()
{
    if (items == 0)
        error();
    if (++front >= n) front = 0;
    int v = buf[front];
    items--;
    return v;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Producer-Consumer on an n -element Buffer



- **Requires a mutex and two counters:**
 - `mutex`: enforces mutually exclusive access to the buffer and counters
 - `slots`: counts the available slots in the buffer
 - `items`: counts the available items in the buffer
- **Makes use of general semaphores**
 - Will range in value from 0 to n

sbuf Package - Declarations

```
#include "csapp.h"

typedef struct {
    int *buf; /* Buffer array */
    int n; /* Maximum number of slots */
    int front; /* is first item */
    int rear; /* item */
    sem_t mutex; /* Protects access to buf */
    sem_t slots; /* Counts available slots */
    sem_t items; /* Counts available items */
} sbuf_t;

void sbuf_init(sbuf_t *sp, int n);
void sbuf_deinit(sbuf_t *sp);
void sbuf_insert(sbuf_t *sp, int item);
int sbuf_remove(sbuf_t *sp);
```

sbuf.h

sbuf Package - Implementation

Initializing and deinitializing a shared buffer:

```

/* Create an empty, bounded, shared FIFO buffer with n slots */
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n;
    sp->front = sp->rear = 0; /* front == rear */
    Sem_init(&sp->mutex, 0, 1); /* Bin for locking */
    Sem_init(&sp->slots, 0, n); /* Ini has n empty slots */
    Sem_init(&sp->items, 0, 0); /* Ini has zero items */
}

/* Clean up buffer sp */
void sbuf_deinit(sbuf_t *sp)
{
    Free(sp->buf);
}

```

sbuf.c

sbuf Package - Implementation

Inserting an item into a shared buffer:

```

/* Insert item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots); /* Acquire lock for available slot */
    P(&sp->mutex); /* Acquire lock for buffer */
    if (++sp->rear >= sp->n) /* In case of wrap-around, rear = 0 */
        sp->rear = 0;
    sp->buf[sp->rear] = item; /* Insert item into buffer */
    V(&sp->mutex); /* Unlock the buffer */
    V(&sp->items); /* Announce available item */
}

```

sbuf.c

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

sbuf Package - Implementation

Removing an item from a shared buffer:

```

/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items);
    P(&sp->mutex);
    if (++sp->front >=
        sp->front = 0;
    item = sp->buf[sp->front];
    V(&sp->mutex);
    V(&sp->slots);
    return item;
}

```

Assignment Project Exam Help
<https://eduassistpro.github.io/>
 Add WeChat edu_assist_pro

sbuf.c

Demonstration

- See program `produce-consume.c` in code directory
- 10-entry shared circular buffer
- 5 producers
 - Agent i generates numbers from $20*i$ to $20*i - 1$.
 - Puts them in bu
- 5 consumers
 - Each retrieves 20 elements from b
- Main program
 - Makes sure each value between 0 and 99 retrieved once

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Summary

- Programmers need a clear model of how variables are shared by threads.
- Variables shared by multiple threads must be protected to ensure mutual exclusion.
<https://eduassistpro.github.io/>
- Semaphores are a fundamental tool for enforcing mutual exclusion.
Add WeChat edu_assist_pro